

Relational Algebra

$$y = g(x)$$

Secant Lines

Tangent line

$$x+h$$

$$f'(x) = \lim_{h \rightarrow 0} f(x+h) - f(x)$$

$$f'(x) = \lim_{h \rightarrow 0} \frac{(x+h)^2 + 2(x+h)}{h}$$

$$= \lim_{h \rightarrow 0} \frac{x^2 + 2xh + h^2 + 2x + 2h}{h}$$

$$= \lim_{h \rightarrow 0} \frac{2xh + h^2 + 2h}{h}$$

$$\frac{g(x+h) - g(x)}{h}$$

$$= \lim_{h \rightarrow 0} h(2x + h)$$

Key = attribute set that **uniquely** identifies each row in the relation/table



Keys

Superkey

- set of attributes of table for which every row has distinct set of values

Candidate key

- Minimal such set of attributes

Primary key

- DB Chosen Candidate key
- Plays a very important role
 - E.g. relations typically sorted by this

醫事機構代碼	醫事機構名稱	醫事機構地址	醫事機構電話	成人口罩剩餘數	兒童口
145080011	衛生福利部花蓮醫院豐濱	花蓮縣豐濱鄉豐濱村光豐路 41 號	(03)8358141	3304	
291010010	連江縣立醫院	連江縣南竿鄉復興村 217 號	(083)623995	0	
2101010013	松山健康服務中心	臺北市松山區八德路 4 段 694 號 1、2 樓	(02)27653147	1	
2101020019	大安健康服務中心	臺北市大安區辛亥路 3 段 15 號	(02)27390997	0	
2101090011	大同健康服務中心	臺北市大同區昌吉街 52 號 1 至 2 樓	(02)25948971	0	
2101100227	中山健康服務中心	臺北市中山區松江路 367 號 1 樓	(02)25013363	2	
2101110027	內湖健康服務中心	臺北市內湖區民權東路 6 段 99 號	(02)27908387	0	
2101120014	南港健康服務中心	臺北市南港區南港路 1 段 360 號 1 樓	(02)27868756	0	
2101150012	士林健康服務中心	臺北市士林區中正路 439 號 1 樓	(02)28836268	0	
2101161033	北投健康服務中心	臺北市北投區石牌路 2 段 111 號	(02)28261026	0	
2101170050	信義健康服務中心	臺北市信義區信義路 5 段 15 號	(02)87804152	0	
2101180038	中正健康服務中心	臺北市中正區牯嶺街 24 號	(02)23210168	0	
2101191068	萬華健康服務中心	臺北市萬華區東園街 152 號	(02)23395384	0	
2101200017	政大(文山健康服務中心)	臺北市文山區木柵路 3 段 220 號	(02)22343501	0	
2303240012	臺北市立聯合醫院	臺北市大同區中華路二段 222 號	(02)22223811	2028	

Example of FaceMask dataset

- Superkey
 - {代碼}, {名稱}, {電話}, {代碼, 名稱}, {代碼, 電話}, {代碼, 口罩剩餘數}
- Candidate key:
 - {代碼, 電話} is not a miminal set. Since {電話} could be removed from {代碼, 電話} while the remining attribute {代碼} is still a key
 - {代碼}, {名稱}, {電話}, {地址}

Introduction

- Query languages are specialized languages for asking questions or queries, that involve the data in a database.
- Relational algebra : queries in terms of operators
- Every operator in relational algebra accepts (one or two) relation instances as arguments and returns a relation instance as the result.

$$(1 + 2 = 3) \quad +(1,2) = 3$$



Introduction

$$(1 + 2 = 3) \quad +(1,2) = 3$$

- A relational algebra expression is recursively defined to be a relation.
- Relational algebra is a procedural query language
 - Defines a step-by-step procedure for computing the answer.

Basic Operators

- There are six basic operators in relation algebra:
 - select (σ)
 - project (Π)
 - union (U)
 - set different ($-$)
 - Cartesian product (x)
 - rename (ρ)

$$y = g(x)$$

Secant Lines

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$
$$= \lim_{h \rightarrow 0} \frac{x^2 + 2xh + h^2 - x^2}{h}$$
$$f'(x) = \lim_{h \rightarrow 0} \frac{g(x+h) - g(x)}{h}$$
$$= \lim_{h \rightarrow 0} \frac{h(2x+h)}{h}$$

Selection (σ)

σ

Relation R

a	b	c	d
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

$\sigma_{a=b \wedge d>5} (R)$

a	b	c	d
α	α	1	7
β	β	23	10

Selection (σ)

- The selection operation σ specifies the tuples to retain through a *selection condition*.
- Selection condition is a Boolean combination (an expression using logical connectives \wedge and \vee) of terms that have the form
 - attribute op constant (e.g. $A = 3$)
 - attribute1 op attribute2 (e.g. $X < Y$)
(op is one of the comparison operators
 $<$, \leq , $=$, \neq , $>$, \geq).

Projection (Π)

Π

Relation R

a	b	c
α	10	1
α	20	1
β	30	1
β	40	2

$\Pi_{a,c} (R)$

a	c
α	1
α	1
β	1
β	2



Eliminate
Duplicated
rows

a	c
α	1
β	1
β	2

Projection (Π)

- The projection operator Π allows us to extract columns from a relation.
- The subscript specifies the fields to be retained.
(The other fields are ‘projected out’).
- The schema of the result of a projection is determined by the fields that are projected.
- Duplicated row will be eliminated in the final result.
 - This follows from the definition of a relation as a **set** of tuples.

Employee

f_name	l_name	id	gender	salary	superid	dno
Joseph	Chan	999999	M	29500	654321	4
Victor	Wong	001100	M	30000	888555	5
Carrie	Kwan	898989	F	26000	654321	4
Joyce	Fong	345345	F	12000	777888	4

Find all employees who works in department 4 and whose salary is greater than 25000.

$\sigma_{dno=4 \wedge salary > 25000} (\text{Employee})$

f_name	l_name	id	gender	salary	superid	dno
Joseph	Chan	999999	M	29500	654321	4
Carrie	Kwan	898989	F	26000	654321	4

- Find the employee names and department number of all employees
- Find the gender and department number of all employees

$\Pi_{l_name, f_name, dno} (\text{Employee})$

f_name	l_name	dno
Joseph	Chan	4
Victor	Wong	5
Carrie	Kwan	4
Joyce	Fong	4

$\Pi_{gender, dno} (\text{Employee})$

gender	dno
M	4
M	5
F	4

- Compositing operations
- Substituting an expression where a relation is expected
- The result of an relational-algebra expression is always a relation.

$$\Pi_{\text{lname}, \text{fname}} (\sigma_{\text{dno}=4 \wedge \text{salary} > 25000} (\text{Employee}))$$

f_name	l_name
Joseph	Chan
Carrie	Kwan

f_name	l_name	id	gender	salary	superid	dno
Joseph	Chan	999999	M	29500	654321	4
Victor	Wong	001100	M	30000	888555	5
Carrie	Kwan	898989	F	26000	654321	4
Joyce	Fong	345345	F	12000	777888	4

Union

R

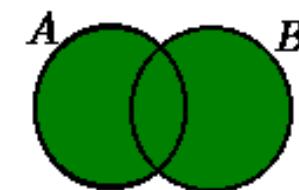
	a	b
a		1
α		2
β		1

S

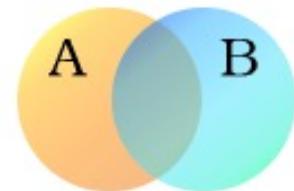
	a	b
α		2
β		3

R U S

	a	b
α		1
α		2
β		1
β		3



- $R \cup S$ returns a relation instance containing all tuples that occur in either relation R or relation S (or both).
- The union operation is **commutative**: $R \cup S = S \cup R$
- Duplicated tuples are eliminated.
- R and S must be *union-compatible*.
 - They have the same number of fields,
 - The corresponding fields have the same domains.
 - The field names are not so important
- The schema of the result is defined to be identical to the schema of R.
 - The fields of $R \cup S$ inherit names from R.

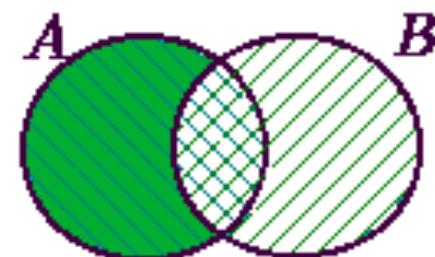


Set Difference

	a	b
R		
α	1	
α	2	
β	1	

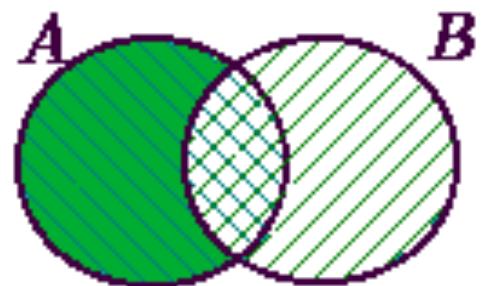
	a	b
S		
α	2	
β	3	

	a	b
$R - S$		
α	1	
β	1	



Set Difference

- Set difference $R - S$
returns a relation instance containing all the tuples that occur in R but not in S.
- Set difference is not commutative:
in general, $R - S \neq S - R$.
- The relations R and S must be union-compatible.
- The schema of the result is defined to be identical to the schema of R.



Intersection

R

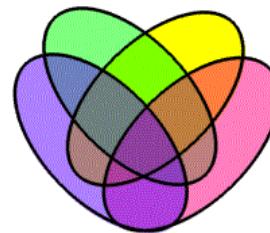
	a	b
a	1	
a	2	
β	1	

S

	a	b
α	2	
β	3	

$R \cap S$

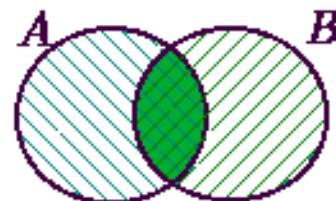
	a	b
α	2	



Intersection

- $R \cap S$ returns a relation instance containing all tuples that occur in both R and S.
- The relations R and S must be union-compatible.
- The schema of the result is defined to be identical to the schema of R.
- Intersection is not considered a basic operation, as it can be derived from the basic operations as shown below:

$$R \cap S = R - (R - S)$$



Cartesian-Product (\times)

$R \times S$

R	
a	b
α	1
β	2

S		
c	d	e
α	10	+
β	10	+
β	20	-
γ	10	-

a	b	c	d	e
α	1	α	10	+
α	1	β	10	+
α	1	β	20	-
α	1	γ	10	-
β	2	α	10	+
β	2	β	10	+
β	2	β	20	-
β	2	γ	10	-

Cartesian product (Cross product)

- $R \times S$ returns a relation instance whose schema contains all the fields of R followed by all the fields of S.
- The result contains one tuple $\langle r, s \rangle$ (concatenation of tuples r and s) for each pair of tuples $r \in R, s \in S$.

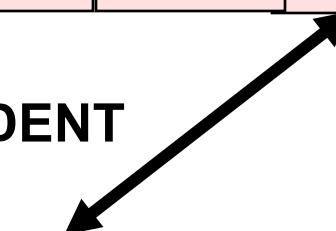
Query: for each **female** employee, retrieve employee's name and a list of the names of her dependents.

EMPLOYEE

f_name	l_name	<u>id</u>	bdate	addr	gender	salary	super_id	dno
--------	--------	-----------	-------	------	--------	--------	----------	-----

DEPENDENT

<u>eid</u>	<u>dependent-name</u>	gender	bdate	relationship
------------	-----------------------	--------	-------	--------------



Female_emps $\leftarrow \sigma_{\text{gender}='F'}$ (Employee)

Empnames $\leftarrow \Pi_{\text{fname}, \text{lname}, \text{id}}$ (Female_emps)

Female_emps

f_name	l_name	id	bdate	address	gender	salary	superid	dnc
Alicia	Chan	998877	2-Jul-70	231, Cai Road	F	9500	654321	4
Jennifer	Wong	654321	20-June-60	342, Cheung Road	F	30000	888555	4
Joyce	Fong	345345	19-Dec-80	23, Young Road	F	12000	777888	5

Empnames

f_name	l_name	id
Alicia	Chan	998877
Jennifer	Wong	654321
Joyce	Fong	345345

Dependents

eid	dep_name	gender	bdate	relationship
334455	Alice	F	5-Apr-90	Daughter
334455	Theodore	M	3-Mar-92	Son
654321	Abner	M	29-Feb-94	Son
123456	Alice	F	2-Nov-97	Daughter

`Emp_dependents ← Empnames x Dependents`

Emp_dependents

f_name	l_name	id	eid	dep_name	gender	bdate	relationship
Alicia	Chan	998877	334455	Alice	F	5-Apr-90	Daughter
Alicia	Chan	998877	334455	Theodore	M	3-Mar-92	Son
Alicia	Chan	998877	654321	Abner	M	29-Feb-94	Son
Alicia	Chan	998877	123456	Alice	F	2-Nov-97	Daughter
Jennifer	Wong	654321	334455	Alice	F	5-Apr-90	Daughter
Jennifer	Wong	654321	334455	Theodore	M	3-Mar-92	Son
Jennifer	Wong	654321	654321	Abner	M	29-Feb-94	Son
Jennifer	Wong	654321	123456	Alice	F	2-Nov-97	Daughter
Joyce	Fong	345345	334455	Alice	F	5-Apr-90	Daughter
Joyce	Fong	345345	334455	Theodore	M	3-Mar-92	Son
Joyce	Fong	345345	654321	Abner	M	29-Feb-94	Son
Joyce	Fong	345345	123456	Alice	F	2-Nov-97	Daughter

Query: for each female employee, retrieve employee's name and a list of the names of her dependents

Actual_dependents $\leftarrow \sigma_{id=eid} (\text{Emps_dependents})$

Result $\leftarrow \Pi_{\text{fname}, \text{lname}, \text{dep_name}} (\text{Actual_dependents})$

Actual_dependents

f_name	l_name	id	eid	dep_name	gender	bdate	relationship
Jennifer	Wong	654321	654321	Abner	M	29-Feb-94	Son

Result

f_name	l_name	dep_name
Jennifer	Wong	Abner

Join

- Because the sequence of operations, π followed by σ is quite common, a special operation, called the “join” operation  was created to specify these as a single operation.
- Join can be defined as a cross-product followed by selections and sometimes with projections.

Query: for each **female** employee, retrieve employee's name and a list of the names of her dependents.

Female_emps $\leftarrow \sigma_{\text{gender}='F'} (\text{Employee})$

Emphnames $\leftarrow \pi_{\text{fname}, \text{lname}, \text{id}} (\text{Female_emps})$

Actual_dependents $\leftarrow \text{emphnames} \bowtie_{\text{id}=\text{eid}} \text{Dependents}$

Result $\leftarrow \pi_{\text{fname}, \text{lname}, \text{dep_name}} (\text{Actual_dependents})$

Female_emps

f_name	l_name	id	bdate	address	gender	salary	superid	dno
Alicia	Chan	998877	2-Jul-70	231, Cai Road, HK	F	9500	654321	4
Jennifer	Wong	654321	20-June-60	342, Cheung Road, HK	F	30000	888555	4
Joyce	Fong	345345	19-Dec-80	23, Young Road, HK	F	12000	777888	5

Emppnames



Dependents

f_name	l_name	id
Alicia	Chan	998877
Jennifer	Wong	654321
Joyce	Fong	345345

eid	dep_name	gender	bdate	relationship
334455	Alice	F	5-Apr-90	Daughter
334455	Theodore	M	3-Mar-92	Son
654321	Abner	M	29-Feb-94	Son
123456	Alice	F	2-Nov-97	Daughter

Actual_dependents

f_name	l_name	id	eid	dep_name	gender	bdate	relationship
Jennifer	Wong	654321	654321	Abner	M	29-Feb-94	Son

Result

f_name	l_name	dep_name
Jennifer	Wong	Abner

Conditional Join

- The most general version of join operation accepts a **join condition c**.
- The join condition is identical to a selection condition in form.
- The operation is defined as follows:

$$R \bowtie_C S = \sigma_C(R \times S)$$

S

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.5

R

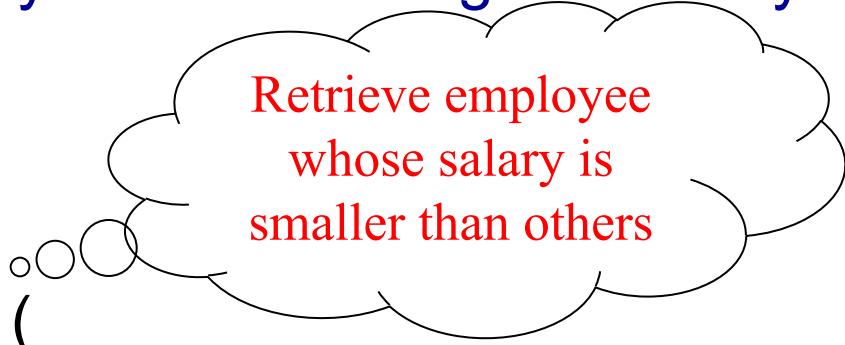
sid	bid	day
22	101	10/10/96
58	103	11/12/96

$$S \bowtie_{S.sid < R.sid} R$$

(sid)	sname	rating	age	(sid)	bid	day
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.5	58	103	11/12/96

Example

Query: Find the names of employees with the highest salary.

$$\Pi_{lname, fname} (\text{Employee}) -$$
$$\Pi_{Employee.lname, Employee.fname} ($$
$$\text{Employee} \bowtie_{Employee.salary < F.salary} \rho (F, \text{Employee}))$$
$$\text{Employee} \bowtie_{Employee.salary > F.salary} \rho (F, \text{Employee})$$


Retrieve employee
whose salary is
smaller than others

Cannot work ? Why ?

Employee

eid	Iname	salary
22	Dustin	7
31	Lubber	8
58	Rusty	10

Employee as F

eid	Iname	salary
22	Dustin	7
31	Lubber	8
58	Rusty	10

Employee \bowtie Employee.salary < F.salary F

eid	Iname	salary	(eid)	Iname	salary
22	Dustin	7	31	Lubber	8
31	Dustin	7	58	Rusty	10
31	Lubber	8	58	Rusty	10

Equi-Join

A condition join where the condition contains only equalities.

The conditions of the join have the form $R.name1 = S.name2$, where we want to join R with S .

In the resulting relation, $S.name2$ will be dropped.

	a	b	c	
R	α	1	α	
	β	5	γ	
	γ	4	β	
	α	1	γ	
	δ	2	β	

	b	e	f	
S	1	X	α	
	3	X	β	
	1	X	γ	
	2	Y	δ	
	3	Y	ϵ	

$R \bowtie_{R.b=S.b} S$

	a	b	c	e	f
α	1	α	X	α	
α	1	α	X	γ	
α	1	γ	X	α	
α	1	γ	X	γ	
δ	2	β	Y	δ	

Natural Join

An Equijoin on all common fields.

For $R \bowtie S$, the resulting schema contains the attributes of R followed by the attributes in S that are not in R .

If two relations have **no attributes in common**, the result is simply the **cross-product**.

	R		
a	b	c	d
α	1	α	X
β	2	γ	X
γ	4	β	γ
α	1	γ	γ
δ	2	β	γ

	S		e
	b	d	e
1		X	α
3	X		β
1	X		γ
2	y		δ
3	y		ϵ

$R \bowtie S$				
a	b	c	d	e
α	1	α	X	α
α	1	α	X	γ
δ	2	β	γ	δ

Division

A	x	y
	α	1
	α	2
	α	3
	β	1
	γ	1
	δ	1
	δ	3
	δ	4
	ε	1
	ε	2

B	y
	1
	2

A / B	x
	α
	ε

The division operation A/B is the set of all x values
(in the form of unary tuples)

such that for every y value in a tuple of B ,
there is a tuple $\langle x, y \rangle$ in A .

A

x	y
α	1
α	2
α	3
β	1
γ	1
δ	1
δ	3
δ	4
ε	1
ε	2

B

y
1
2

A / B

x
α
ε

Another way to understand division:

For each x value in A , consider the set of y values that appear in tuples of A with that x value.

If this set contains all y values in B , the x value is in the result of A/B .

Examples of Division A/B

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

A

pno
p2

B1

pno
p2
p4

B2

pno
p1
p2
p4

B3

sno
s1
s2
s3
s4

A/B1

sno
s1
s4

A/B2

sno
s1

A/B3

Employee	fname	Iname	id	bdate	address	salary	sid	dno
----------	-------	-------	----	-------	---------	--------	-----	-----

Works_on	id	pno
----------	----	-----

Query: Retrieve the names of employees who work on **all** the projects that 'Chris Peng' works on.

$PENG \leftarrow \sigma_{\text{fname}=\text{"Chris"} \wedge \text{Iname}=\text{"Peng"}} (\text{Employee})$

$PENG_pnos \leftarrow \Pi_{pno} (\text{Works_on} \bowtie PENG)$

fname	Iname	id	bdate	address	salary	sid	dno	pno	pno
-------	-------	----	-------	---------	--------	-----	-----	-----	-----

$\text{Result_id} \leftarrow \text{Works_on} / PENG_pnos$

$\text{Result} \leftarrow \Pi_{\text{fname}, \text{Iname}} (\text{Result_id} \bowtie \text{Employee})$



User (ID, PID, Name)

Inbound (PID, Origin, flight_no, date)

Airport (IATA code, city, country)

Query 1: Find PID of users who traveled UK and were back Taiwan on 2020/03/16

— Solution 1

$$\Pi_{PID} (\sigma_{country=UK} (Airport) \bowtie \sigma_{date=2020/03/16} (Inbound))$$

— Solution 2:

$$\Pi_{PID} (\sigma_{country=UK \text{ and } date=2020/03/16} (Airport \bowtie Inbound))$$

Query 2: Find the ID of users who traveled to UK and were back Taiwan on 2020/03/16

— Solution 1:

$$\Pi_{ID} ((\sigma_{country=UK} (Airport) \bowtie \sigma_{date=2020/03/16} (Inbound)) \bowtie User)$$

— Solution 2:

$$\Pi_{ID} (\Pi_{PID} ((\Pi_{LATA\ code} (\sigma_{country=UK} Airport)) \bowtie \sigma_{date=2020/03/16} (Inbound)) \bowtie User)$$

Query 3: Find the ID of users who were back from UK or Japan

– Solution 1:

$$\rho (\text{TempAirport}, (\sigma_{\text{country}=\text{UK} \vee \text{country}=\text{Japan}} \text{Airport})) \\ \Pi_{\text{ID}} (\text{TempAirport} \bowtie \text{Inbound} \bowtie \text{User})$$

– Solution 2

$$\Pi_{\text{ID}} (((\sigma_{\text{country}=\text{UK}} \text{Airport}) \bowtie \text{Inbound}) \bowtie \text{User}) \cup \\ \Pi_{\text{ID}} (((\sigma_{\text{country}=\text{Japan}} \text{Airport}) \bowtie \text{Inbound}) \bowtie \text{User})$$

Query 4: Find the PID of users who were in the same plant of one confirmed case with his PID as TW0001 on 2020/02/28

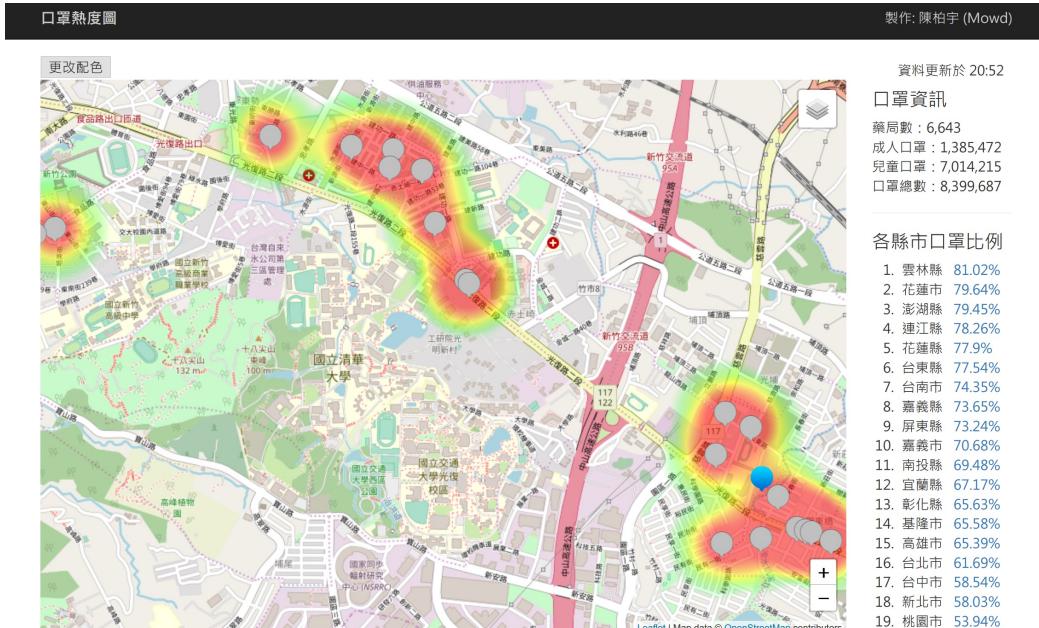
-Solution 1

$$\Pi_{PID, flight_no} (\sigma_{Date=2020/02/28} Inbound) / \Pi_{flight_no} (\sigma_{PID=TW001 \text{ and } Date=2020/02/28} Inbound)$$

-Any other solution ?

Covid 2019 HackMD

Open Data Taiwan FB Group



Additional Relational Operations

- AGGREGATE FUNCTIONS

- Functions such as SUM, COUNT, AVERAGE, MIN, MAX are often applied to sets of values or sets of tuples in database applications
- <**grouping attributes**> F <function list> (R)
 - The grouping attributes are optional
 - e.g. Retrieve the average salary of all employees (no grouping needed):

$R(\text{AVGSAL}) \leftarrow F_{\text{AVERAGE SALARY}} (\text{EMPLOYEE})$

Query 1: Find the average number of adult face mask in each clinic.

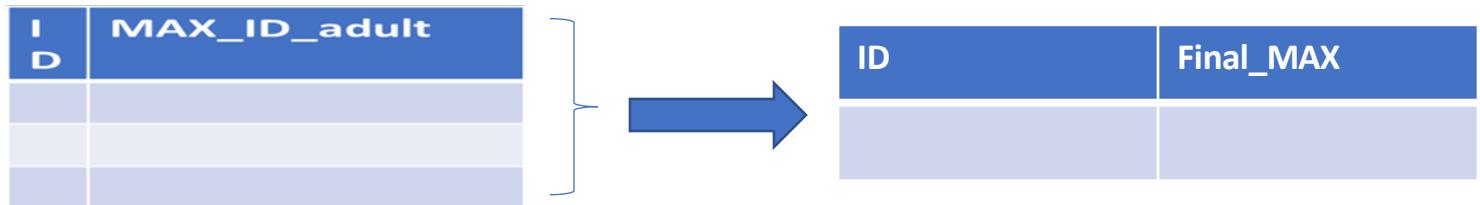
$$R(ID, \text{ave_adult}) =_{ID} F_{\text{avg}(\text{adult_no})} (\text{Maskdata})$$

Query 2: Find the phone number of the clinic ID that has the maximal number of adult face mask

Step 1: $R(ID, MAX_ID_adult) =_{ID} F_{MAX(adult_no)}(Maskdata)$

ID	MAX_ID_adult

Step 2: $R(ID, Final_MAX) =_{ID} F_{MAX(MAX_ID_adult)}(Maskdata)$



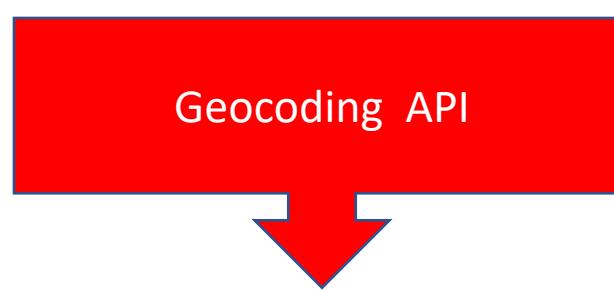
Step 3: $\Pi_{tel}(R(ID, Final_MAX) \bowtie Maskdata)$

Query 3: Given one location (longitude, latitude), find the nearest clinic ID that has face mask

	field1	field2
1	過濾	過濾
2	2312010014	新竹市東區衛生所
2	2312041028	新竹市北區衛生所
3	2312050018	新竹市香山衛生所
4	5912010046	中美藥局
5	5912010206	康寧藥局
6	5912010215	李藥局
7	5912010313	金松大藥局
8	5912010322	南門藥局
9	5912010331	大家藥局
10	5912010484	聖心藥局
11	5912010724	倫洋藥局
12	5912010779	誠真藥局

新竹市
新竹市東區民族路 4 0 之 2 號
新竹市北區國華街六十九號一樓
新竹市香山區牛埔里育德街 1 8 8 號 2 樓
新竹市東區南大里西大路一〇二巷三號一樓
新竹市東區親仁里東大路 1 段 1 7 9 號
新竹市東區新莊里長春街 3 7 號
新竹市東區立功里光復路二段 3 1 0 號
新竹市東區關帝里南門街 8 2 號
新竹市東區南門里武昌街 6 號 1 樓
新竹市東區南大里食品路 2 6 9 號
新竹市東區建中路 3 8 號 1 樓
新竹市東區牛埔里牛埔路五段 2 9 號

field4	fi
過濾	過濾
03)5236158	0
03)5353969	600
03)5388109	600
03)5283828	597
03)5321797	0
03)5779177	600
03)5719019	0
03)5222665	0
035)234035	600
03)5610697	597
03)5749482	0
03)5278233	4



ID	Name	(X, Y)	Tel
2312010014	新竹市東區衛生所	(122.1, 34.123)	
2312041028	新竹市北區衛生所	(122.1, 35.12131)	

Spatial
Query in
DB will
rescue
you !



[MySQL 5.7 Reference Manual](#) / [Functions and Operators](#) / [Spatial Analysis Functions](#)

12.16 Spatial Analysis Functions

- [12.16.1 Spatial Function Reference](#)
- [12.16.2 Argument Handling by Spatial Functions](#)
- [12.16.3 Functions That Create Geometry Values from WKT Values](#)
- [12.16.4 Functions That Create Geometry Values from WKB Values](#)
- [12.16.5 MySQL-Specific Functions That Create Geometry Values](#)
- [12.16.6 Geometry Format Conversion Functions](#)
- [12.16.7 Geometry Property Functions](#)
- [12.16.8 Spatial Operator Functions](#)
- [12.16.9 Functions That Test Spatial Relations Between Geometry Objects](#)
- [12.16.10 Spatial Geohash Functions](#)
- [12.16.11 Spatial GeoJSON Functions](#)
- [12.16.12 Spatial Convenience Functions](#)

Summary

- The relational model has rigorously defined query languages that are simple and powerful.
- **Relational algebra** is operational; useful as an internal representation for query evaluation steps.
- Typically there are several ways of expressing a given query; a query optimizer should choose an efficient version.
- Reading assignment: Chapter X
-