


# Indexing Structures

## B+-tree, Hashing

# Range Searches


- *“Find all students with  $\text{gpa} > 3.0$ ”*
  - If records are sorted on gpa, do binary search to find first such student, then scan to find others.
  - Cost of binary search can be quite high.
- Simple idea: Create an ‘index’ file.



☞ *Can do binary search on (smaller) index file!*

# B+ Tree: Most Widely Used Index


- General concept of a tree:



- I. Height of a node: its distance to the root
- II. If a higher level node is connected to a lower level node, then the higher level node is called a **parent** (grandparent, ancestor, etc.) of the lower level node

## B+ Tree (cont.)


- Balanced tree: all leaves at the same level
  - Example:



- Structure of a B+ tree:
  - It is **balanced**: all leaf nodes are at the same level
  - Each node has a special structure


# Example: A B+ tree with order of 1

- Each node must hold at least 1 entry, and at most 2 entries

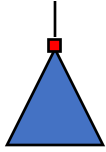


- Given search key values 27 how to find the rids?
  - Search begins at the root, and key comparisons direct it to a leaf
  - Note that key values are sorted at each level


# B+ tree: Internal node structure



Each  $P_i$  is a pointer to a child node, each  $K_i$  is a search key value  
 Pointers outnumber search key values by *exactly one*.

- $K_1 < K_2 < \dots < K_n$
- If the node is not the root, we require  $d \leq n \leq 2d$  where  $d$  is a pre-determined value for this B+ tree, called its *order*
- If the node is the root,  $1 \leq n \leq 2d$
- For any search key value  $K$  in the subtree  pointed by  $P_i$ ,
  - If  $P_i = P_0$ , we require that  $K < K_1$
  - If  $P_i = P_1, \dots, P_{n-1}$ , we require that  $K_i \leq K < K_{i+1}$
  - If  $P_i = P_n$ , we require  $K_n \leq K$


## B+ tree: leaf node structure



- Each  $r_i$  is a pointer to a record that contains search key value  $K_i$
- $L$  points to the left neighbor, and  $R$  points to the right neighbor
- $K_1 < K_2 < \dots < K_n$
- $d \leq n \leq 2d$  where  $d$  is the order of this B+ tree
- We will use  $K_i^*$  for the pair  $(K_i, r_i)$  and omit  $L$  and  $R$  for simplicity in the remaining slides

## Example: a B+ tree with order 2

- Search for  $5^*$ ,  $15^*$ , all data entries  $\geq 24^*$  ...
- The last one is a range search, we need to do the sequential scan, starting from the first leaf containing a value  $\geq 24$ .







# Inserting a Data Entry into a B+ Tree

- Find correct leaf  $L$ .
- Put data entry onto  $L$ .
  - If  $L$  has enough space, done!
  - Else, must split  $L$  (into  $L$  and a new node  $L2$ )
    - Redistribute entries evenly, put **middle key** in  $L2$
    - **copy up** middle key.
    - Insert index entry pointing to  $L2$  into parent of  $L$ .
- This can happen recursively
  - To split index node, redistribute entries evenly,
  - but **push up** middle key. (Contrast with leaf splits.)
- Splits “grow” tree; root split increases height.
  - Tree growth: gets wider or one level taller at top.

# Inserting 16\*, 8\* into Example B+ tree



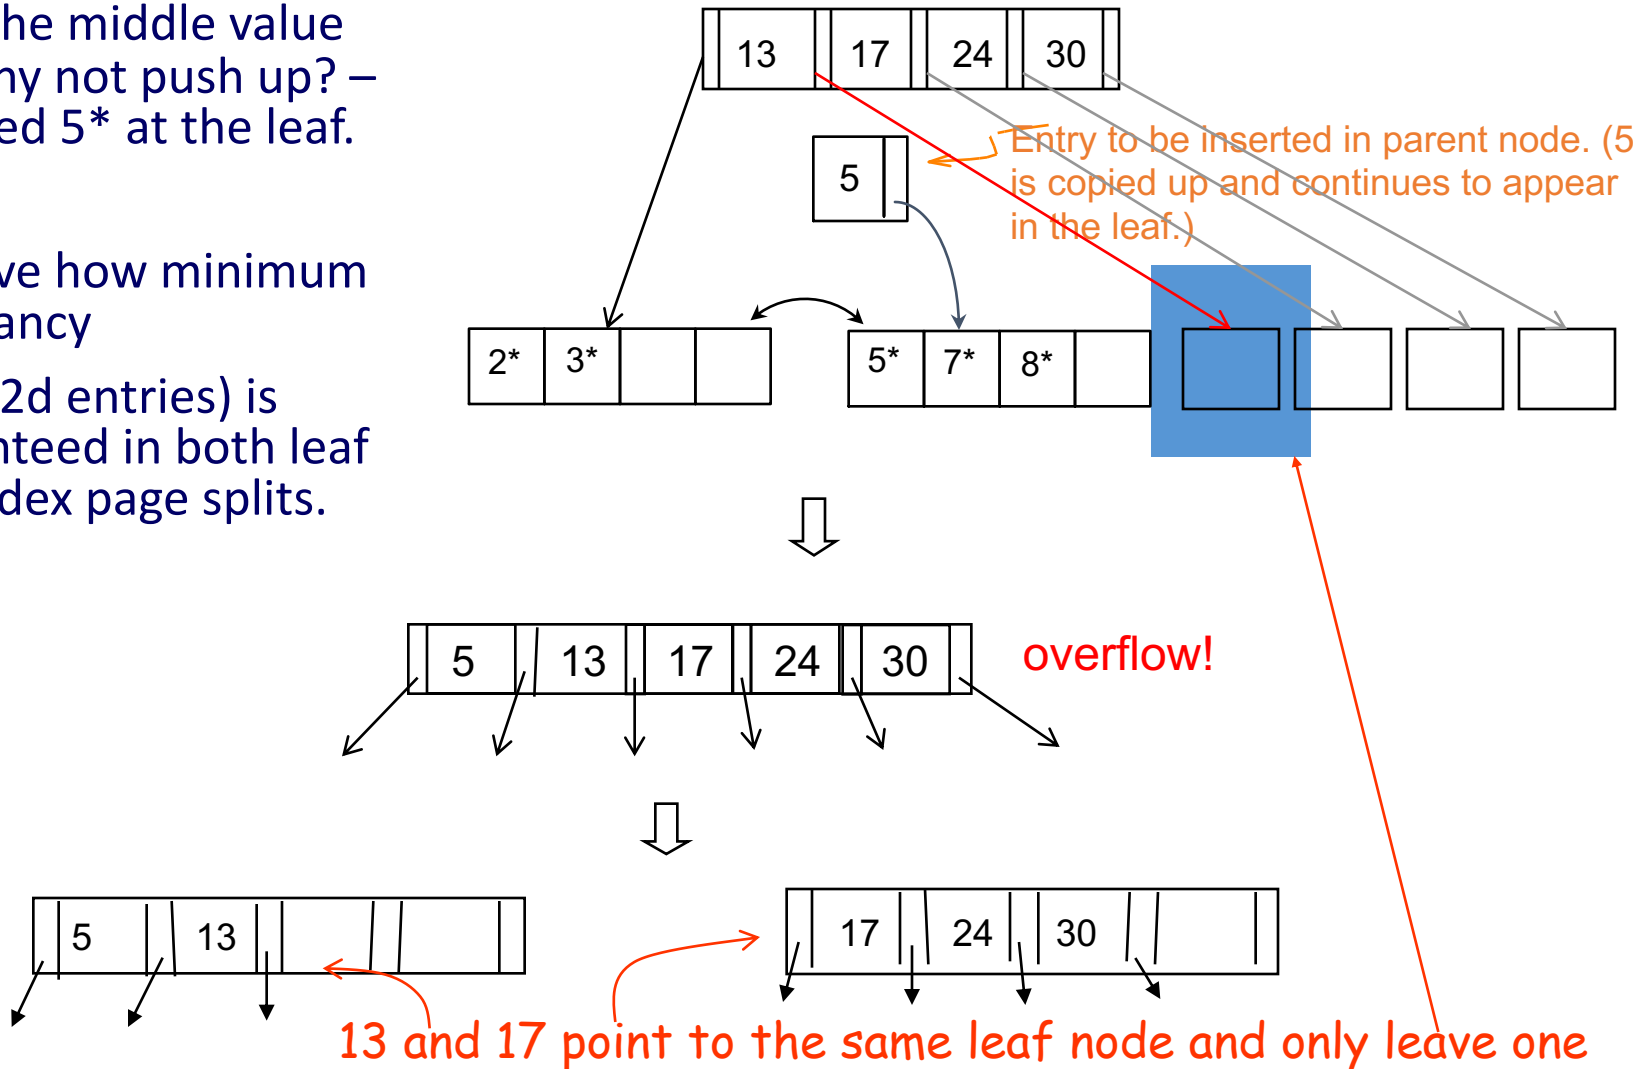
Overflow




One more child generated, must add one more pointer to its parent, thus one more key value as well.

## Inserting 8\* into Example B+ Tree (order 2)


- **Copy** the middle value up. Why not push up? – we need 5\* at the leaf.
- Observe how minimum occupancy  
(2 to 2d entries) is guaranteed in both leaf and index page splits.




## Insertion into B+ tree (cont.)



- We delete this pointer!
  - For internal nodes, we only have one copy of key values. Thus, 17 is deleted and push up to the parent node.
- $17 \leq k$



## Example B+ Tree After Inserting 8\*

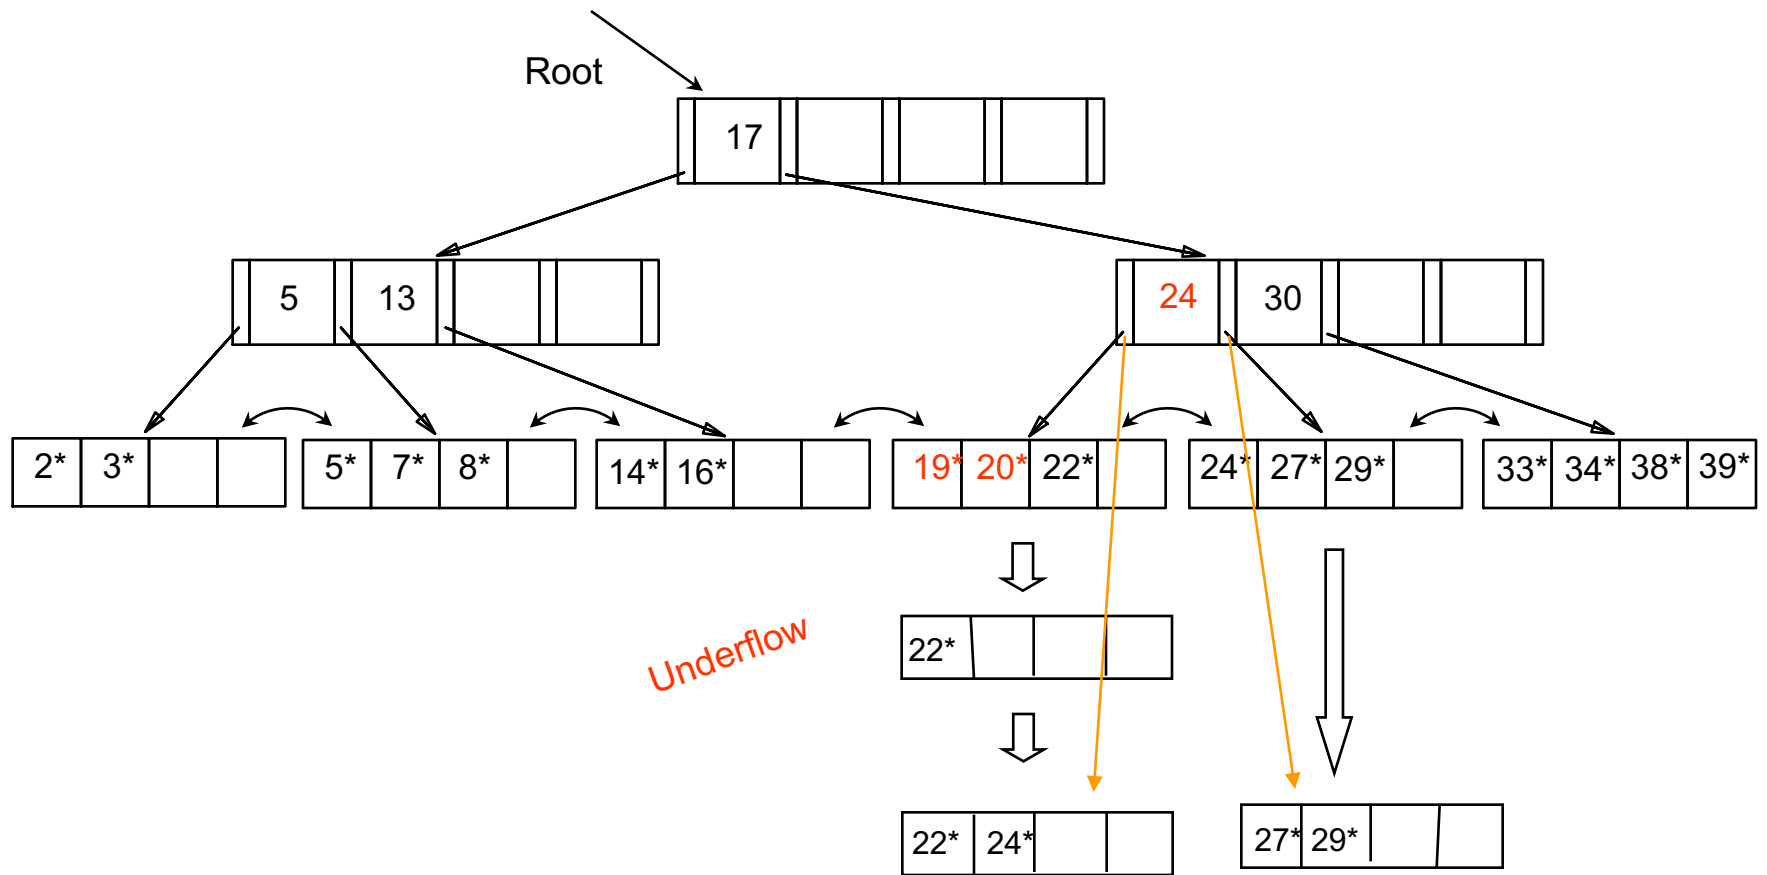


- the root was split, leading to increase in height.
- In this example, we can avoid splitting by **re-distributing** entries.

# Deleting a Data Entry from a B+ Tree

- Start at root, find leaf L where entry belongs.
- Remove the entry.
  - If L is at least half-full, done!
  - If L has only  $d-1$  entries,
    - Try to **re-distribute**, borrowing from sibling (adjacent node **with the same parent as L**).
    - If re-distribution fails, merge L and sibling.
- If merge occurred, must delete entry (pointing to L or sibling) from parent of L.
- Merge could propagate to root, decreasing height.


# Delete 19\* and 20\*



Have we still forgotten something?

The left child key values should be smaller than 24.


## Deleting 19\* and 20\* (cont.)




- Notice how 27 is **copied up**.
- Now we want to delete 24
- Underflow again! But can we redistribute this time? (**Leave it to you**)




# Deleting 24\*



Observe the two  
leaf nodes are  
merged, and 27 is  
discarded from  
their parent, but...




## Deleting 24\* (Cont.)

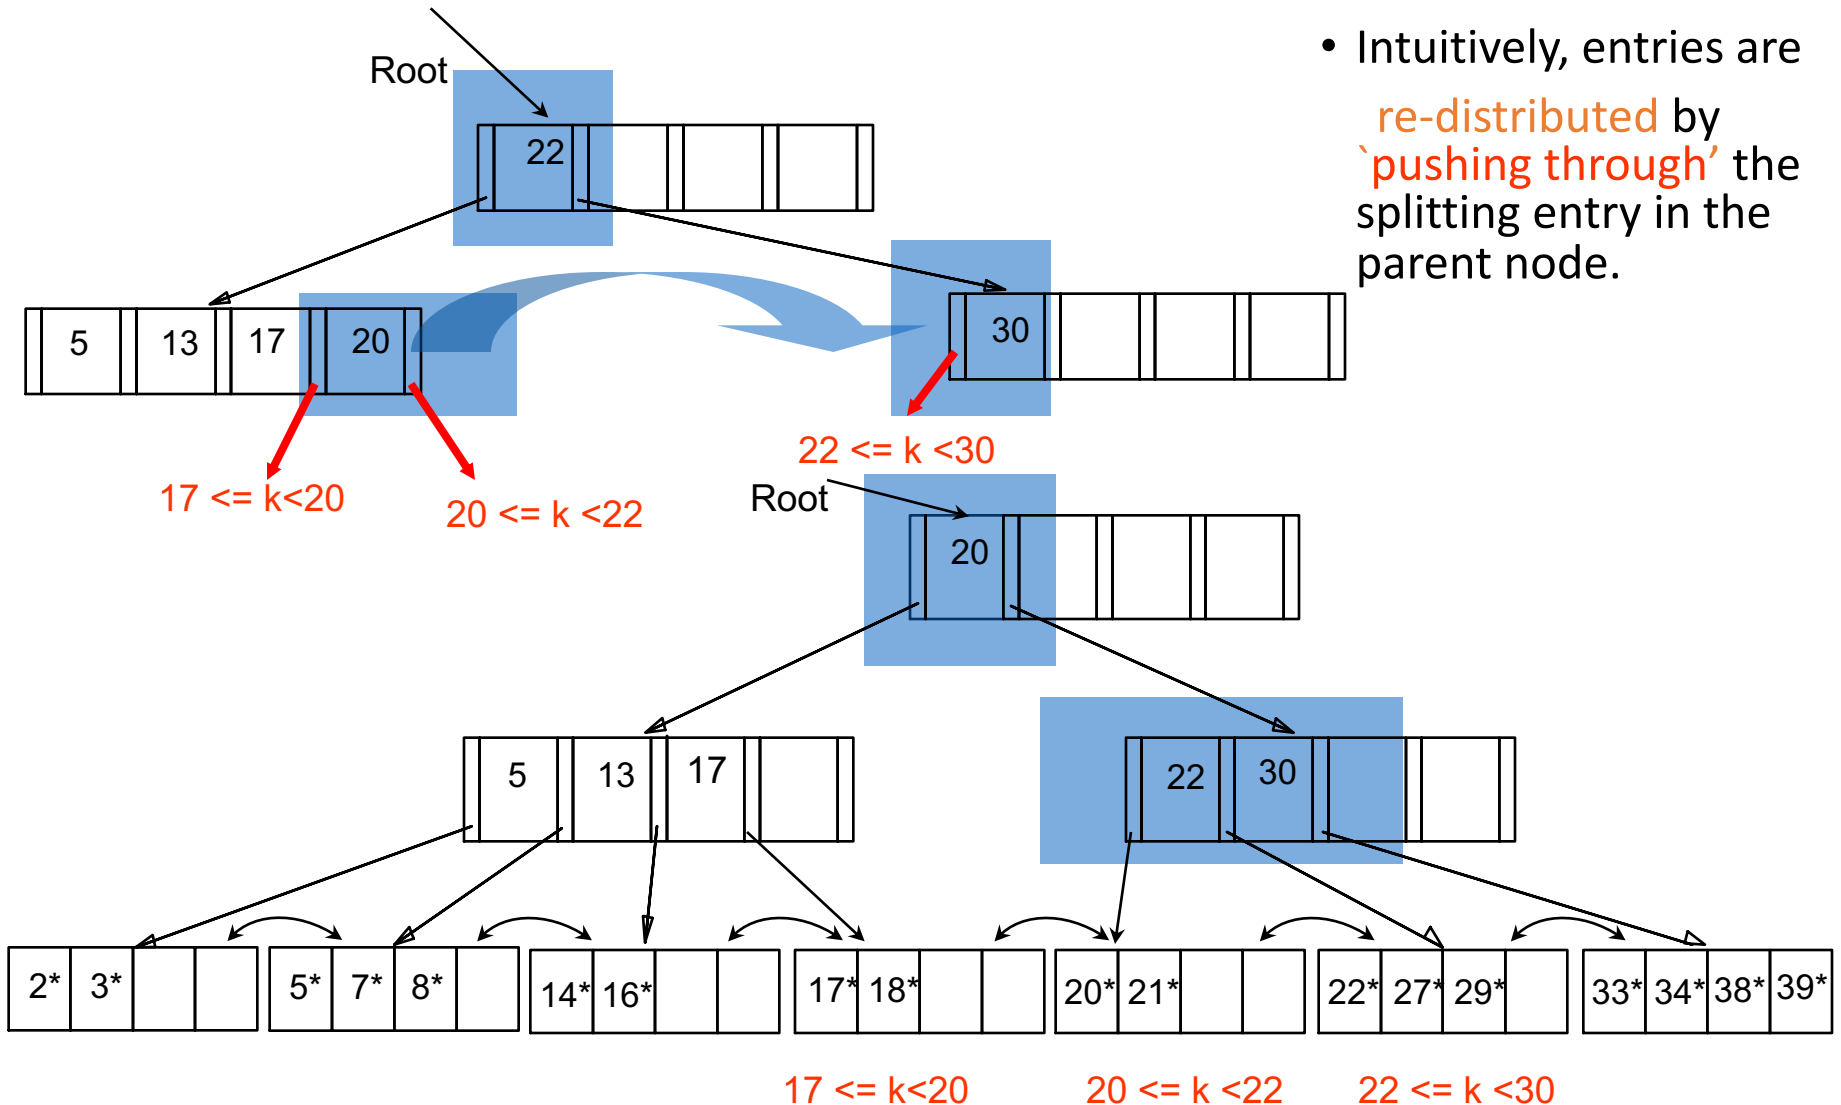


# Example of Non-leaf Re-distribution

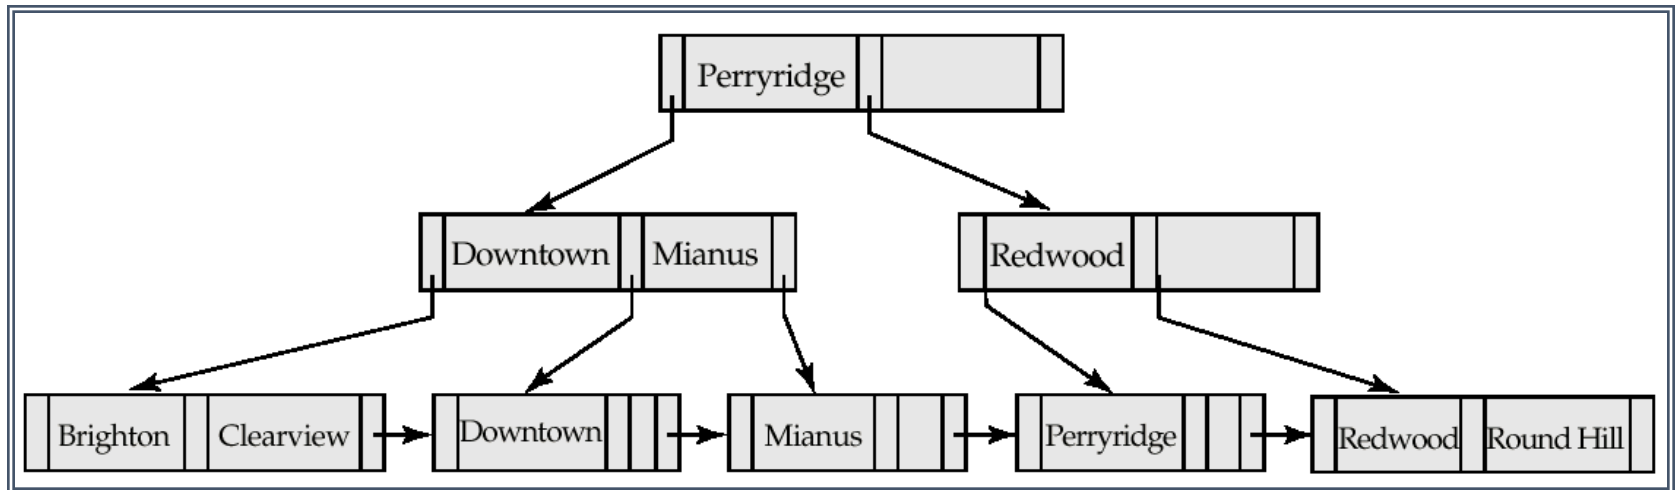
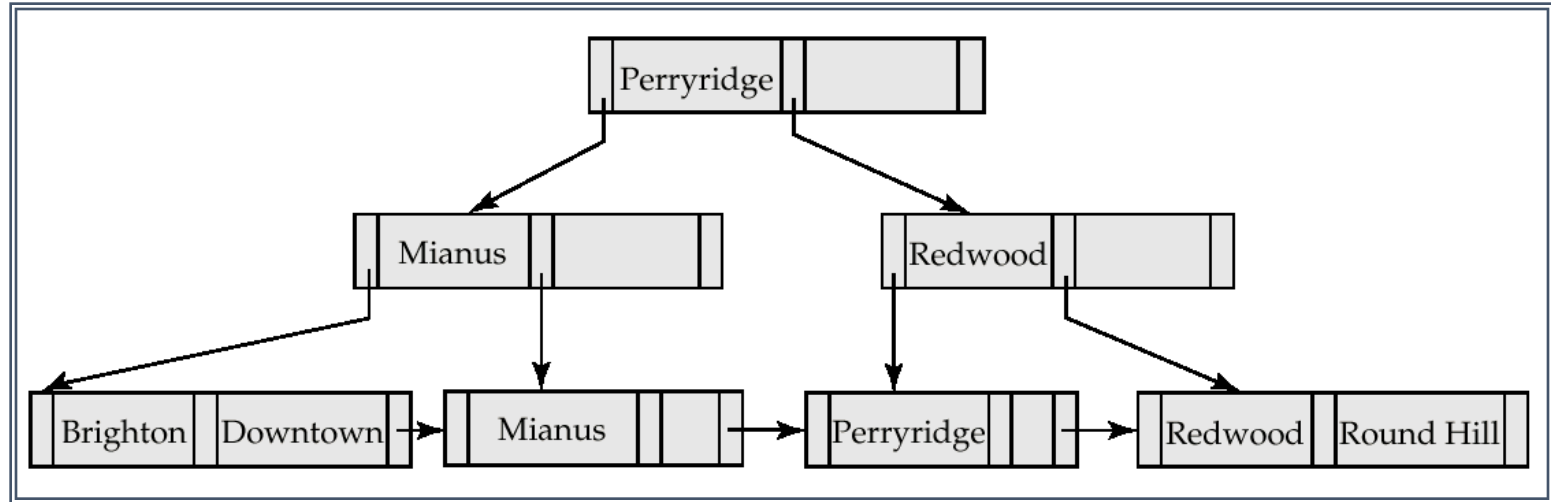
- Tree is shown below *during deletion* of 24\*.
- In contrast to previous example, re-distribute entry from left child of root to right child.



# Non-leaf Re-distribution (Cont.)



# B<sup>+</sup>-Trees: Insertion

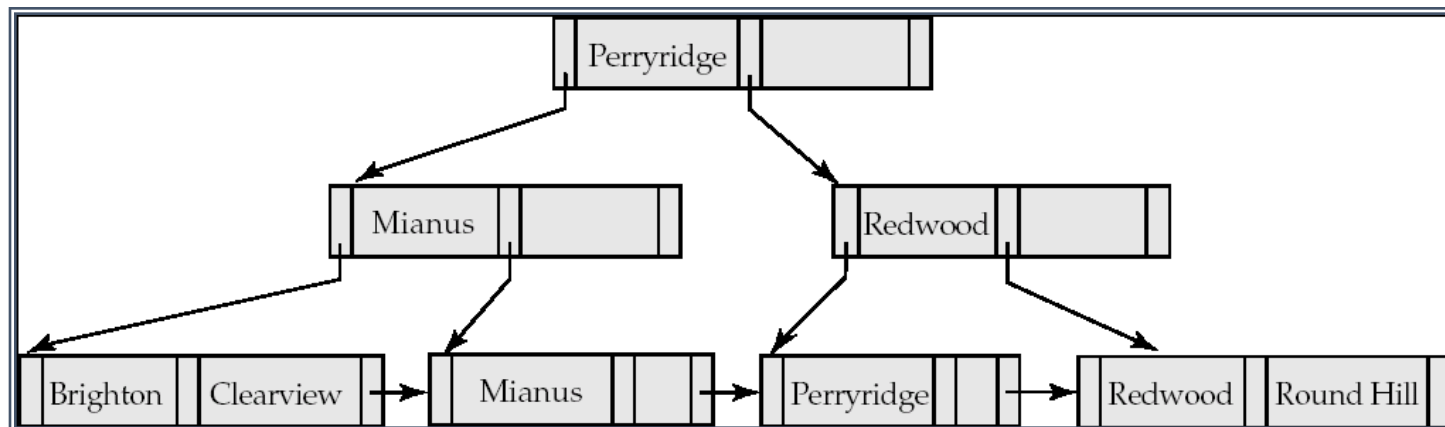
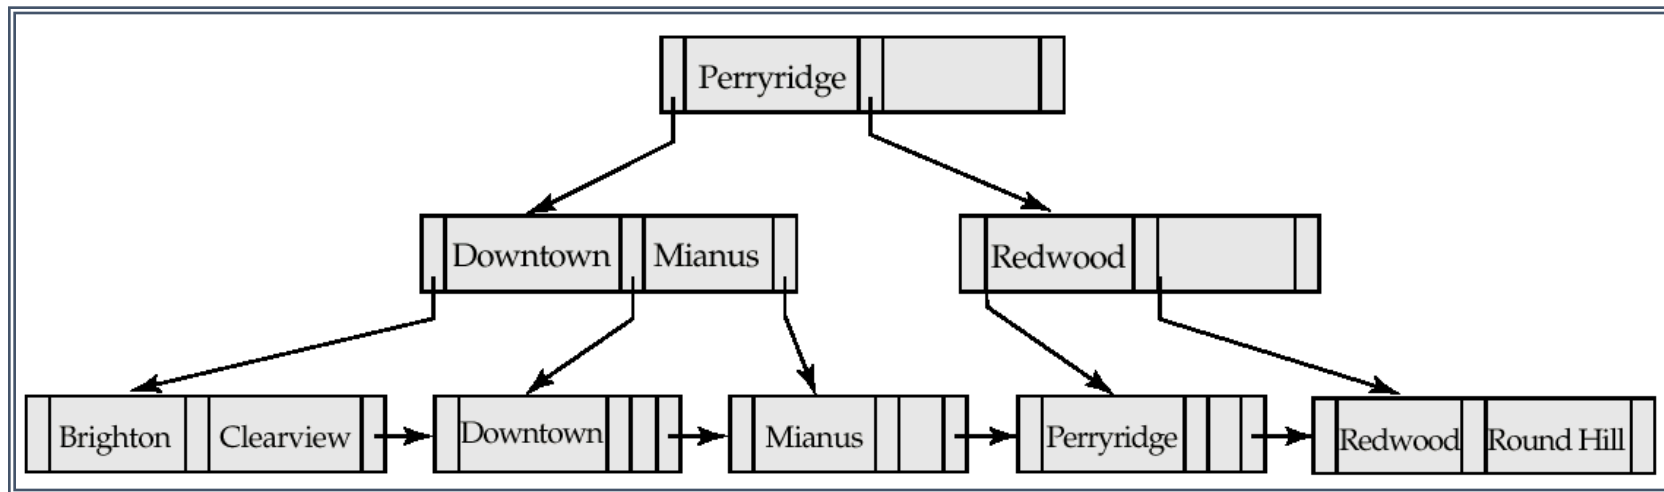


B<sup>+</sup>-Tree before and after insertion of “Clearview”

# Updates on B<sup>+</sup>-Trees: Deletion

- Find the record, delete it.
- Remove the corresponding (search-key, pointer) pair from a leaf node
  - Note that there might be another tuple with the same search-key
  - In that case, this is not needed
- Issue:
  - The leaf node now may contain too few entries
    - Why do we care ?
  - Solution:
    1. See if you can borrow some entries from a sibling
    2. If all the siblings are also just barely full, then *merge (opposite of split)*
  - May end up merging all the way to the root
  - In fact, may reduce the height of the tree by one

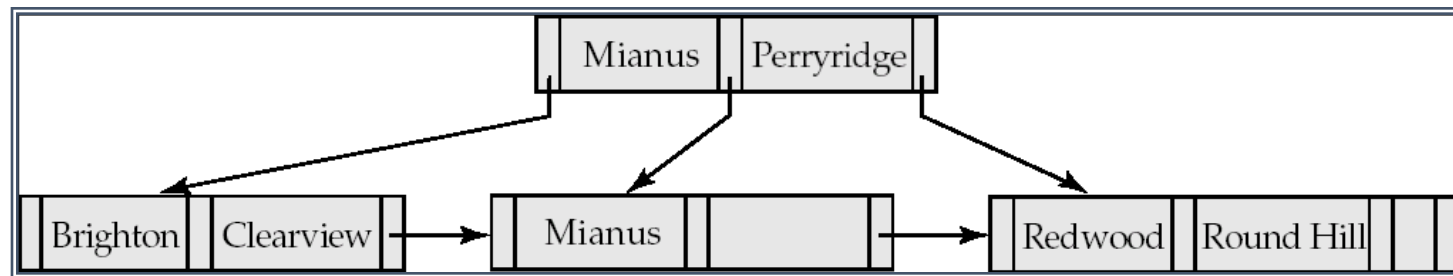
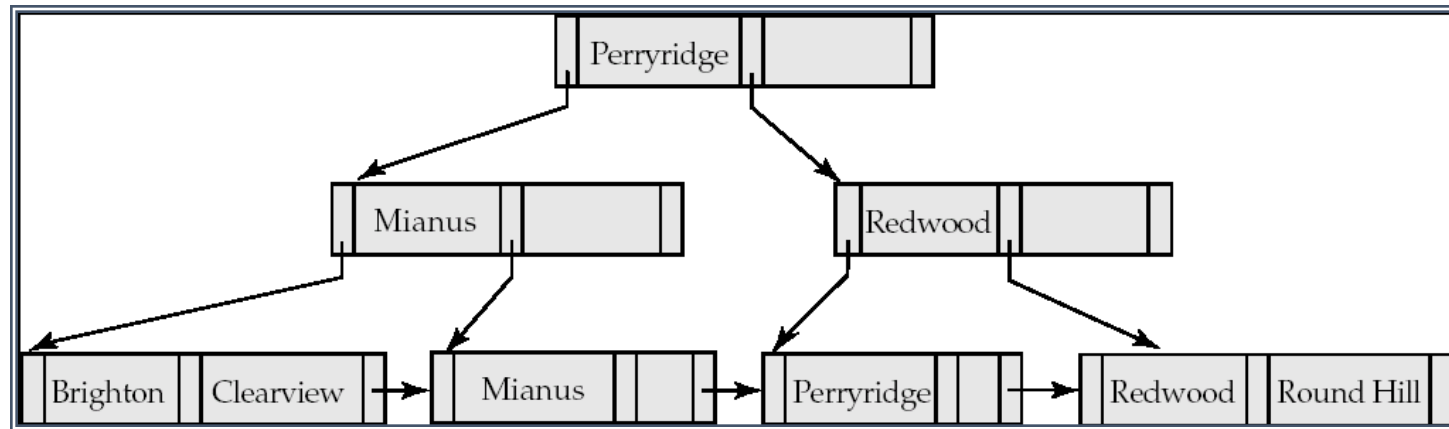
# Examples of B<sup>+</sup>-Tree Deletion



Before and after deleting "Downtown"

- Deleting "Downtown" causes merging of under-full leaves
  - leaf node can become empty only for  $n=3$ !

# Examples of B<sup>+</sup>-Tree Deletion



Deletion of "Perryridge" from result of previous example



# B+ Trees in Practice

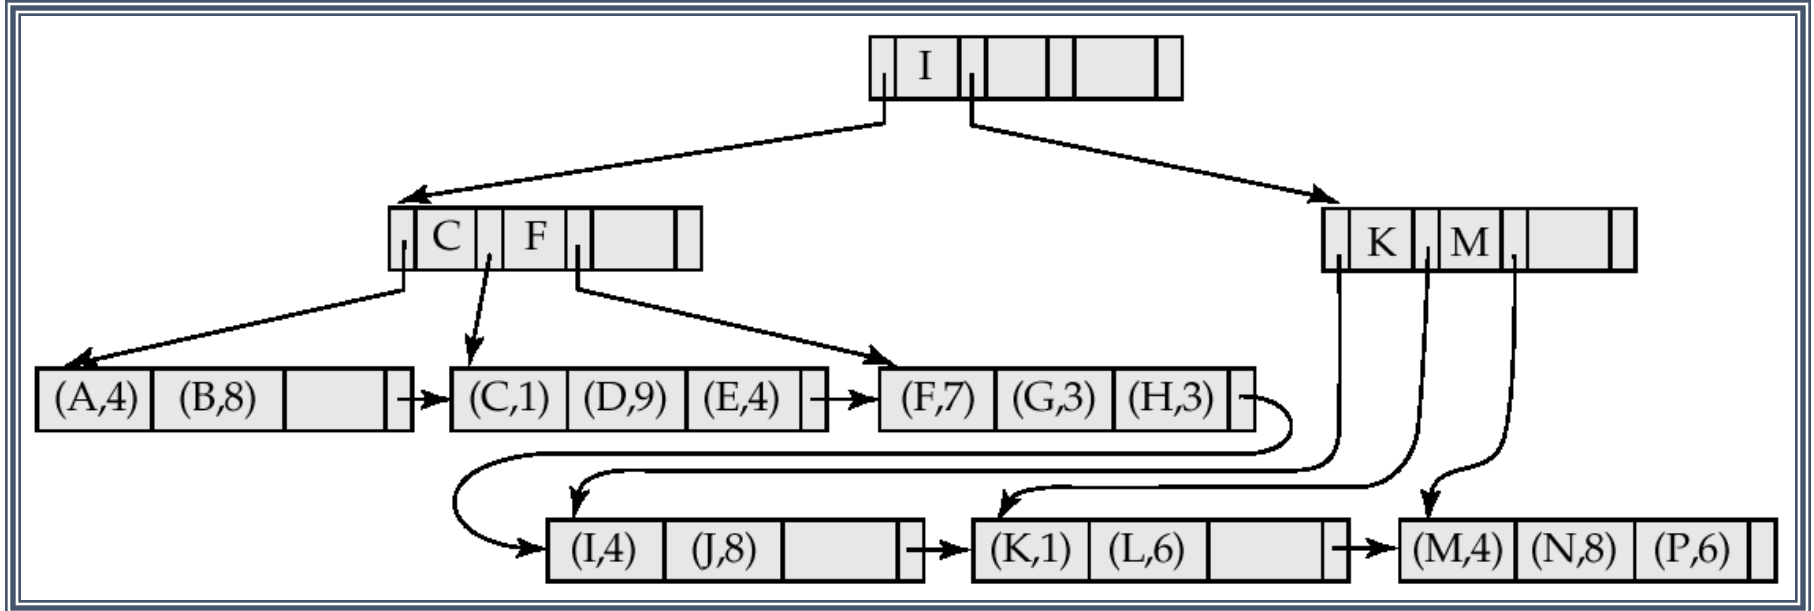
- Typical order: 100. Typical fill-factor: 67%.
  - average fanout = 133
- Typical capacities:
  - Height 3:  $133^3 = 2,352,637$  entries
  - Height 4:  $133^4 = 312,900,700$  entries
- Can often hold top levels in buffer pool:
  - Level 1 = 1 page = 8 Kbytes
  - Level 2 = 133 pages = 1 Mbyte
  - Level 3 = 17,689 pages = 133 MBytes

# B+ Trees: Summary

- Searching:
  - $\log_d(n)$  – Where  $d$  is the order, and  $n$  is the number of entries
- Insertion:
  - Find the leaf to insert into
  - If full, split the node, and adjust index accordingly
  - Similar cost as searching
- Deletion
  - Find the leaf node
  - Delete
  - May not remain half-full; must adjust the index accordingly

# B+-Tree File Organization

- Store the records at the leaves
- Sorted order etc..



# Hash-based File Organization

Store record with search key  $k$   
in block number  $h(k)$

e.g. for a person file,  
 $h(SSN) = SSN \% 4$


*Blocks called “buckets”*

What if the block becomes full ?  
Overflow pages

Uniformity property:

Don't want all tuples to map to  
the same bucket

$h(SSN) = SSN \% 2$  would be bad



# Hash-based File Organization

Hashed on “branch-name”

Hash function:

$$a = 1, b = 2, \dots, z = 26$$

$$h(abz)$$

$$= (1 + 2 + 26) \% 10$$

$$= 9$$

bucket 0			
bucket 1			
bucket 2			
bucket 3	A-217	Brighton	750
	A-305	Round Hill	350
bucket 4	A-222	Redwood	700
bucket 5	A-102	Perryridge	400
	A-201	Perryridge	900
	A-218	Perryridge	700
bucket 6			
bucket 7	A-215	Mianus	700
bucket 8	A-101	Downtown	500
	A-110	Downtown	600
bucket 9			

# Hash Indexes

Extends the basic idea

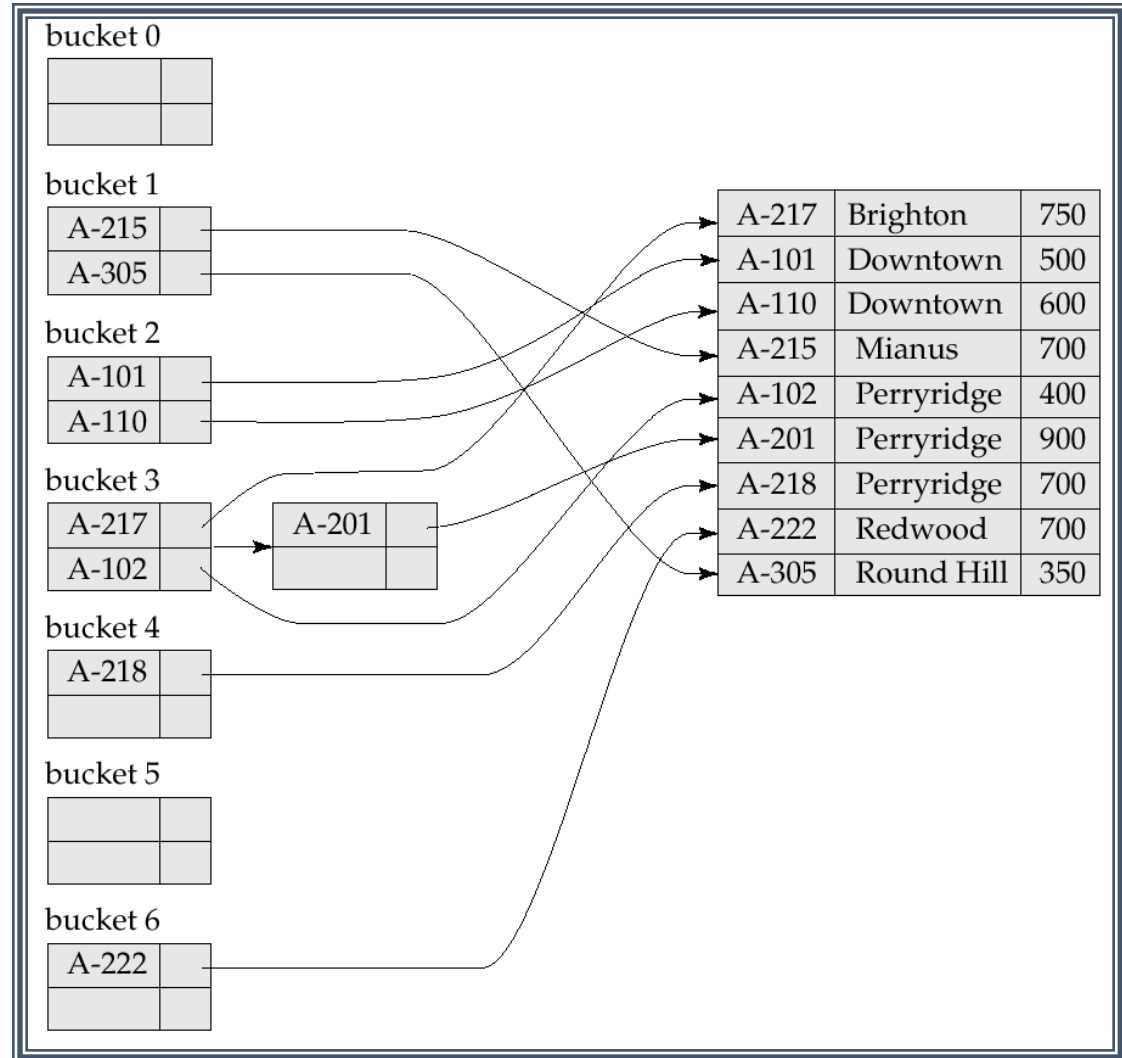
Search:

Find the block with  
search key

Follow the pointer

Range search ?

$a < X < b$  ?

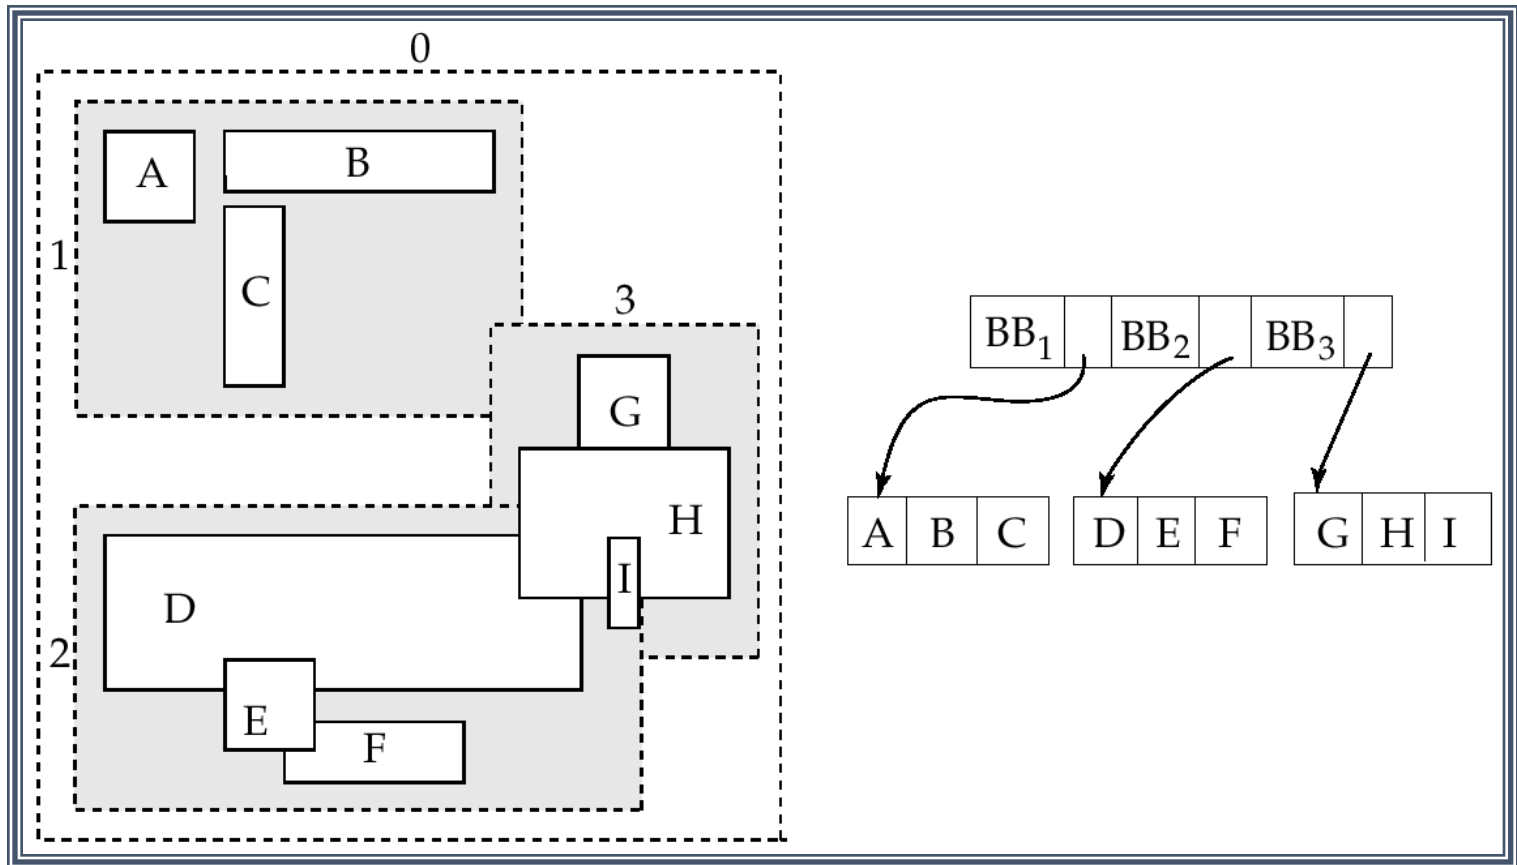


# Hash Indexes

- Very fast search on equality
- Can't search for “ranges” at all
  - Must scan the file
- Inserts/Deletes
  - Overflow pages can degrade the performance
- Two approaches
  - Dynamic hashing
  - Extendible hashing

# R-Trees

For spatial data (e.g. maps, rectangles, GPS data etc)





# Conclusions

- Indexing Goal: "Quickly find the tuples that match certain conditions"
- Equality and range queries most common
  - Hence B+-Trees the predominant structure for on-disk representation
  - Hashing is used more commonly for in-memory operations
- Many many more types of indexing structures exists
  - For different types of data
  - For different types of queries
    - E.g. "nearest-neighbor" queries