
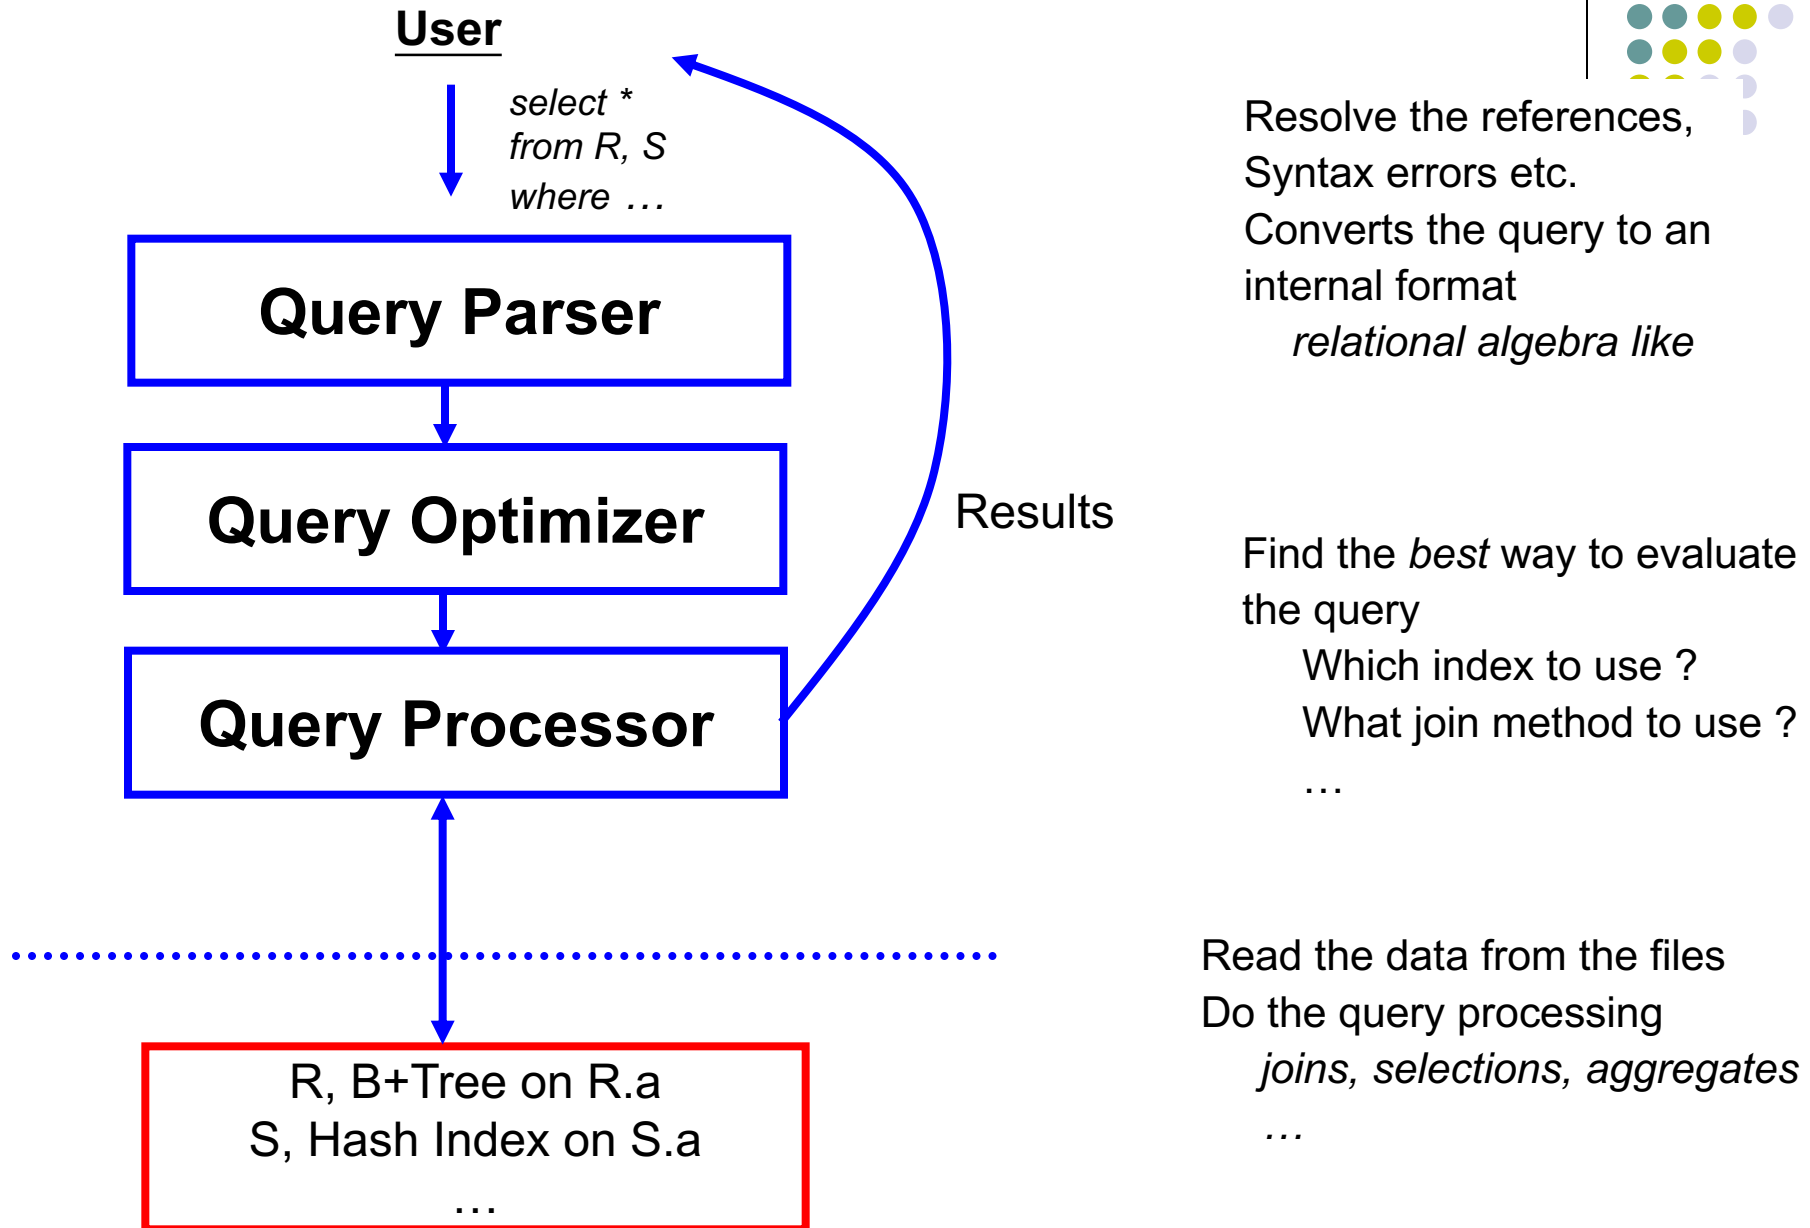



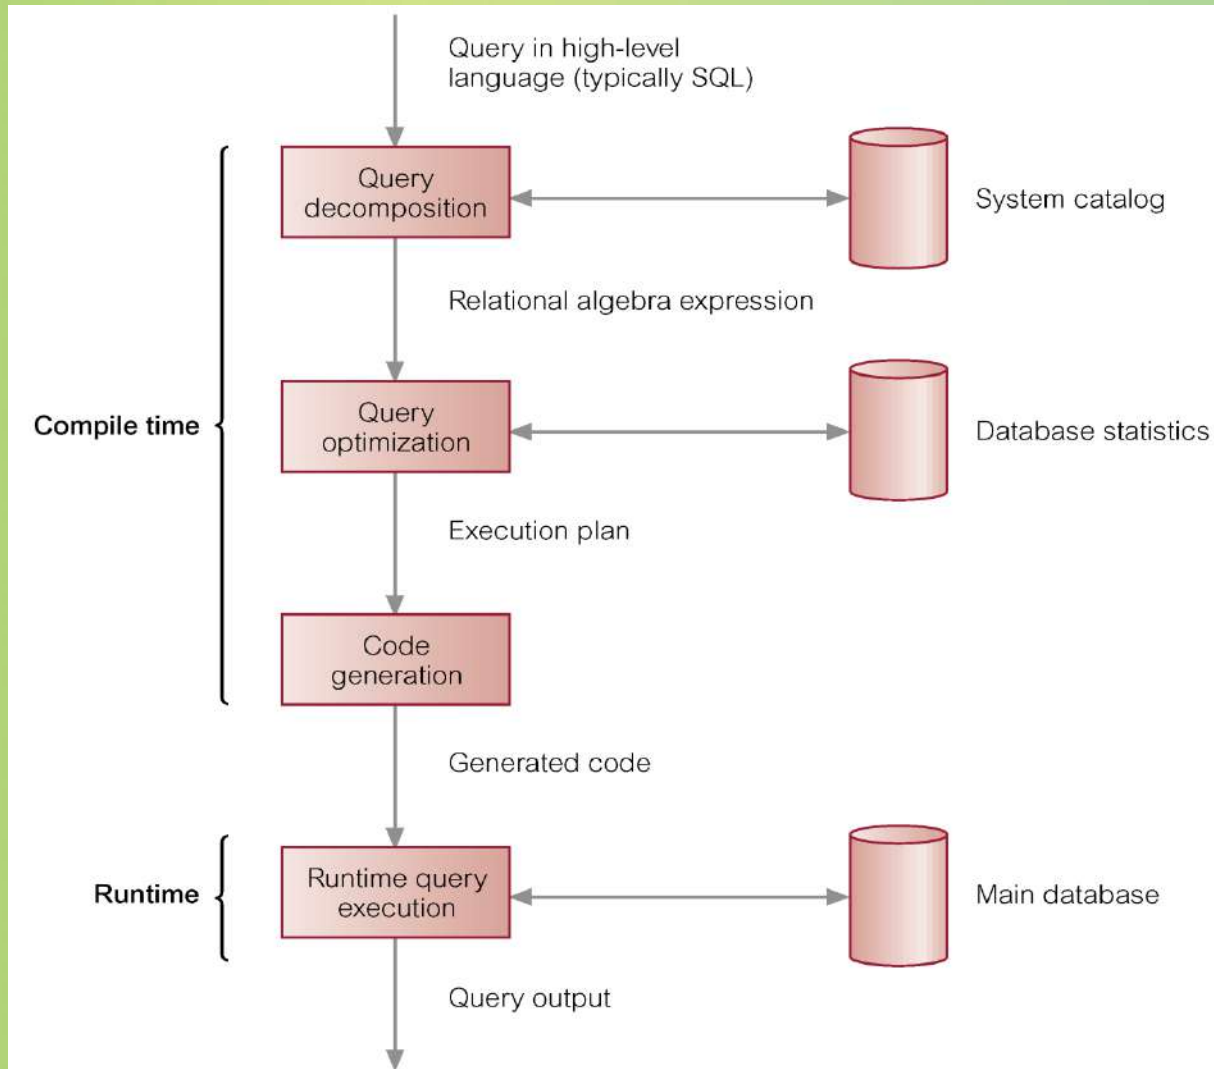
Query Processing



Overview



Phases of Query Processing




Dynamic versus Static Optimization

- **Two times when first three phases of QP can be carried out:**
 - dynamically every time query is run;
 - statically when query is first submitted.
- **Advantages of dynamic QO arise from fact that information is up to date.**
- **Disadvantages are that performance of query is affected, time may limit finding optimum strategy.**


Dynamic versus Static Optimization

- Advantages of static QO are removal of runtime overhead, and more time to find optimum strategy.
- Disadvantages arise from fact that chosen execution strategy may no longer be optimal when query is run.
- Could use a hybrid approach to overcome this.




“Cost”

- Complicated to compute
- We will focus on disk:
 - Number of I/Os ?
 - Not sufficient
 - Number of seeks matters a lot... why ?
 - t_T – time to transfer one block
 - t_S – time for one seek
 - Cost for b block transfers plus S seeks
$$b * t_T + S * t_S$$
 - Measured in *seconds*




Selection Operation

- select * from person where SSN = “123”
- Option 1: Sequential Scan
 - Read the relation start to end and look for “123”
 - Can always be used (not true for the other options)
 - Cost
 - Let b_r = Number of relation blocks
 - Then:
 - 1 seek and b_r block transfers
 - So:
 - $t_s + b_r * t_T$ sec
 - Improvements:
 - If SSN is a key, then can stop when found
 - So on average, $b_r/2$ blocks accessed




Selection Operation

- select * from person where SSN = “123”
- Option 2 : Binary Search:
 - Pre-condition:
 - *The relation is sorted on SSN*
 - *Selection condition is an equality*
 - E.g. can't apply to “Name like ‘%424%’”
 - Do binary search
 - Cost of finding the *first* tuple that matches
 - $\lceil \log_2(b_r) \rceil * (t_T + t_S)$
 - All I/Os are random, so need a seek for all




Selection Operation

- select * from person where SSN = “123”
- Option 3 : Use Index
 - Pre-condition:
 - *An appropriate index must exist*
 - Use the index
 - Find the first leaf page that contains the search key
 - Retrieve all the tuples that match by following the pointers
 - If primary index, the relation is sorted by the search key
 - Go to the relation and read blocks sequentially
 - If secondary index, must follow all points using the index




Selection Operation

- Selections involving ranges
 - *select * from accounts where balance > 100000*
 - *select * from matches where matchdate between '10/20/06' and '10/30/06'*
 - Option 1: Sequential scan
 - Option 2: Using an appropriate index
 - Can't use hash indexes for this purpose
 - Cost formulas:
 - Range queries == “equality” on “non-key” attributes



Selection Operation

- Complex selections
 - Conjunctive: *select * from accounts where balance > 100000 and SSN = "123"*
 - Disjunctive: *select * from accounts where balance > 100000 or SSN = "123"*
 - Option 1: **Sequential scan**
 - *(Conjunctive only)* Option 2: **Using an appropriate index on one of the conditions**
 - E.g. Use SSN index to evaluate SSN = "123". Apply the second condition to the tuples that match
 - Or do the other way around (if index on balance exists)
 - Which is better ?
 - *(Conjunctive only)* Option 3: **Choose a multi-key index**
 - Not commonly available



Selection Operation

- Complex selections
 - Conjunctive: *select * from accounts where balance > 100000 and SSN = "123"*
 - Disjunctive: *select * from accounts where balance > 100000 or SSN = "123"*
- **Option 4: Conjunction or disjunction of *record identifiers***
 - Use indexes to find all RIDs that match each of the conditions
 - Do an intersection (for conjunction) or a union (for disjunction)
 - Sort the records and fetch them in one shot
 - Called "Index-ANDing" or "Index-ORing"
- Heavily used in commercial systems

Example 23.1 - Different Strategies

Find all Managers who work at a London branch.

```
SELECT *  
FROM Staff s, Branch b  
WHERE s.branchNo = b.branchNo AND  
(s.position = 'Manager' AND b.city = 'London');
```

Example 23.1 - Different Strategies

• Three equivalent RA queries are:

(1) $\sigma_{(\text{position}='Manager') \wedge (\text{city}='London') \wedge (\text{Staff.branchNo}=\text{Branch.branchNo})}$
(Staff X Branch)

(2) $\sigma_{(\text{position}='Manager') \wedge (\text{city}='London')}$ (
Staff \bowtie Staff.branchNo=Branch.branchNo Branch)

(3) ($\sigma_{\text{position}='Manager'}(\text{Staff})$) \bowtie Staff.branchNo=Branch.branchNo
($\sigma_{\text{city}='London'}(\text{Branch})$)

Example 23.1 - Different Strategies

• Assume:

- 1000 tuples in Staff; 50 tuples in Branch;
- 50 Managers; 5 London branches;
- no indexes or sort keys;
- results of any intermediate operations stored on disk;
- cost of the final write is ignored;
- tuples are accessed one at a time.

Example 23.1 - Cost Comparison

- Cost (in disk accesses) are:

(1) $(1000 + 50) + 2 * (1000 * 50) = 101\ 050$

(2) $2 * 1000 + (1000 + 50) = 3\ 050$

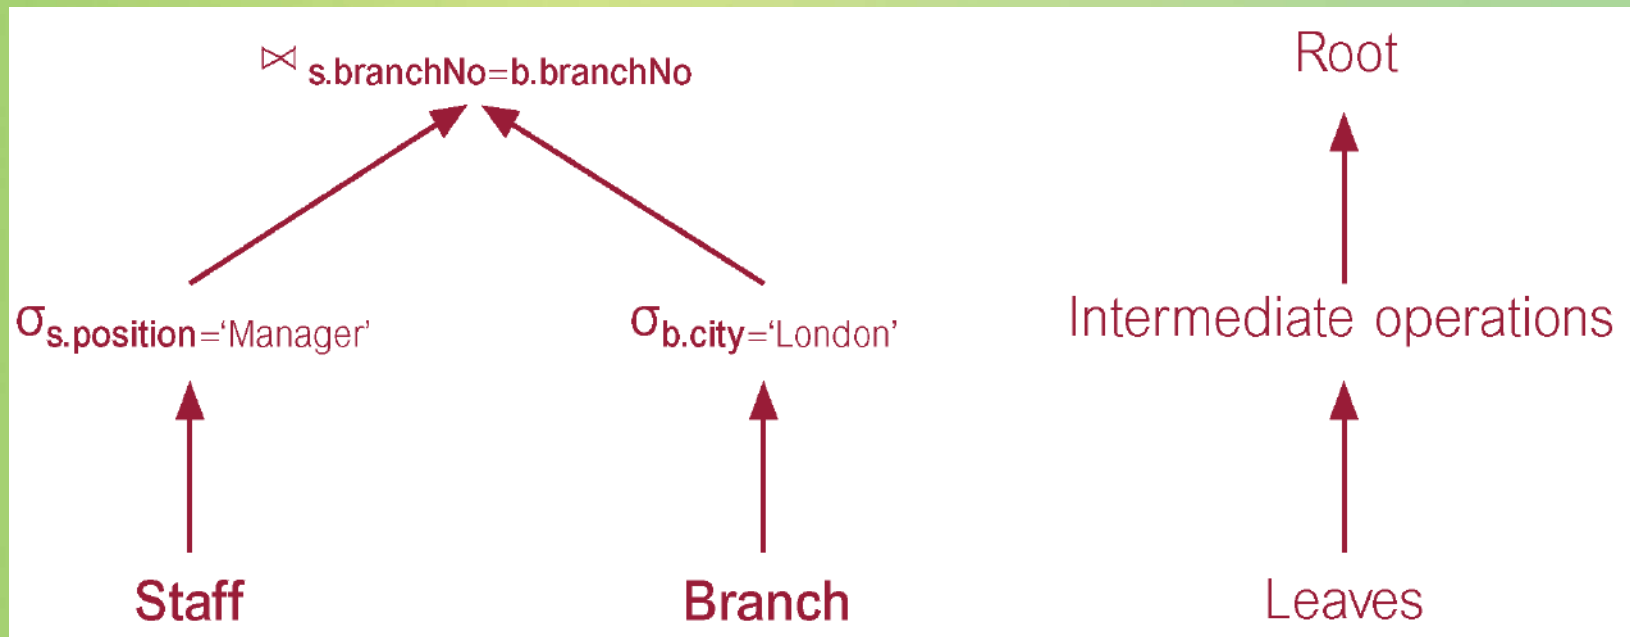
(3) $1000 + 2 * 50 + 5 + (50 + 5) = 1\ 160$


- Cartesian product and join operations much more expensive than selection, and third option significantly reduces size of relations being joined together.

Analysis

- Finally, query transformed into some internal representation more suitable for processing.
- Some kind of query tree is typically chosen, constructed as follows:
 - Leaf node created for each base relation.
 - Non-leaf node created for each intermediate relation produced by RA operation.
 - Root of tree represents query result.
 - Sequence is directed from leaves to root.


Example 23.1 - R.A.T.






Query Processing

- Overview
- Selection operation
- **Join operators**
- Sorting
- Other operators
- Putting it all together...



Join




- *select * from R, S where $R.a = S.a$*
 - R called *outer relation*
 - S called *inner relation*
 - Called an “*equi-join*”
- *select * from R, S where $|R.a - S.a| < 0.5$*
 - Not an “*equi-join*”
- Option 1: Nested-loops
 - *for each tuple r in R*
 - *for each tuple s in S*
 - *check if $r.a = s.a$ (or whether $|r.a - s.a| < 0.5$)*
- Can be used for any join condition




Nested-loops Join

- Cost ? Depends on the actual values of parameters, especially memory
- $b_r, b_s \rightarrow$ Number of blocks of R and S
- $n_r, n_s \rightarrow$ Number of tuples of R and S
- Case 1: Minimum memory required = 3 blocks
 - One to hold the current R block, one for current S block, one for the result being produced
 - Blocks transferred:
 - Must scan R tuples once: b_r
 - For each tuple in R , must scan S : $n_r * b_s$
 - Seeks
 - $n_r + b_r$


Example






Nested-loops Join

- Case 1: Minimum memory required = 3 blocks
 - Blocks transferred: $n_r * b_s + b_r$
 - Seeks: $n_r + b_r$ In this case, disk blocks in R is not ordered and data tuples in a disk block are not ordered.
- Example:
 - Number of records -- $R: n_r = 10,000, S: n_s = 5000$
 - Number of blocks -- $R: b_r = 400, S: b_s = 100$
- Then:
 - blocks transferred: $10000 * 100 + 400 = 1,000,400$
 - seeks: 10400
- What if we were to switch R and S ?
 - 2,000,100 block transfers, 5100 seeks



Nested-loops Join

- Case 2: S fits in memory
 - Blocks transferred: $b_s + b_r$
 - Seeks: 2 (Assume that R is also sequential read into a block)
- Example:
 - Number of records -- R: $n_r = 10,000$, S: $n_s = 5000$
 - Number of blocks -- R: $b_r = 400$, S: $b_s = 100$
- Then:
 - blocks transferred: $400 + 100 = 500$
 - seeks: 2
- This is orders of magnitude difference



Block Nested-loops Join

- Simple modification to “nested-loops join”

- Block at a time


for each block B_r in R

for each block B_s in S

for each tuple r in B_r


for each tuple s in B_s

check if $r.a = s.a$ (or whether $|r.a - s.a| < 0.5$)



Block Nested-loops Join


- Case 1: Minimum memory required = 3 blocks
 - Blocks transferred: $b_r * b_s + b_r$
 - Seeks: $b_r + b_r$
- Case 2: S fits in memory
 - Blocks transferred: $b_s + b_r$
 - Seeks: 2 (Assume that R is also sequential read into a block)
- What about in between ?
 - There are 50 blocks, but S is 100 blocks



Block Nested-loops Join


- Case 3: 50 blocks (S = 100 blocks) ?
 - for each block **in** R*
 - for each group of 48 blocks **in** S*
 - for each tuple **r** **in** one block*
 - for each tuple **s** **in** each group of 48 blocks in S*
 - check if** $r.a = s.a$ (or whether $|r.a - s.a| < 0.5$)*
- Why is this good ?
 - We only have to read S a total of $b_s/48$ times (instead of b_s times)
 - Blocks transferred: $b_r * (b_s / 48) + b_r$
 - Seeks: 2 (if R and S are ordered files)

Index Nested-loops Join




- *select * from R, S where R.a = S.a*
 - Called an “*equi-join*”
- Nested-loops
 - for each tuple r in R*
 - for each tuple s in S*
 - check if $r.a = s.a$ (or whether $|r.a - s.a| < 0.5$)*
- Suppose there is an index on S.a
- Why not use the index instead of the inner loop ?
 - for each tuple r in R*
 - use the index to find S tuples with $S.a = r.a$*

Index Nested-loops Join



- Cost of the join:
 - $b_r (t_T + t_S) + n_r * c$
 - $c == \text{the cost of index access}$
 - *Computed using the formulas discussed earlier*


Index Nested-loops Join



- Restricted applicability
 - An appropriate index must exist
 - What about $|R.a - S.a| < 5$?
- Great for queries with joins and selections


*select **
from accounts, customers
where accounts.customer-SSN = customers.customer-SSN and
accounts.acct-number = "A-101"
- Only need to access one SSN from the other relation

Notes




- Block Nested-loops join
 - Can always be applied to various join conditions
 - If the smaller relation fits in memory, then cost:
 - $b_r + b_s$
 - This is the best we can hope if we have to read the relations once each
 - CPU cost of the inner loop is high
 - Typically used when the smaller relation is really small (few tuples) and index nested-loops can't be used
- Index Nested-loops join
 - Only applies if an appropriate index exists
 - Very useful when we have selections that return small number of tuples
 - **select** balance **from** customer, accounts **where** customer.name = "j. s." and customer.SSN = accounts.SSN

Hash Join




- Case 1: Smaller relation (S) fits in memory
- Nested-loops join:
 - for each tuple r in R*
 - for each tuple s in S*
 - check if $r.a = s.a$*
- Cost: $b_r + b_s$ transfers, 2 seeks (R and S are sorted)
- The inner loop is not exactly cheap (high CPU cost)
- Hash join:
 - read S in memory and build a hash index on it*
 - for each tuple r in R*
 - use the hash index on S to find tuples such that $S.a = r.a$*

Hash Join




- Case 1: Smaller relation (S) fits in memory
- Hash join:
 - read S in memory and build a hash index on it*
 - for each tuple r in R*
 - use the hash index on S to find tuples such that $S.a = r.a$*
- Cost: $b_r + b_s$ transfers, 2 seeks (unchanged)
- Why good ?
 - CPU cost is much better (even though we don't care about it too much)
 - Performs much better than nested-loops join when S doesn't fit in memory

Hash Join




- Case 2: Smaller relation (S) doesn't fit in memory
- Two “phases”
- Phase 1:
 - Read the relation R block by block and partition it using a hash function, $h1(a)$
 - Create one partition for each possible value of $h1(a)$
 - Write the partitions to disk
 - R gets partitioned into $R1, R2, \dots, Rk$
 - Similarly, read and partition S , and write partitions $S1, S2, \dots, Sk$ to disk
 - Only requirement:
 - Each S partition fits in memory

Hash Join



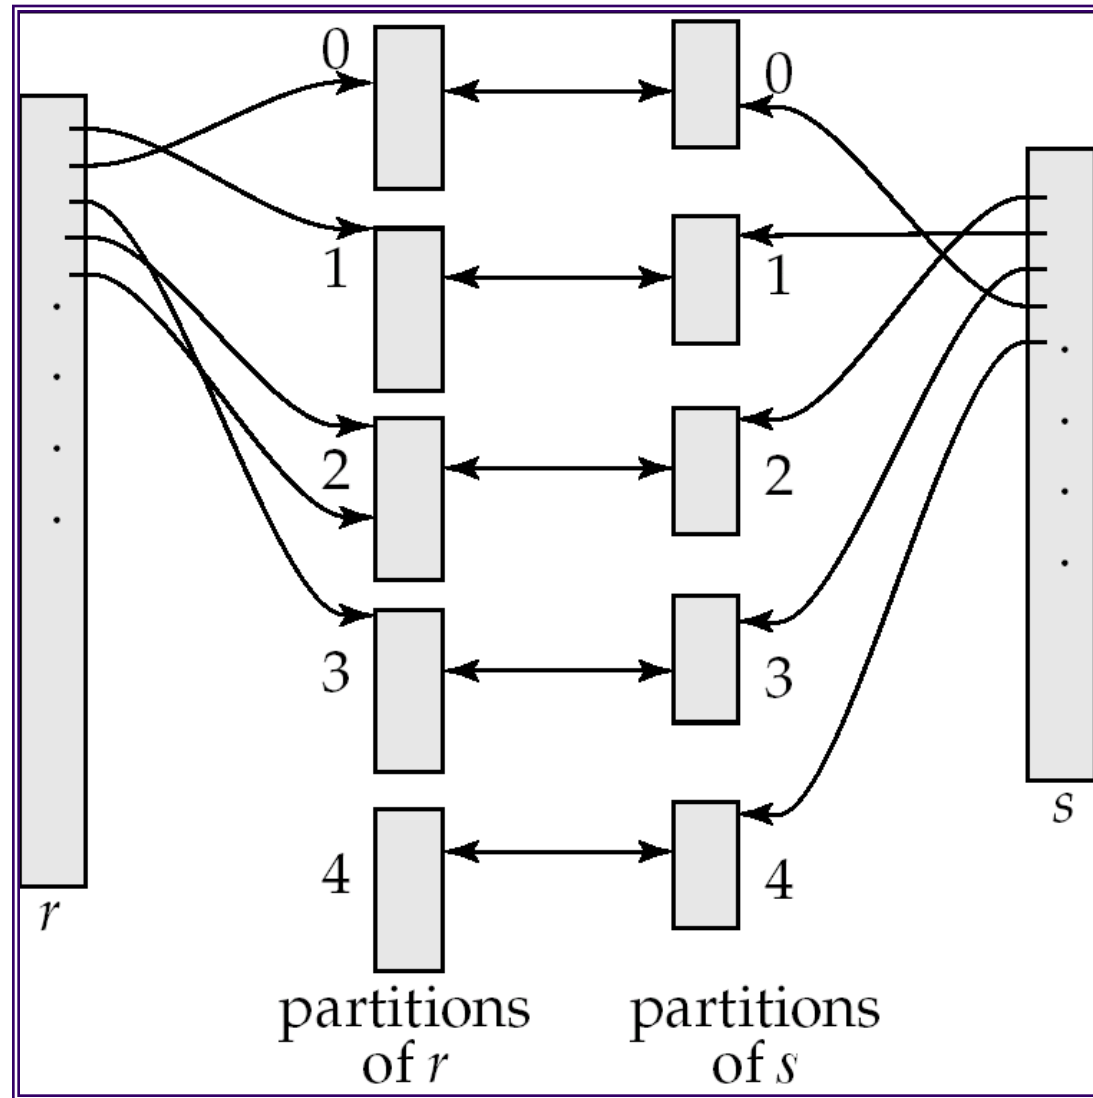

- Case 2: Smaller relation (S) doesn't fit in memory
- Two “phases”
- Phase 2:
 - Read S1 into memory, and build a hash index on it (S1 fits in memory)
 - Using a different hash function, $h_2(a)$
 - Read R1 block by block, and use the hash index to find matches.
 - Repeat for S2, R2, and so on.

Hash Join




- Case 2: Smaller relation (S) doesn't fit in memory
- Two “phases”:
- Phase 1:
 - Partition the relations using one hash function, $h_1(a)$
- Phase 2:
 - Read S_i into memory, and build a hash index on it (S_i fits in memory)
 - Read R_i , and use the hash index to find matches.
- Cost
 - $3(b_r + b_s) + 4 * n_h$ block transfers + $2(\lceil b_r / b_b \rceil + \lceil b_s / b_b \rceil)$ seeks (Read 13.5.5.4 in the reference text book)
 - Where b_b is the size of each output buffer
 - Much better than Nested-loops join under the same conditions

Hash Join




Hash Join: Issues




- How to guarantee that the partitions of S all fit in memory ?
 - $S = 10000$ blocks, Memory = $M = 100$ blocks
 - Use a hash function that hashes to 100 different values ?
 - Eg. $h1(a) = a \% 100$?
 - Problem: Impossible to guarantee uniform split
 - Some partitions will be larger than 100 blocks, some will be smaller
 - Use a hash function that hashes to $100 * f$ different values
 - f is called fudge factor, typically around 1.2
 - So we may consider $h1(a) = a \% 120$.
 - This is okay IF a is uniformly distributed

Hash Join: Issues




- Memory required ?
 - $R = 10000$ blocks, Memory = $M = 100$ blocks
 - 120 different partitions
 - During phase 1:
 - Need 1 block for storing R (*block by block*)
 - Need 120 blocks for storing each partition of R
 - At least 121 blocks of memory
 - We only have 100 blocks
- Typically need $\text{SQRT}(|R| * f)$ blocks of memory
- If R is 10000 blocks, and $f = 1.2$, need 110 blocks of memory
- If memory = 10000 blocks = $10000 * 4 \text{ KB} = 40\text{MB}$
 - Then, R can be as large as $10000 * 10000 / 1.2$ blocks = 333 GB

Hash Join: Issues




- What if we don't have enough memory ?
 - Recursive Partitioning
 - Rarely used, but can be done
- What if the hash function turns out to be bad ?
 - We used $h1(a) = a \% 100$
 - Turns out all *values of a* are multiple of 100
 - So $h1(a)$ is always = 0
- Called *hash-table overflow*
- Overflow avoidance: Use a good hash function
- Overflow resolution: Repartition using a different hash function

Hybrid Hash Join




- Motivation:
 - $R = 10000$ blocks, $S = 101$ blocks, $M = 100$ blocks
 - S doesn't fit in memory
- Approach: Use a hash function such that $S1 = 90$ blocks, and $S2 = 10$ blocks
(Try to keep more data records of S in memory) (Why S ?)
- Steps:
 - Read $S1$, and partition it
 - Write $S2$ to disk
 - Keep $S1$ in memory, and build a hash table on it
 - Read $R1$, and partition it
 - Write $R2$ to disk
 - Probe using $R1$ directly into the hash table
 - Saves huge amounts of I/O

So far



- Block Nested-loops join
 - Can always be applied irrespective of the join condition
- Index Nested-loops join
 - Only applies if an appropriate index exists
 - Very useful when we have selections that return small number of tuples
 - `select balance from customer, accounts where customer.name = "j. s." and customer.SSN = accounts.SSN`
- Hash joins
 - Join algorithm of choice when the relations are large
 - Only applies to equi-joins (since it is hash-based)
- Hybrid hash join
 - An optimization on hash join that is always implemented

Merge-Join (Sort-merge join)




- Pre-condition:
 - The relations must be sorted by the join attribute
 - If not sorted, can sort first, and then use this algorithms
- Called “sort-merge join” sometimes


```
select *  
from r, s  
where r.a1 = s.a1
```

Step:

1. Compare the tuples at *pr* and *ps*
2. Move pointers down the list
 - Depending on the join condition
3. Repeat




Merge-Join (Sort-merge join)



- Cost:
 - If the relations sorted, then just
 - $b_r + b_s$ block transfers, some seeks depending on memory size
 - What if not sorted ?
 - Then sort the relations first
 - In many cases, still very good performance
 - Typically comparable to hash join
- Observation:
 - The final join result will also be sorted on $a1$
 - This might make further operations easier to do
 - E.g. duplicate elimination


Group By and Aggregation



```
select a, count(b)  
from R  
group by a;
```

- Hash-based algorithm
- Steps:
 - Create a hash table on a , and keep the $count(b)$ so far
 - Read R tuples one by one
 - For a new R tuple, " r "
 - Check if $r.a$ exists in the hash table
 - If yes, increment the count
 - If not, insert a new value


Group By and Aggregation



```
select a, count(b)  
from R  
group by a;
```

- Sort-based algorithm
- Steps:
 - Sort R on a
 - Now all tuples in a single group are continuous
 - Read tuples of R (*sorted*) one by one and compute the aggregates


Duplicate Elimination



*select distinct a
from R ;*

- Best done using sorting – Can also be done using hashing
- Steps:
 - Sort the relation R
 - Read tuples of R in sorted order
 - $prev = null$;
 - for each tuple r in R (sorted)
 - if $r \neq prev$ then
 - Output r
 - $prev = r$
 - else
 - Skip r


Set operations




*(select * from R) **union** (select * from S) ;*
*(select * from R) **intersect** (select * from S) ;*
*(select * from R) union all (select * from S) ;*
*(select * from R) intersect all (select * from S) ;*

- Remember the rules about duplicates
- “union all”: just append the tuples of *R* and *S*
- “**union**”: append the tuples of *R* and *S*, and do **duplicate elimination**
- “*intersection*”: similar to joins
 - Find tuples of *R* and *S* that are identical on all attributes
 - Can use **hash-based** or **sort-based** algorithm

Evaluation of Expressions




select customer-name
from account a, customer c
where a.SSN = c.SSN and
a.balance < 2500




- Two options:
 - Materialization
 - Pipelining

Evaluation of Expressions



- Materialization
 - Evaluate each expression separately
 - Store its result on disk in *temporary relations*
 - Read it for next operation
- Pipelining
 - Evaluate multiple operators simultaneously
 - Skip the step of going to disk
 - Usually faster, but requires more memory
 - Also not always possible..
 - E.g. Sort-Merge Join

Pipelining



- Iterator Interface
- Each operator implements:
 - *init(): Initialize the state*
 - *get_next(): get the next tuple from the operator*
 - *close(): Finish and clean up*
- Sequential Scan:
 - *init(): open the file*
 - *get_next(): get the next tuple from file*
 - *close(): close the file*
- Execute by repeatedly calling *get_next()* at the root

Transformation Rules for RA Operations

Conjunctive Selection operations can cascade into individual Selection operations (and vice versa).

$$\sigma_{p \wedge q \wedge r}(R) = \sigma_p(\sigma_q(\sigma_r(R)))$$

● **Sometimes referred to as cascade of Selection.**

$$\begin{aligned} \sigma_{\text{branchNo}='B003' \wedge \text{salary}>15000}(\text{Staff}) = \\ \sigma_{\text{branchNo}='B003'}(\sigma_{\text{salary}>15000}(\text{Staff})) \end{aligned}$$

Transformation Rules for RA Operations

Commutativity of Selection.

$$\sigma_p(\sigma_q(R)) = \sigma_q(\sigma_p(R))$$

● For example:

$$\sigma_{\text{branchNo}='B003'}(\sigma_{\text{salary}>15000}(\text{Staff})) = \sigma_{\text{salary}>15000}(\sigma_{\text{branchNo}='B003'}(\text{Staff}))$$

Transformation Rules for RA Operations

In a sequence of Projection operations, only the last in the sequence is required.

$$\Pi_L \Pi_M \dots \Pi_N(R) = \Pi_L(R)$$

• For example:

$$\Pi_{\text{IName}} \Pi_{\text{branchNo, IName}}(\text{Staff}) = \Pi_{\text{IName}}(\text{Staff})$$

Transformation Rules for RA Operations

Commutativity of Selection and Projection.

- If predicate p involves only attributes in projection list, Selection and Projection operations commute:

$$\Pi_{A_i, \dots, A_m}(\sigma_p(R)) = \sigma_p(\Pi_{A_i, \dots, A_m}(R))$$

where $p \in \{A_1, A_2, \dots, A_m\}$

- For example:

$$\Pi_{fName, lName}(\sigma_{lName='Beech'}(Staff)) = \sigma_{lName='Beech'}(\Pi_{fName, lName}(Staff))$$

Transformation Rules for RA Operations

Commutativity of Theta join (and Cartesian product).

$$R \bowtie_p S = S \bowtie_p R$$

$$R \times S = S \times R$$

Rule also applies to Equijoin and Natural join. For example:

$$\text{Staff} \bowtie_{\text{staff.branchNo}=\text{branch.branchNo}} \text{Branch} =$$

$$\text{Branch} \bowtie_{\text{staff.branchNo}=\text{branch.branchNo}} \text{Staff}$$

Transformation Rules for RA Operations

Commutativity of Selection and Theta join (or Cartesian product).

- If selection predicate involves only attributes of one of join relations, Selection and Join (or Cartesian product) operations commute:

$$\sigma_p(R \bowtie_r S) = (\sigma_p(R)) \bowtie_r S$$

$$\sigma_p(R \times S) = (\sigma_p(R)) \times S$$

where $p \in \{A_1, A_2, \dots, A_n\}$

Transformation Rules for RA Operations

- If selection predicate is conjunctive predicate having form $(p \wedge q)$, where p only involves attributes of R , and q only attributes of S , Selection and Theta join operations commute as:

$$\sigma_{p \wedge q}(R \bowtie_r S) = (\sigma_p(R)) \bowtie_r (\sigma_q(S))$$

$$\sigma_{p \wedge q}(R \times S) = (\sigma_p(R)) \times (\sigma_q(S))$$

Transformation Rules for RA Operations

• For example:

$$\sigma_{\text{position}='Manager' \wedge \text{city}='London'}(\text{Staff} \bowtie \text{Branch}) =$$
$$(\sigma_{\text{position}='Manager'}(\text{Staff})) \bowtie_{\text{Staff.branchNo}=\text{Branch.branchNo}} (\sigma_{\text{city}='London'}(\text{Branch}))$$

Transformation Rules for RA Operations

Commutativity of Projection and Theta join (or Cartesian product).

- If projection list is of form $L = L_1 \cup L_2$, where L_1 only has attributes of R , and L_2 only has attributes of S , provided join condition only contains attributes of L , Projection and Theta join commute:

$$\Pi_{L_1 \cup L_2}(R \bowtie_r S) = (\Pi_{L_1}(R)) \bowtie_r (\Pi_{L_2}(S))$$

Transformation Rules for RA Operations

- If join condition contains additional attributes not in L ($M = M_1 \cup M_2$ where M_1 only has attributes of R, and M_2 only has attributes of S), a final projection operation is required:

$$\Pi_{L_1 \cup L_2}(R \bowtie_r S) = \Pi_{L_1 \cup L_2}((\Pi_{L_1 \cup M_1}(R)) \bowtie_r (\Pi_{L_2 \cup M_2}(S)))$$

Transformation Rules for RA Operations

- For example:

$$\Pi_{\text{position, city, branchNo}} (\text{Staff} \bowtie_{\text{Staff.branchNo=Branch.branchNo}} \text{Branch}) =$$

$$(\Pi_{\text{position, branchNo}} (\text{Staff})) \bowtie_{\text{Staff.branchNo=Branch.branchNo}} ($$

$$\Pi_{\text{city, branchNo}} (\text{Branch}))$$

- and using the latter rule:

$$\Pi_{\text{position, city}} (\text{Staff} \bowtie_{\text{Staff.branchNo=Branch.branchNo}} \text{Branch}) =$$

$$\Pi_{\text{position, city}} ((\Pi_{\text{position, branchNo}} (\text{Staff})) \bowtie_{\text{Staff.branchNo=Branch.branchNo}}$$

$$(\Pi_{\text{city, branchNo}} (\text{Branch})))$$

Transformation Rules for RA Operations

**Commutativity of Union and Intersection
(but not set difference).**

$$R \cup S = S \cup R$$

$$R \cap S = S \cap R$$

Transformation Rules for RA Operations

Commutativity of Selection and set operations (Union, Intersection, and Set difference).

$$\sigma_p(R \cup S) = \sigma_p(S) \cup \sigma_p(R)$$

$$\sigma_p(R \cap S) = \sigma_p(S) \cap \sigma_p(R)$$

$$\sigma_p(R - S) = \sigma_p(S) - \sigma_p(R)$$

Transformation Rules for RA Operations

Commutativity of Projection and Union.

$$\Pi_L(R \cup S) = \Pi_L(S) \cup \Pi_L(R)$$

Associativity of Union and Intersection (but not Set difference).

$$(R \cup S) \cup T = S \cup (R \cup T)$$

$$(R \cap S) \cap T = S \cap (R \cap T)$$

Transformation Rules for RA Operations

Associativity of Theta join (and Cartesian product).

- **Cartesian product and Natural join are always associative:**

$$(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$$

$$(R \times S) \times T = R \times (S \times T)$$

- **If join condition q involves attributes only from S and T , then Theta join is associative:**

$$(R \bowtie_p S) \bowtie_{q \wedge r} T = R \bowtie_{p \wedge r} (S \bowtie_q T)$$

Transformation Rules for RA Operations

● For example:

(Staff ⋈_{Staff.staffNo=PropertyForRent.staffNo} PropertyForRent)

⋈_{ownerNo=Owner.ownerNo ∧ staff.lName=Owner.lName} Owner =

Staff ⋈_{staff.staffNo=PropertyForRent.staffNo ∧ staff.lName=lName}


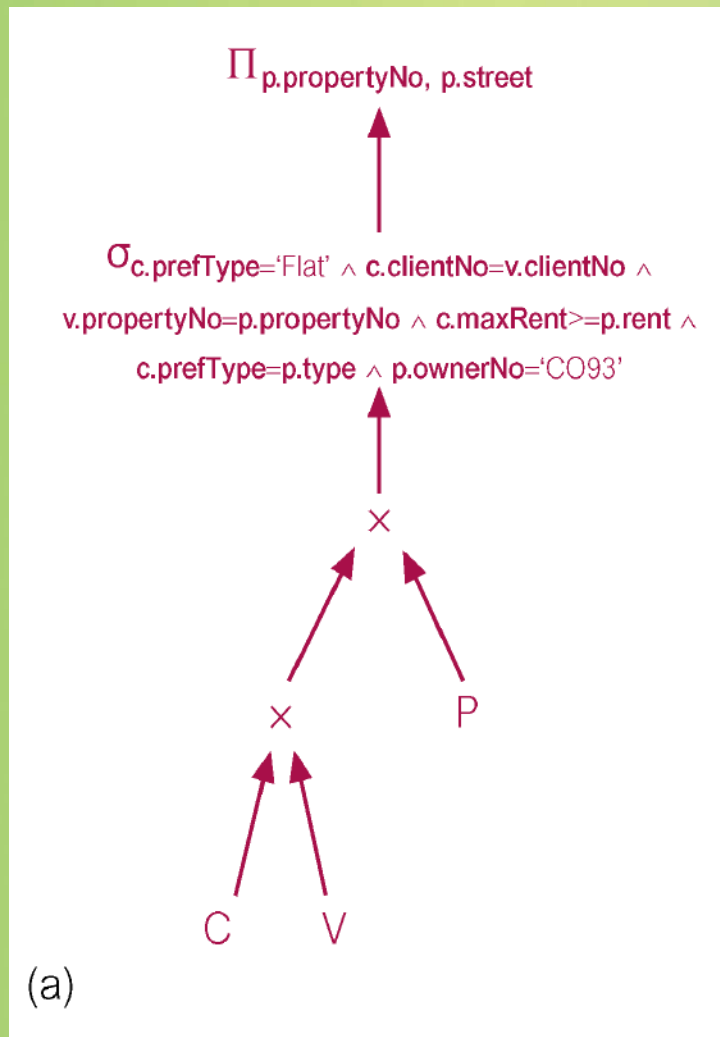
(PropertyForRent ⋈_{ownerNo} Owner)

Example 23.3 Use of Transformation Rules

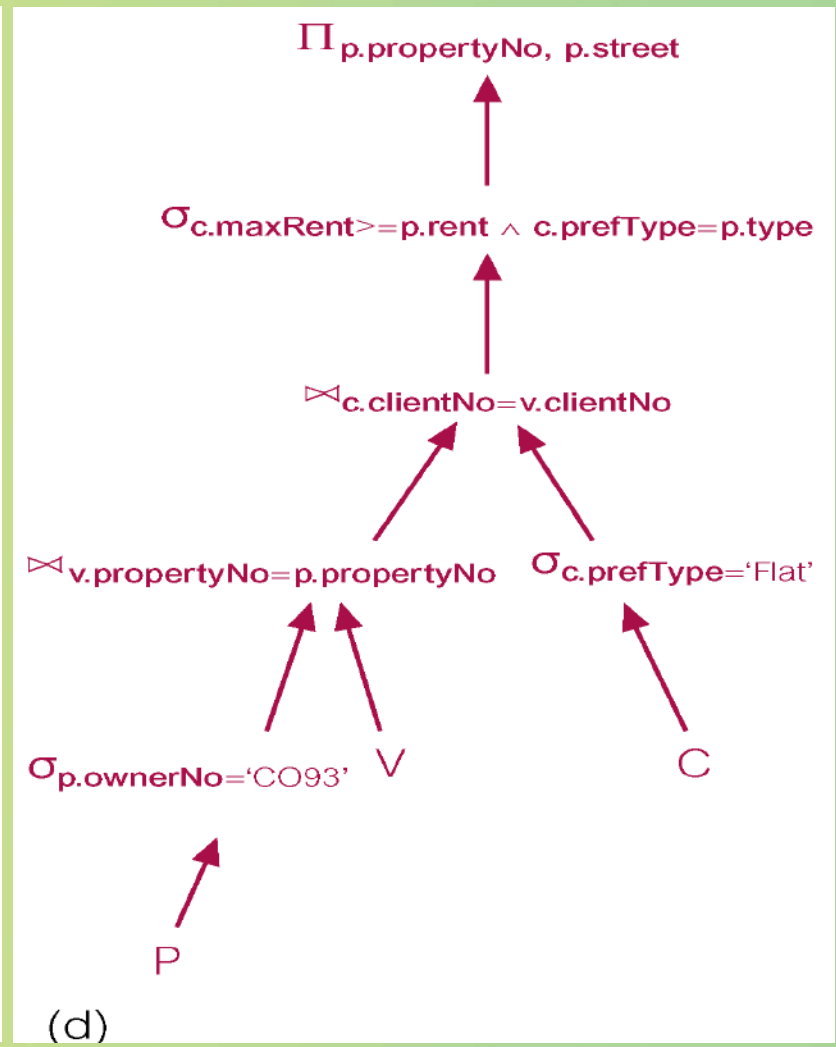
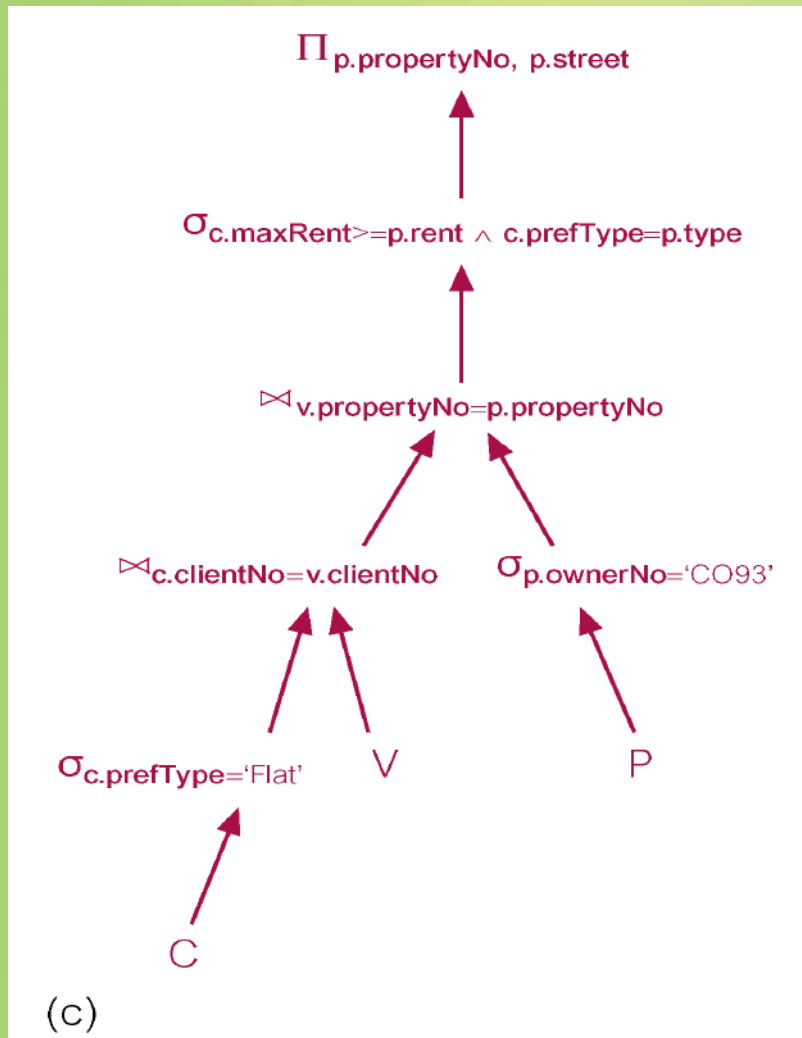
For prospective renters of flats, find properties that match requirements and owned by CO93.

```
SELECT p.propertyNo, p.street  
FROM Client c, Viewing v, PropertyForRent p  
WHERE  c.prefType = 'Flat' AND  
        c.clientNo = v.clientNo AND  
        v.propertyNo = p.propertyNo AND  
        c.maxRent >= p.rent AND  
        c.prefType = p.type AND  
        p.ownerNo = 'CO93';
```

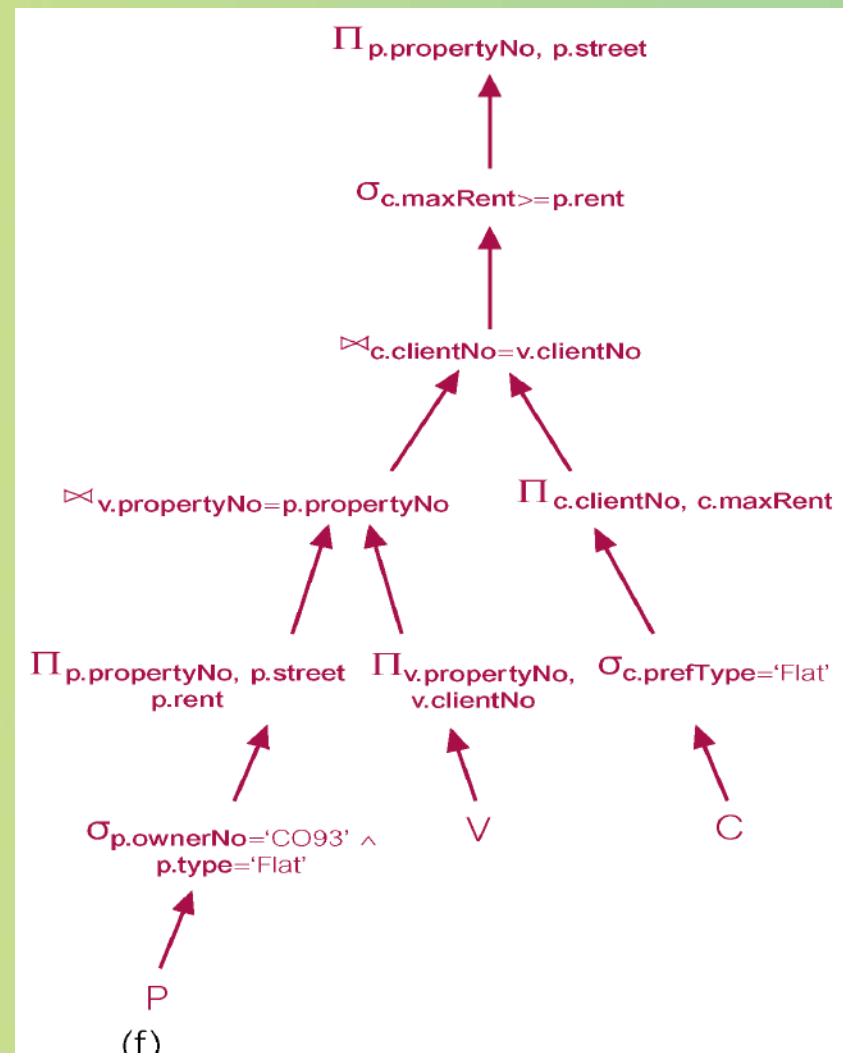
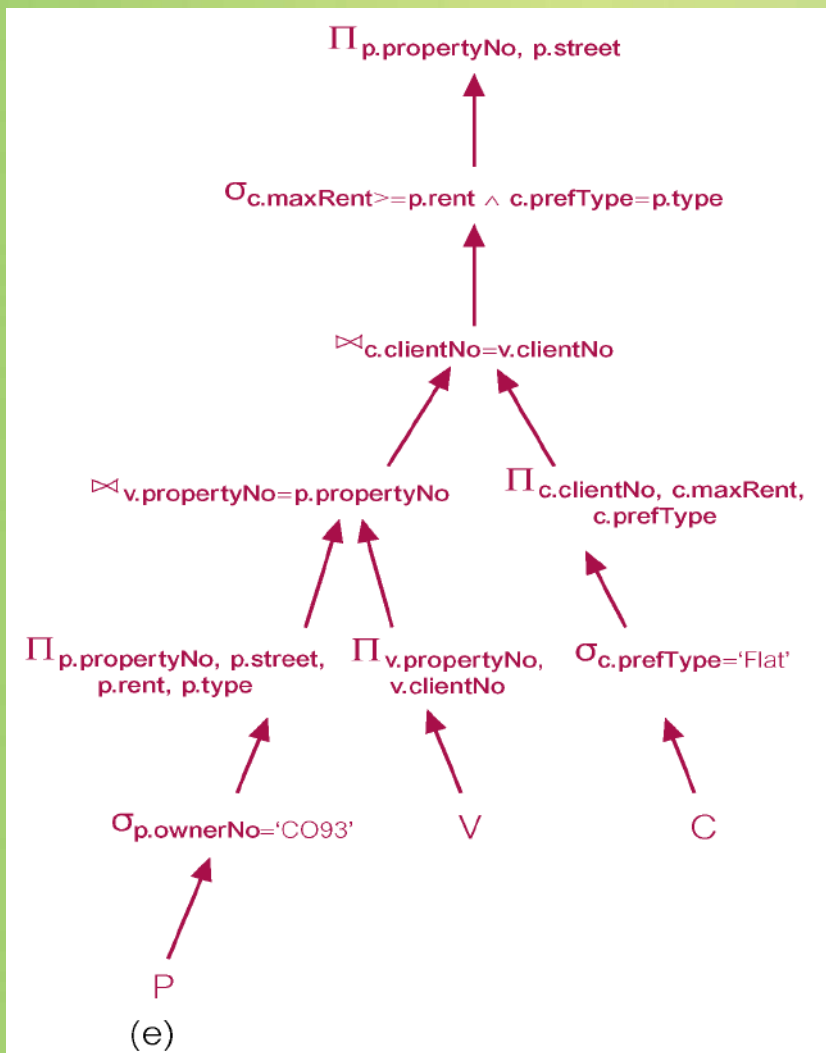
Example 23.3 Use of Transformation Rules



Example 23.3 Use of Transformation Rules



Example 23.3 Use of Transformation Rules



Heuristic Processing Strategies

- **Perform Selection operations as early as possible.**
 - **Keep predicates on same relation together.**
- **Combine Cartesian product with subsequent Selection whose predicate represents join condition into a Join operation.**
- **Use associativity of binary operations to rearrange leaf nodes so leaf nodes with most restrictive Selection operations executed first.**

Heuristical Processing Strategies

- **Perform Projection as early as possible.**
 - Keep projection attributes on same relation together.
- **Compute common expressions once.**
 - If common expression appears more than once, and result not too large, store result and reuse it when required.
 - Useful when querying views, as same expression is used to construct view each time.

Cost Estimation for RA Operations

- Many different ways of implementing RA operations.
- Aim of QO is to choose most efficient one.
- Use formulae that estimate costs for a number of options, and select one with lowest cost.
- Consider only cost of disk access, which is usually dominant cost in QP.
- Many estimates are based on cardinality of the relation, so need to be able to estimate this.

Database Statistics

- Success of estimation depends on amount and currency of statistical information DBMS holds.
- Keeping statistics current can be problematic.
- If statistics updated every time tuple is changed, this would impact performance.
- DBMS could update statistics on a periodic basis, for example nightly, or whenever the system is idle.

Typical Statistics for Relation R

nTuples(R) - number of tuples in R.

bFactor(R) - blocking factor of R.

nBlocks(R) - number of blocks required to store R:

$$\text{nBlocks(R)} = \lceil \text{nTuples(R)} / \text{bFactor(R)} \rceil$$

Typical Statistics for Attribute A of Relation R

$n\text{Distinct}_A(R)$ - number of distinct values that appear for attribute A in R.

$\min_A(R)$, $\max_A(R)$ - minimum and maximum possible values for attribute A in R.

$SC_A(R)$ - *selection cardinality* of attribute A in R.

Average number of tuples that satisfy an equality condition on attribute A.

Statistics for Multilevel Index I on Attribute A

$nLevels_A(I)$ - number of levels in I.

$nLfBlocks_A(I)$ - number of leaf blocks in I.

References

- SQL query optimization
<http://redbook.cs.berkeley.edu/redbook3/lec7.html>