

Transactions

Overview

- Transaction: A sequence of database actions enclosed within special tags

T1: transfers \$50 from A to B

begin T1

read(A)

A = A -50

write(A)

read(B)

B=B+50

write(B)

end T1

Overview (cont'd)

- Properties:
 - **Atomicity**: Entire transaction or nothing
 - **Consistency**: Transaction, executed completely, takes database from one consistent state to another
 - **Isolation**: Concurrent transactions appear to run in isolation
 - **Durability**: Effects of committed transactions are not lost
- Consistency: Transaction programmer needs to guarantee that
 - DBMS can do a few things, e.g., enforce constraints on the data
- Rest: DBMS guarantees

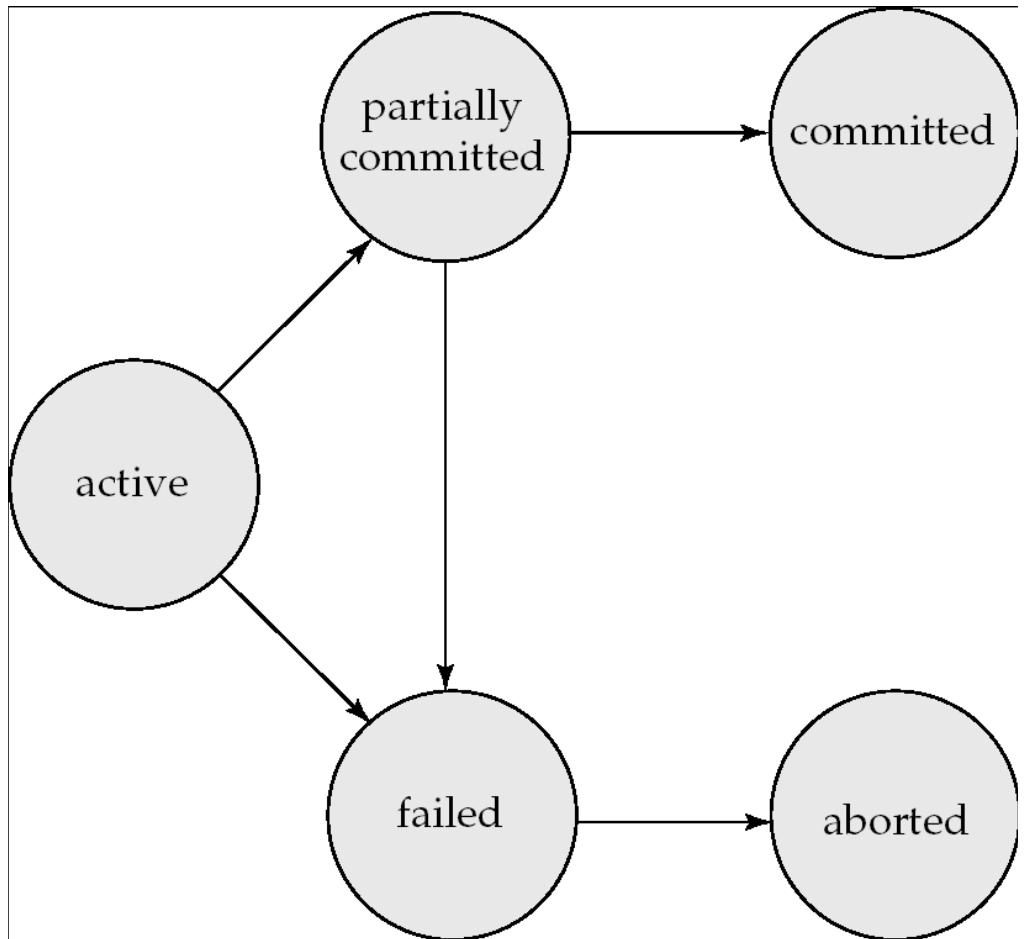
Assumptions and Goals

- Assumptions:
 - The system can crash at any time
 - Similarly, the power can go out at any point
 - Contents of the main memory won't survive a crash, or power outage
 - **Disks are durable. They might stop, but data is not lost.**
 - Disks only guarantee *atomic sector writes*, nothing more
 - Transactions are by themselves consistent
- Goals:
 - Guaranteed durability, atomicity
 - As much concurrency as possible, while not compromising isolation and/or consistency
 - Two transactions updating the same account balance... NO
 - Two transactions updating different account balances... YES

Next

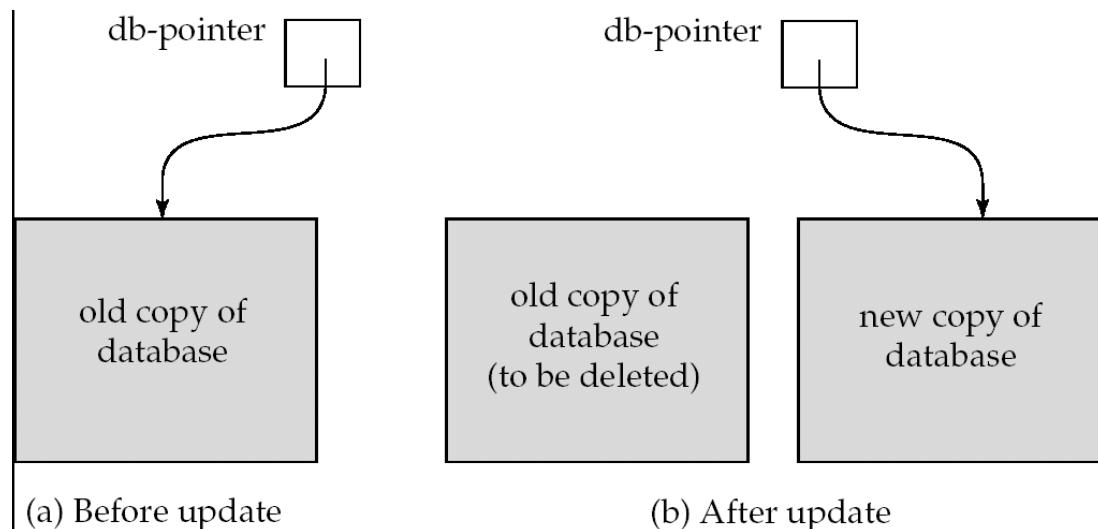
- States of a transaction
- A simple solution called *shadow copy*
 - Satisfies Atomicity, Durability, and Consistency, but no Concurrency
 - Very inefficient

Transaction states



Shadow Copy

- Make updates on a copy of the database.
- Switch pointers atomically after done.
 - Some *text editors* work this way



Shadow Copy

Atomicity:

- As long as the DB pointer switch is atomic.
- Okay if DB pointer is in a single block

Concurrency:

- No.

Isolation:

- No concurrency, so isolation is guaranteed.

Durability:

- Assuming disk is durable (we will assume this for now).

Very inefficient:

- Databases tend to be very large. Making extra copies not feasible. Further, no concurrency.

Next

- Concurrency control schemes
 - A CC scheme is used to guarantee that concurrency does not lead to problems
 - For now, we will assume durability is not a problem
 - No crashes
 - Though transactions may still abort
- Schedules
- When is concurrency okay ?
 - Serial schedules
 - Serializability

A Schedule

Transactions:

T1: transfers \$50 from A to B

T2: transfers 10% of A to B

Database constraint: A + B is constant (*checking+saving accts*)

T1	T2	Effect:	Before	After
read(A) A = A - 50 write(A)	read(A) tmp = A * 0.1 A = A - tmp write(A)	A	100	45
read(B) B=B+50 write(B)	read(B) B = B+ tmp write(B)	B	50	105

Each transaction obeys the constraint.

This schedule does too.

Schedules

- A *schedule* is simply a (possibly interleaved) execution sequence of transaction instructions
- *Serial Schedule*: A schedule in which transactions appear one after the other
 - No interleaving
- Serial schedules satisfy isolation and consistency
 - Since each transaction by itself does not introduce inconsistency

Example Schedule

- Another “serial” schedule:

T1	T2	Effect:	Before	After
read(A) A = A -50 write(A) read(B) B=B+50 write(B)	read(A) tmp = A*0.1 A = A - tmp write(A) read(B) B = B+ tmp write(B)	A B	100 50	40 110
		Consistent ? Constraint is satisfied.		

Another schedule

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A * 0.1 A = A - tmp write(A)
read(B) B = B + 50 write(B)	read(B) B = B + tmp write(B)

Is this schedule okay ?

Lets look at the final effect

Effect:	Before	After
A	100	45
B	50	105

Consistent.

So this schedule is okay too.

Another schedule

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A * 0.1 A = A - tmp write(A)
read(B) B = B + 50 write(B)	read(B) B = B + tmp write(B)

Is this schedule okay ?

Let's look at the final effect...

Effect:	Before	After
A	100	45
B	50	105

Further, the effect same as the serial schedule 1.

Called Serializable

Example Schedules (Cont.)

A “bad” schedule

T1	T2	Effect:	Before	After
read(A) A = A -50	read(A) tmp = A*0.1 A = A – tmp write(A) read(B)	A	100	50
write(A) read(B) B=B+50 write(B)	B = B+ tmp write(B)	B	50	60
		<u>Not consistent</u>		

Serializability

- A schedule is called *serializable* if its final effect is the same as that of a *serial schedule*
- Serializability → schedule is fine and does not result in inconsistent database
 - Since serial schedules are fine
- Non-Serializable schedules are likely to result in inconsistent databases
- We will ensure serializability
 - Typically relaxed in real high-throughput environments

Serializability

- Not possible to look at all $n!$ serial schedules to check if the effect is the same
 - Instead we ensure serializability by allowing or not allowing certain schedules
- Conflict serializability
- View serializability
 - View serializability allows more schedules

Conflict Serializability

Two read/write instructions “conflict” if

- They are by different transactions
- They operate on the same data item
- **At least one** is a “write” instruction

Why do we care ?

- If two read/write instructions don’t conflict, they can be “swapped” without any change in the final effect
- However, if they conflict they CAN’T be swapped without changing the final effect

Conflicting Instructions

- Instructions I_i and I_j of transactions T_i and T_j respectively, **conflict** if and only if there exists some item Q accessed by both I_i and I_j , and at least one of these instructions wrote Q .
- 1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict.
 2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict
 4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict
- Intuitively, a conflict between I_i and I_j forces a (logical) temporal order between them.
 - If I_i and I_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

Equivalence by Swapping

T1	T2	T1	T2
read(A) A = A -50 write(A)	read(A) tmp = A*0.1 A = A - tmp write(A)	read(A) A = A -50 write(A)	read(A) tmp = A*0.1 A = A - tmp
read(B) B=B+50 write(B)	read(B) B = B+ tmp write(B)	read(B) B=B+50 write(B)	write(A)
			read(B) B = B+ tmp write(B)

Effect:	Before	After
A	100	45
B	50	105

==

Effect:	Before	After
A	100	45
B	50	105

Equivalence by Swapping

T1	T2
<pre>read(A) A = A -50 write(A)</pre>	<pre>read(A) tmp = A * 0.1 A = A - tmp write(A)</pre>
<pre>read(B) B=B+50 write(B)</pre>	<pre>read(B) B = B + tmp write(B)</pre>

T1	T2
<pre>read(A) A = A -50 write(A)</pre>	<pre>read(A) tmp = A * 0.1 A = A - tmp write(A)</pre>
<pre>read(B) B=B+50 write(B)</pre>	<p>read(B)</p> <p>write(B)</p> <p>B = B + tmp</p> <p>write(B)</p>

Effect:	Before	After
A	100	45
B	50	105

! ==

Effect:	Before	After
A	100	45
B	50	55

Conflict Serializability

Conflict-equivalent schedules:

- If S can be transformed into S' through a series of swaps, S and S' are called *conflict-equivalent*
- *conflict-equivalent guarantees same final effect on the database*

A schedule S is conflict-serializable if it is conflict-equivalent to a serial schedule

(a)	T_1	T_2
	<pre>read_item(X); X:=X-N; write_item(X); read_item(Y); Y:=Y+N; write_item(Y);</pre>	<pre>read_item(X); X:=X+M; write_item(X);</pre>

(serial) schedule A

(b)	T_1	T_2
		<pre>read_item(X); X:=X+M; write_item(X);</pre>

(serial) schedule B

(c)	T_1	T_2
	<pre>read_item(X); X:=X-N;</pre>	<pre>read_item(X); X:=X+M;</pre>

(non-Serializable) schedule C

	T_1	T_2
	<pre>read_item(X); X:=X-N; write_item(X);</pre>	<pre>read_item(X); X:=X+M; write_item(X);</pre>

(Serializable) schedule D



Equivalence by Swapping

T1	T2	T1	T2
read(A) A = A -50 write(A)	read(A) tmp = A * 0.1 A = A - tmp write(A)	read(A) A = A -50 write(A)	read(A) tmp = A * 0.1 A = A - tmp
read(B) B=B+50 write(B)	read(B) B=B+50 write(B)	read(B) B=B+50 write(B)	write(A)
	read(B) B = B + tmp write(B)		read(B) B = B + tmp write(B)

Effect:	Before	After
A	100	45
B	50	105

==

Effect:	Before	After
A	100	45
B	50	105

Equivalence by Swapping

T1	T2
read(A) A = A -50 write(A)	read(A) tmp = A *0.1 A = A – tmp write(A)
read(B) B=B+50 write(B)	read(B) B=B+50 write(B)
	read(B) B = B+ tmp write(B)

T1	T2
read(A) A = A -50 write(A)	read(B) B=B+50 write(B)
	read(A) tmp = A *0.1 A = A – tmp write(A)
	read(B) B = B+ tmp write(B)

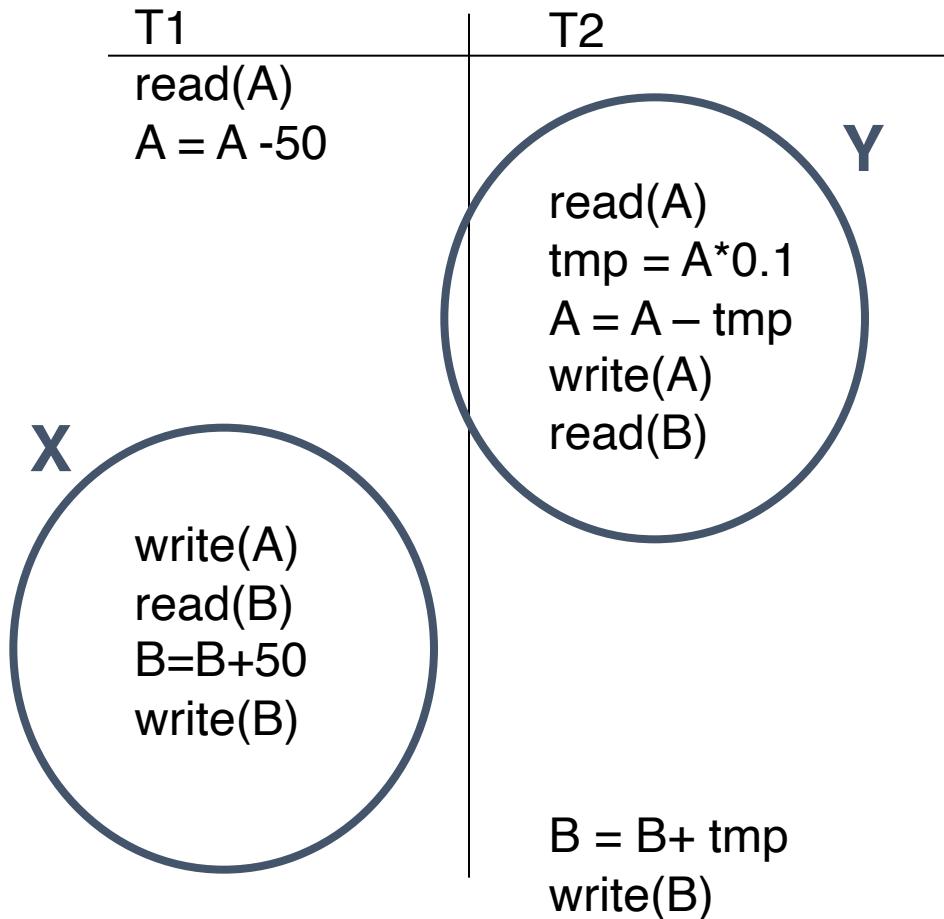
Effect:	Before	After
A	100	45
B	50	105

==

Effect:	Before	After
A	100	45
B	50	105

Example Schedules (Cont.)

A “bad” schedule



Can't move Y below X
read(B) and write(B) conflict

Other options don't work either

Not Conflict Serializable

Serializability

- In essence, following set of instructions is not conflict-serializable:

T_3	T_4
read(Q)	write(Q)
write(Q)	

View-Serializability

- Similarly, the following schedule is not conflict-serializable

T_3	T_4	T_6
read(Q)		
write(Q)	write(Q)	write(Q)

T_3	T_4	T_6
read(Q)	write(Q)	
		write(Q)

- It is serializable in terms of the final results
 - Intuitively, this is because the *conflicting write instructions* don't matter
 - The final write is the only one that matters

Schedules S1 and S2 are **view equivalent** if:

If T_i reads initial value of A in S1, then T_i also reads initial value of A in S2

If T_i reads value of A written by T_j in S1, then T_i also reads value of A written by T_j in S2

If T_i writes final value of A in S1, then T_i also writes final value of A in S2

Other notions of serializability

The schedule below produces same outcome as the serial schedule $\langle T_1, T_5 \rangle$, yet is not conflict equivalent or view equivalent to it.

T_1	T_5
read(A)	
$A := A - 50$	
write(A)	
	read(B)
	$B := B - 10$
	write(B)
	read(B)
	$B := B + 50$
	write(B)
	read(A)
	$A := A + 10$
	write(A)

- Not conflict-serializable or view-Serializable, but Serializable
- Mainly because of the +/- only operations
 - Requires analysis of the actual operations, not just read/write operations
- Most high-performance transaction systems will allow these

Testing for conflict- serializability

Given a schedule, determine if it is conflict-
Serializable

Draw a *precedence-graph* over the
transactions

- A directed edge from T1 and T2, if they have conflicting instructions, and T1's conflicting instruction comes first

If there is a cycle in the graph, not conflict-
Serializable

- Can be checked in at most $O(n+e)$ time, where n is the number of vertices, and e is the number of edges

If there is none, conflict-Serializable

Testing for view-serializability is NP-hard.

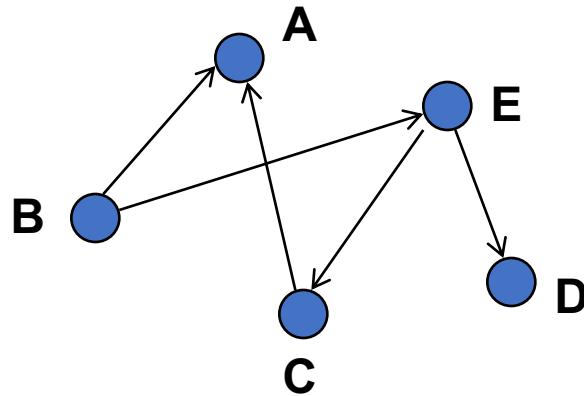
Precedence Graph

- Testing for conflict serializability
- $T = \{T_1, \dots, T_n\}$, a set of transactions.
- The Precedence graph of a schedule S is a directed graph $G(S) = (V, E)$, where
 - $V = \{T_1, \dots, T_n\}$ is a set of vertices;
 - E consists of edges (T_i, T_j) if
 - one of T_i 's operations precedes and conflicts with one of T_j 's operations in S .

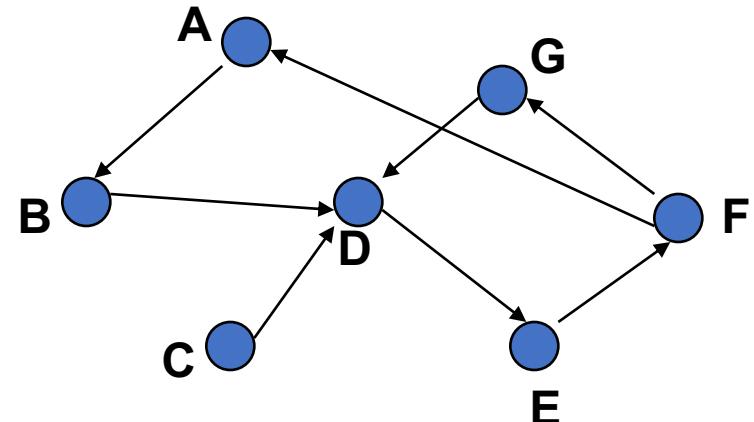
■ Cycle in a graph

- ★ A cycle is a path that starts and terminates at the same node

■ Examples



Does not contain cycle

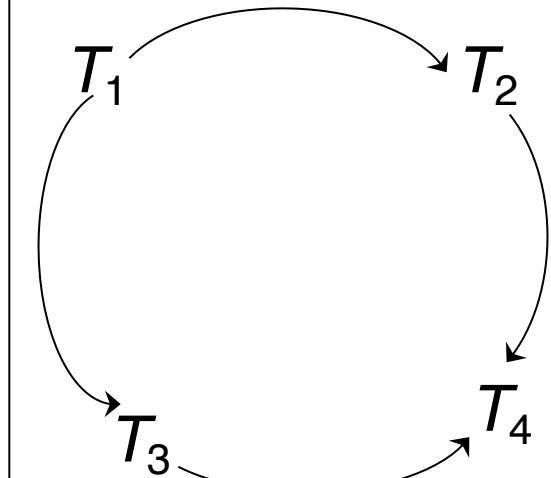


Contains cycles

■ **Theorem:** Schedule is conflict serializable if and only if its dependence graph is acyclic

Example Schedule (Schedule A) + Precedence Graph

T_1	T_2	T_3	T_4	T_5
	read(X)			
read(Y) read(Z)				read(V) read(W) read(W)
	read(Y) write(Y)	write(Z)		
read(U)			read(Y) write(Y) read(Z) write(Z)	
read(U) write(U)				



Recoverability

- Serializability is good for consistency

- But what if transactions fail ?

- T2 has already committed
 - A user might have been notified
- Now T1 abort creates a problem
 - T2 has seen its effect, so just aborting T1 is not enough. T2 must be aborted as well (and possibly restarted)
 - But T2 is *committed*

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A * 0.1 A = A - tmp write(A) COMMIT
read(B) B=B+50 write(B) ABORT	

■ Database must ensure that schedules are recoverable.

Recoverability

- Recoverable schedule: If T1 has read something that T2 has written, T2 must commit before T1
 - Otherwise, if T1 commits, and T2 aborts, we have a problem
- Cascading rollbacks: If T10 aborts, T11 must abort, and hence T12 must abort and so on.

T_{10}	T_{11}	T_{12}
read(A) read(B) write(A)	read(A) write(A)	read(A)

Cascadeless Schedules

- **Cascadeless schedules** — cascading rollbacks cannot occur; for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless

Notes

- We discussed:
 - Serial schedules, serializability
 - Conflict-serializability, view-serializability
 - How to check for conflict-serializability
 - Recoverability, cascade-less schedules
- We haven't discussed:
 - How to guarantee serializability ?
 - Allowing transactions to run, and then aborting them if the schedule wasn't serializable is clearly not the way to go
 - We instead use schemes to guarantee that the schedule will be conflict-serializable