


# Concurrency Control

# Concurrency Control

---

- A database must provide a mechanism that will ensure that all possible schedules are
  - either **conflict or view serializable**, and
  - **recoverable** and preferably cascadeless
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Testing a schedule for serializability *after* it has executed is a little too late!
- **Goal** – to develop concurrency control protocols that will **assure serializability**



# Why Concurrency Control ?

---

- Three concurrency problems
  - Lost update
  - Temporary update (dirty read): uncommitted dependency
  - Incorrect summary
  - Unrepeatable read

# Lost Update Problem

$X=1000, Y=200$

$N=500, M=300$

$X'=500$

$X=500$

T1

read(X);

$X=X-N$ ;

write(X);

read(Y);

$Y=Y+N$ ;

write(Y);

T2

read(X);  $X'=1000$

$X=X+M$ ;  $X'=1300$

write(X);  $X=1300$

# Temporary Update Problem

$X=1000, Y=200$

$N=500, M=300$

$X'=500$

$X=500$

T1

read(X);

$X=X-N$ ;

write(X);

T2

read(X);     $X'=500$

$X=X+M$ ;

write(X);     $X=800$

read(Y);

fails

# Incorrect Summary Problem

$X=1000, Y=200$   
 $N=500$

$X'=500$

$X=500$

T1

read(X);  
 $X=X-N$ ;  
write(X);

read(Y);  
 $Y=Y+N$ ;  
write(Y);

T2

sum=0

read(X);  
sum=sum+X      sum=500  
read(Y)  
sum=sum+Y      sum=700

# Solution for Lost Update Problem

■ Locking techniques

X=1000, Y=200  
N=500, M=300

	<u>T1</u>	<u>T2</u>
	readlock(X)	
	writelock(X)	
	read(X);	
X'=500	X=X-N;	
X=500	write(X)	
	unlock(X)	
		readlock(X)
		writelock(X)
		read(X);
		X=X+M;
		write(X);
		unlock(X)
	read(Y);	
	Y=Y+N;	
	write(Y);	

X'=500  
X'=800  
X=800

# Approach, Assumptions

---

- Approach
  - Guarantee **conflict-serializability** by allowing certain types of concurrency
    - Lock-based
- Assumptions:
  - Durability is not a problem
    - no crashes
    - Though transactions may still abort
- Goal:
  - Serializability
  - Minimize the bad effect of aborts (cascade-less schedules only)





# Lock-based Protocols

---

- A transaction *must* get a *lock* before operating on the data
- Two types of locks:
  - *Shared* (S) locks (also called *read locks*)
    - Obtained if we want to only read an item
  - *Exclusive* (X) locks (also called *write locks*)
    - Obtained for updating a data item

Lock instructions



**New instructions**

# Lock-based Protocols



- Lock requests are made to the *concurrency control manager*
  - It decides whether to *grant* a lock request
- T1 asks for a lock on data item A, and T2 currently has a lock on it ?
  - Depends
- If *compatible*, grant the lock. Otherwise T1 waits in a *queue*.

<u>T2 lock type</u>	<u>T1 lock type</u>	<u>Should allow ?</u>
Shared	Shared	YES
Shared	Exclusive	NO
Exclusive	Shared/Exclusive	NO

# Lock-based Protocols

- How do we actually use this to guarantee serializability/recoverability ?
  - Not enough just to take locks when you need to read/write something

T1

lock-X(B)  
read(B)  
 $B \leftarrow B - 50$   
write(B)  
unlock(B)

Assume  $A=100$ ,  $B=200$

lock-S(A), lock-S(B)  
Display(A+B)  
unlock(A), unlock(B)

$A=100$ ,  $B=150$ ,

lock-X(A)  
read(A)  
 $A \leftarrow A + 50$   
write(A)  
unlock(A)

# 2-Phase Locking Protocol (2PL)



Phase 1:  
Growing  
phase

Transaction may obtain locks



Phase 2:  
Shrinking  
phase

Transaction may only  
release locks


T1

lock-X(B)  
read(B)  
 $B \leftarrow B - 50$   
write(B)  
unlock(B)

lock-X(A)  
read(A)  
 $A \leftarrow A + 50$   
write(A)  
unlock(A)

# 2 Phase Locking

- Example: T1 in 2PL



## 2 Phase Locking

- Can be shown that this achieves *conflict-serializability*
- Guarantees *conflict-serializability*, but not cascade-less recoverability

T1	T2	T3
lock-X(A), lock-S(B) read(A) read(B) write(A) unlock(A), unlock(B)	lock-X(A) read(A) write(A) unlock(A) Commit	lock-S(A) read(A) Commit
<action fails>		

# 2 Phase Locking



---

- Guaranteeing just recoverability:
  - If T2 reads a dirty data of T1 (i.e., T1 has not committed), then T2 **can't** commit unless T1 either commits or aborts
  - If T1 commits, T2 can proceed with committing
  - If T1 aborts, T2 must abort
    - So cascades still happen



# Strict 2PL

- Release *exclusive locks* only at *the very end*, just before commit or abort

T1	T2	T3
lock-X(A), lock-S(B) read(A) read(B) write(A) unlock(A), unlock(B)	lock-X(A) read(A) write(A) unlock(A) Commit	lock-S(A) read(A) Commit
<action fails>		

Strict 2PL  
will not  
allow that

Works. Guarantees cascade-less and recoverable schedules.

# Strict 2PL



---

- Release *exclusive* locks only at the very end, just before commit or abort
  - Read locks are not important
- Rigorous 2PL: Release *both exclusive and read locks* only at the very end
  - The serializability order === the commit order
  - More intuitive behavior for the users
    - No difference for the system

# Strict 2PL



---

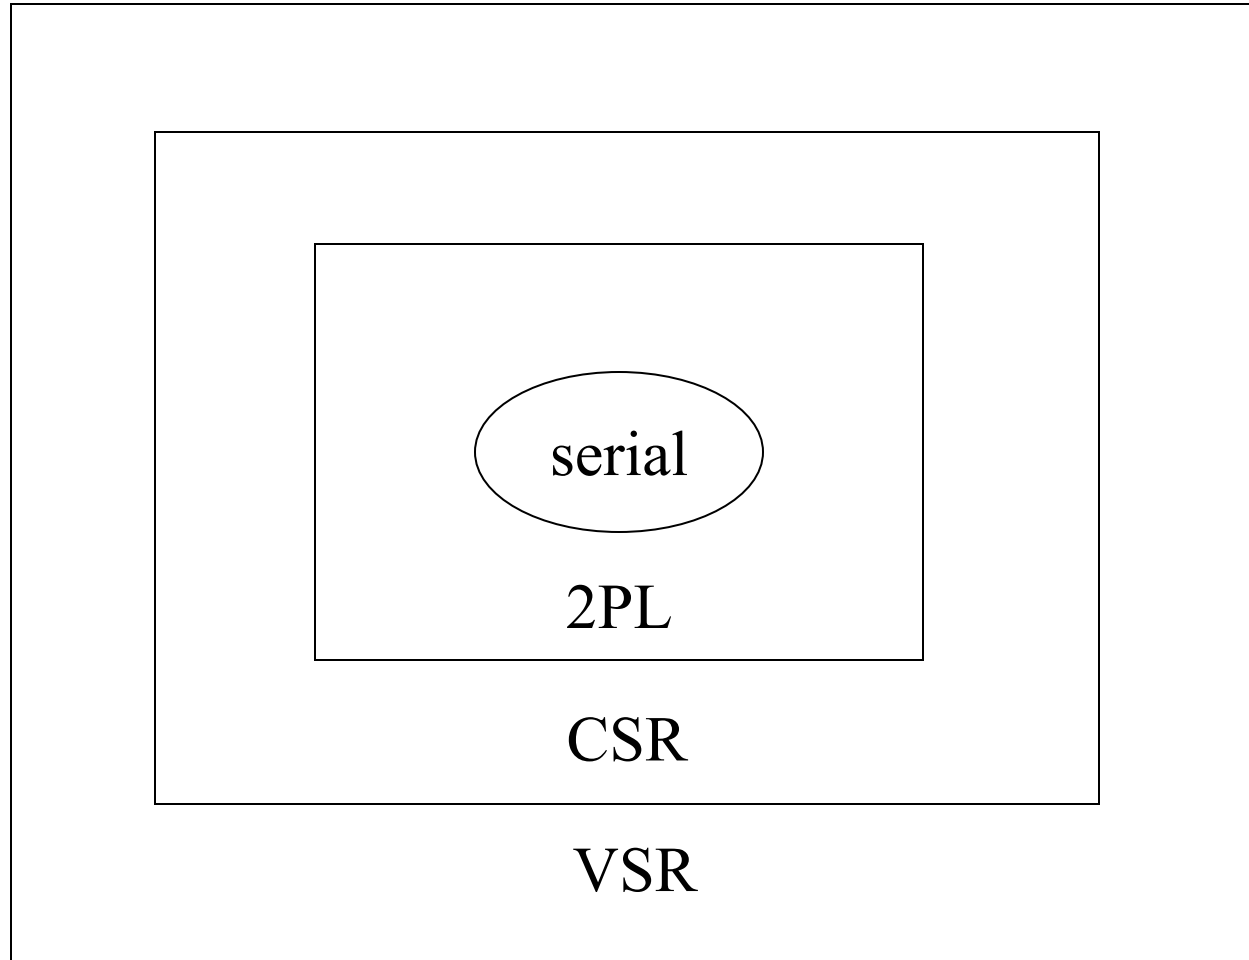
- Lock conversion:
  - Transaction might not be sure what it needs a write lock on
  - Start **with a S lock**
  - *Upgrade to an X lock later* if needed
  - Doesn't change any of the other properties of the protocol

# Recap

---

- Concurrency Control Scheme
  - A way to guarantee serializability, recoverability
- Lock-based protocols
  - Use *locks* to prevent multiple transactions accessing the same data items
- 2 Phase Locking
  - Locks acquired during *growing phase*, released during *shrinking phase*

# Hierarchy of Serializable Schedules



# Other CC Schemes

---

- Time-stamp based
  - Transactions are issued **time-stamps** when they enter the system
  - The time-stamps determine the *serializability* order
  - If T1 entered before T2, then T1 should be before T2 in the serializability order
    - $timestamp(T1) < timestamp(T2)$
  - If **T1 wants to read data item A**
    - If any transaction with **larger time-stamp wrote** that data item, then this operation is not permitted, and T1 is *aborted*
  - If **T1 wants to write data item A**
    - If a transaction with **larger time-stamp already read that** data item or written it, then the write is *rejected* and T1 is aborted
  - Aborted transaction are restarted with a new timestamp
    - Possibility of *starvation*

# Other CC Schemes

- Time-stamp based
  - Example (Timestamps  $T_1 < T_2 < T_3 < T_4 < T_5$ )

$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
read( $Y$ )	read( $Y$ )	write( $Y$ ) write( $Z$ )		write( $X$ )
	read( $X$ ) abort	write( $Z$ ) abort		read( $Z$ )
read( $X$ ) ?				write( $Y$ ) write( $Z$ )

# Other CC Schemes

---

- Time-stamp based
  - As discussed here, has too many problems
    - **Starvation**
    - Non-recoverable
    - Cascading rollbacks required
    - Remember: We can always put more and more restrictions on what the transactions can do to ensure these things
    - The goal is to find the minimal set of restrictions to as to not hinder concurrency



# Other CC Schemes

---

- **Optimistic** concurrency control
  - Also called validation-based
  - Intuition
    - Let the transactions execute as they wish
    - At the very end when they are about to commit, check if there might be any problems/conflicts etc
      - If no, let it commit
      - If yes, abort and restart
  - Optimistic: The hope is that there won't be too many problems/aborts
- Rarely used any more

# More Locking Issues: Deadlocks

- No action proceeds:

Deadlock

- T1 waits for T2 to unlock A
- T2 waits for T1 to unlock B
- 2PL does not prevent deadlock
  - Strict doesn't either

Rollback transactions

Can be costly

T1	T2
lock-X(B) read(B) $B \leftarrow B-50$ write(B)	lock-S(A) read(A) lock-S(B)
lock-X(A)	

# Preventing deadlocks

---

- Solution 1: A transaction must acquire all locks before it begins
  - Not acceptable in most cases
- Solution 2: A transaction must acquire locks in a particular order over the data items
  - Also called *graph-based protocols*
- Solution 3: Use time-stamps; say T1 is older than T2
  - *wait-die scheme*: T1 will wait for T2. T2 will not wait for T1; instead it will abort and restart
  - *wound-wait scheme*: T1 will *wound* T2 (force it to abort) if it needs a lock that T2 currently has; T2 will wait for T1.
- Solution 4: Timeout based
  - Transaction waits a certain time for a lock; aborts if it doesn't get it by then

# Deadlock Prevention

- Assign priorities based on timestamps.
- **small timestamp --- high priority**




T1: S(A), R(A),  
T2: X(B), W(B)      S(B) *wait!*  
T3: S(C), R(C)      X(C) *wait!*  
T4: X(A) *abort!*  
X(B) *abort!*

- Assume  $T_i$  wants a lock that  $T_j$  holds.
- **Wait-Die**: If  $T_i$  has **higher** priority,  $T_i$  waits for  $T_j$ ; otherwise  $T_i$  aborts (commits suicide)
- **Lower priority never waits for higher priority.**

# Deadlock Prevention

- Assign priorities based on timestamps.
- **small** timestamp --- **high** priority




- Assume  $T_i$  wants a lock that  $T_j$  holds.
- Wound-wait: If  $T_i$  has higher priority,  $T_j$  aborts; otherwise  $T_i$  waits
- **Higher priority never waits for lower priority.**

# Deadlock detection and recovery

- Instead of trying to prevent deadlocks, let them happen and deal with them if they happen
- How do you detect a deadlock?
  - Wait-for graph
  - Directed edge from  $T_i$  to  $T_j$ 
    - $T_i$  waiting for  $T_j$

T1	T2	T3	T4
S(V)	X(V) S(W)	X(Z) S(V)	X(W)



Suppose T4 requests lock-S(Z)

# Dealing with Deadlocks

---

- Deadlock detected, now what ?
  - Will need to abort some transaction
  - Prefer to **abort** the one with the minimum work done so far
  - Possibility of starvation
    - If a transaction is aborted too many times, it may be given priority in continuing