

ER Model and Relational Model

Wen-Chih Peng

Dept. of Computer Science

National Chiao Tung University

Introduction

- The **relational model** was first introduced by Ted Codd of IBM Research in 1970,
- The model uses the concept of a **mathematical relation** - which looks like ***a table of values*** –
as its basic building block.
- The **SQL Query Language** was developed by IBM in the 1970's

Why study the relational model ?

- Dominant data model
- Used in almost DB systems
- Supports simple, powerful **querying** of data

Data Modeling

- Goals:
 - Conceptual representation of the data
 - “Reality” meets “bits and bytes”
 - Must make sense, and be usable by other people
- We will study:
 - Entity-relationship Model
 - Relational Model
 - Note the difference !!

Motivation

- Suppose that you are DBA for online banking web site.
- You are asked to create a database that monitors:
 - customers
 - accounts
 - loans
 - branches
 - transactions

Database Design Steps

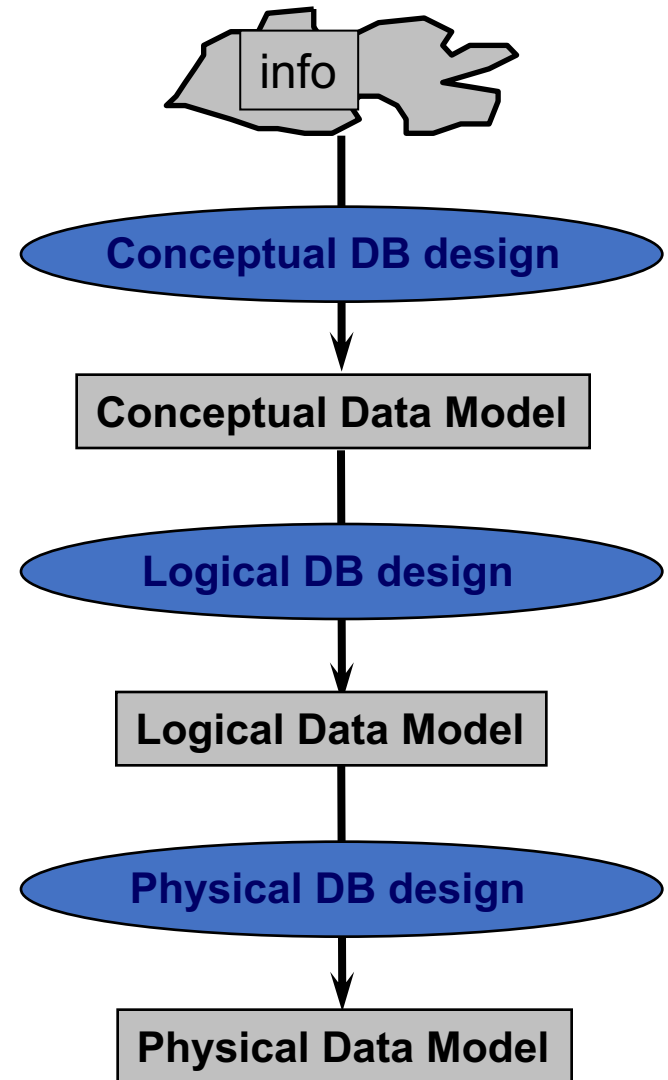
Entity-relationship Model

Typically used for conceptual database design

Three Levels of Modeling

Relational Model

Typically used for logical database design



Entity-Relationship Model

- Two key concepts

- Entities:

- An object that *exists* and is *distinguishable* from other objects
 - Examples: Bob Smith, BofA, CMSC424
 - Have attributes (people have names and addresses)
 - Form entity sets with other entities of the same type that share the same properties
 - Set of all people, set of all classes
 - Entity sets may overlap
 - Customers and Employees

Entity-Relationship Model

- Two key concepts
 - Relationships:
 - Relate 2 or more entities
 - E.g. Bob Smith has account at College Park Branch
 - Form relationship sets with other relationships of the same type that share the same properties
 - Customers have accounts at Branches
 - Can have attributes:
 - has account at may have an attribute *start-date*
 - Can involve more than 2 entities
 - Employee *works at* Branch *at* Job

Relationship Sets

- A **relationship** is an association among several entities

Example:

<u>Hayes</u>	<u>depositor</u>	<u>A-102</u>
<i>customer</i>	<i>relationship</i>	<i>account</i>

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

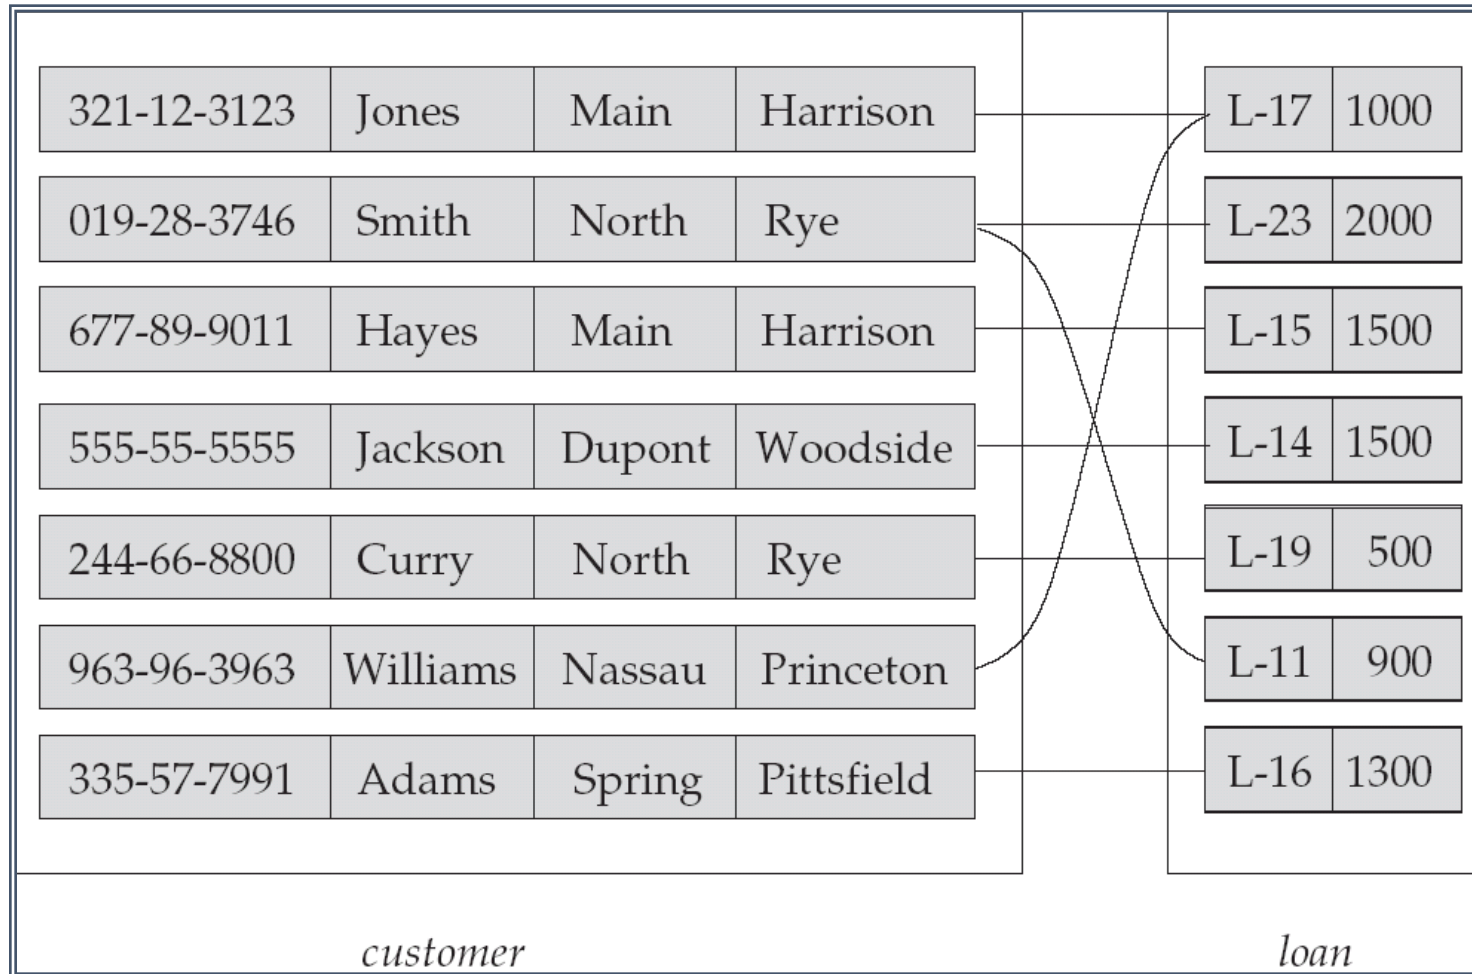
$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where (e_1, e_2, \dots, e_n) is a relationship

- Example:

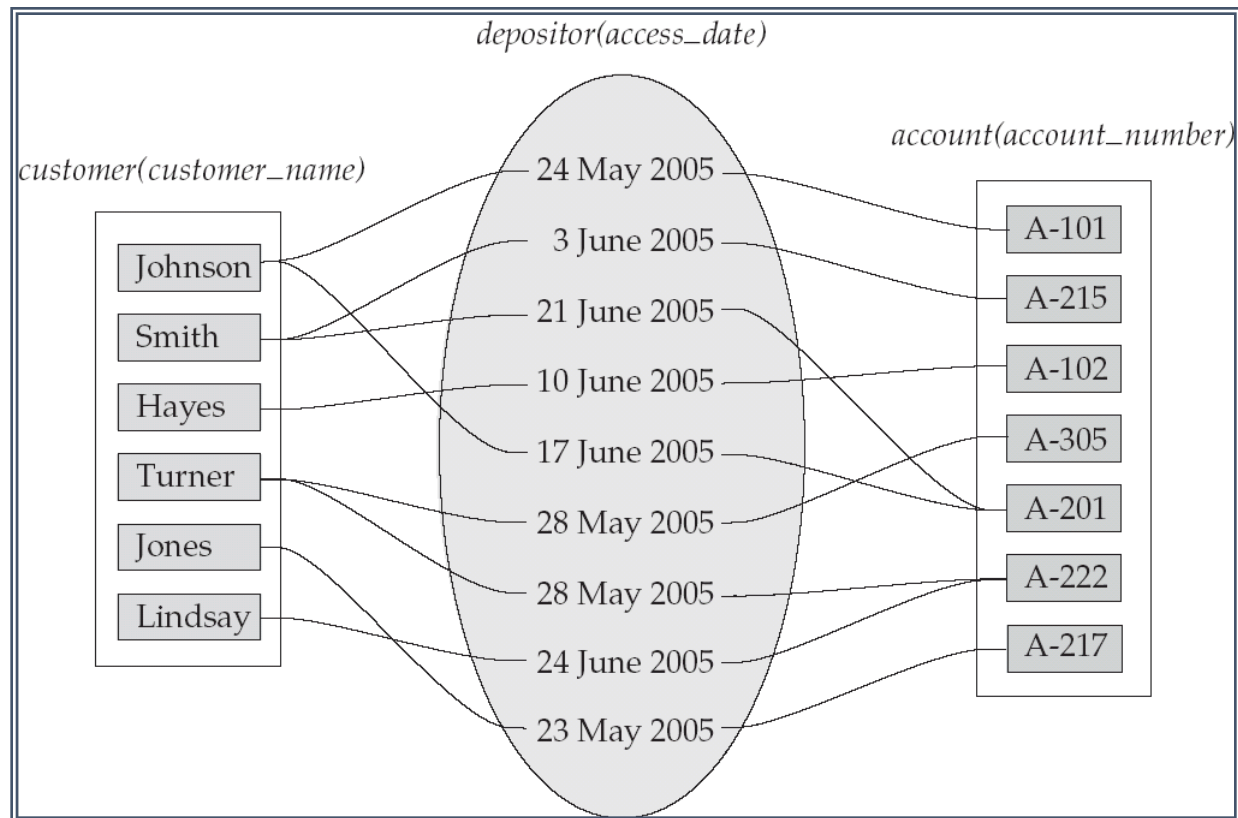
$$(Hayes, A-102) \in depositor$$

Relationship Set *borrower*

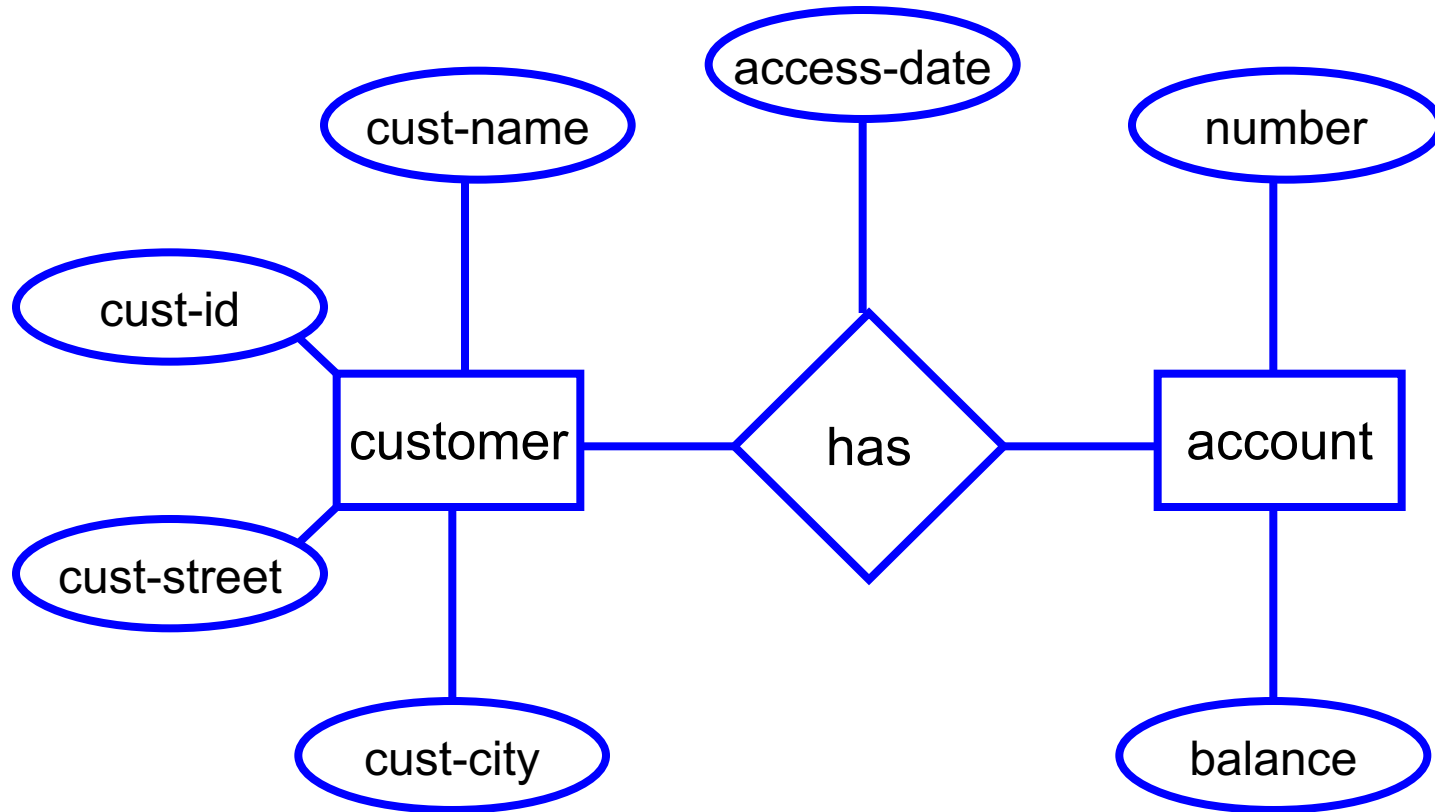


Relationship Sets (Cont.)

- An **attribute** can also be property of a relationship set.
- For instance, the *depositor* relationship set between entity sets *customer* and *account* may have the attribute *access-date*



ER Diagram: Starting Example



- Rectangles: entity sets
- Diamonds: relationship sets
- Ellipses: attributes

Rest of the class

- Details of the ER Model
 - How to represent various types of constraints/semantic information
- Design issues
 - A detailed example

Next: Relationship Cardinalities

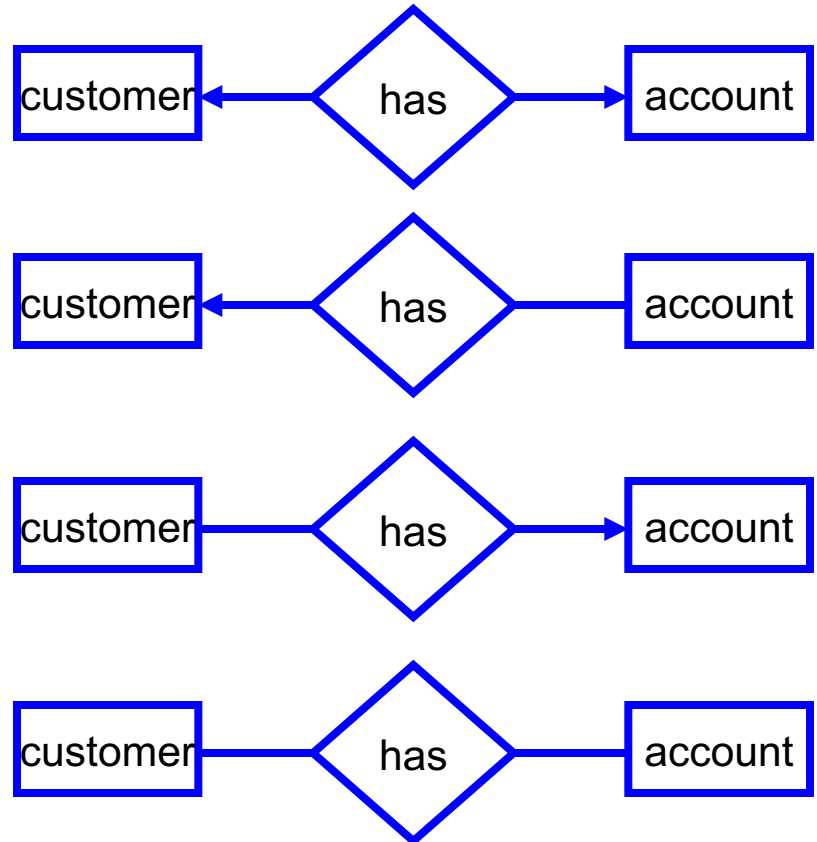
- We may know:
 - One customer can only open one account
 - OR*
 - One customer can open multiple accounts
- Representing this is important
- Why ?
 - Better manipulation of data
 - If former, can store the account info in the customer table
 - Can enforce such a constraint
 - Application logic will have to do it; NOT GOOD
 - Remember: If not represented in conceptual model, the domain knowledge may be lost

Mapping Cardinalities

- Express the number of entities to which another entity can be associated via a relationship set
- Most useful in describing binary relationship sets

Mapping Cardinalities

- One-to-One
- One-to-Many
- Many-to-One
- Many-to-Many

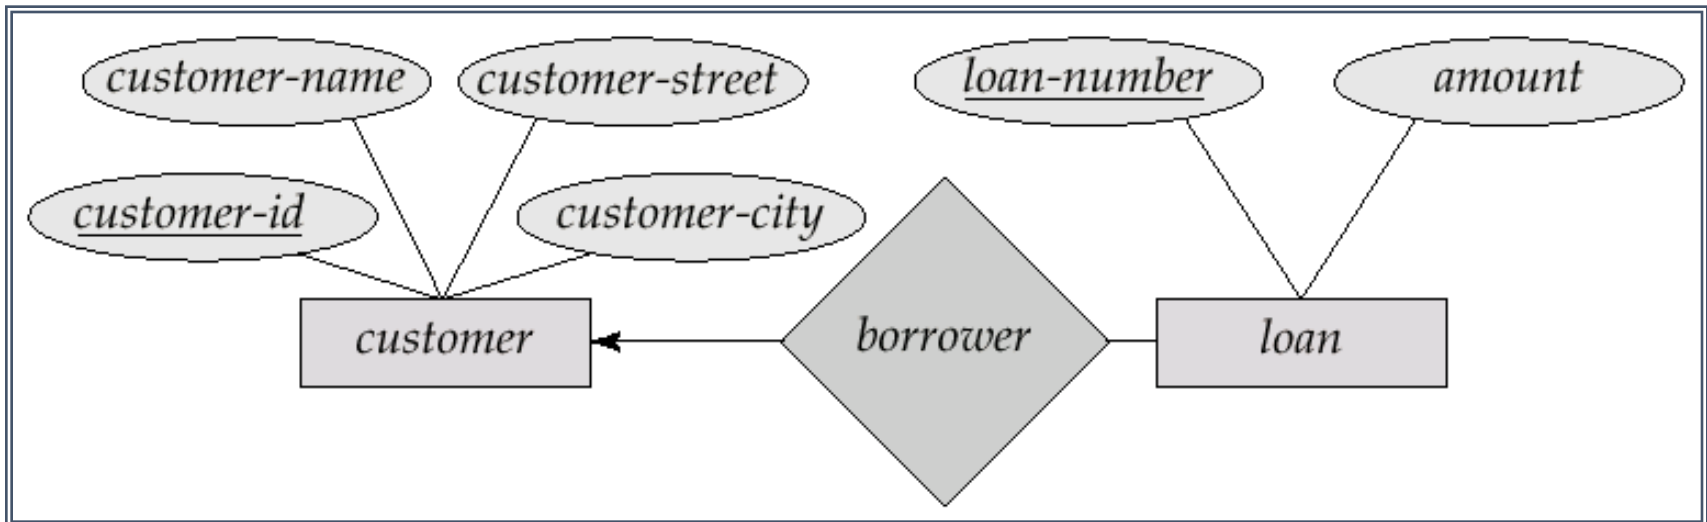


Mapping Cardinalities

- Express the number of entities to which another entity can be associated via a relationship set
- Most useful in describing binary relationship sets
- N-ary relationships ?
 - More complicated
 - Details in the book

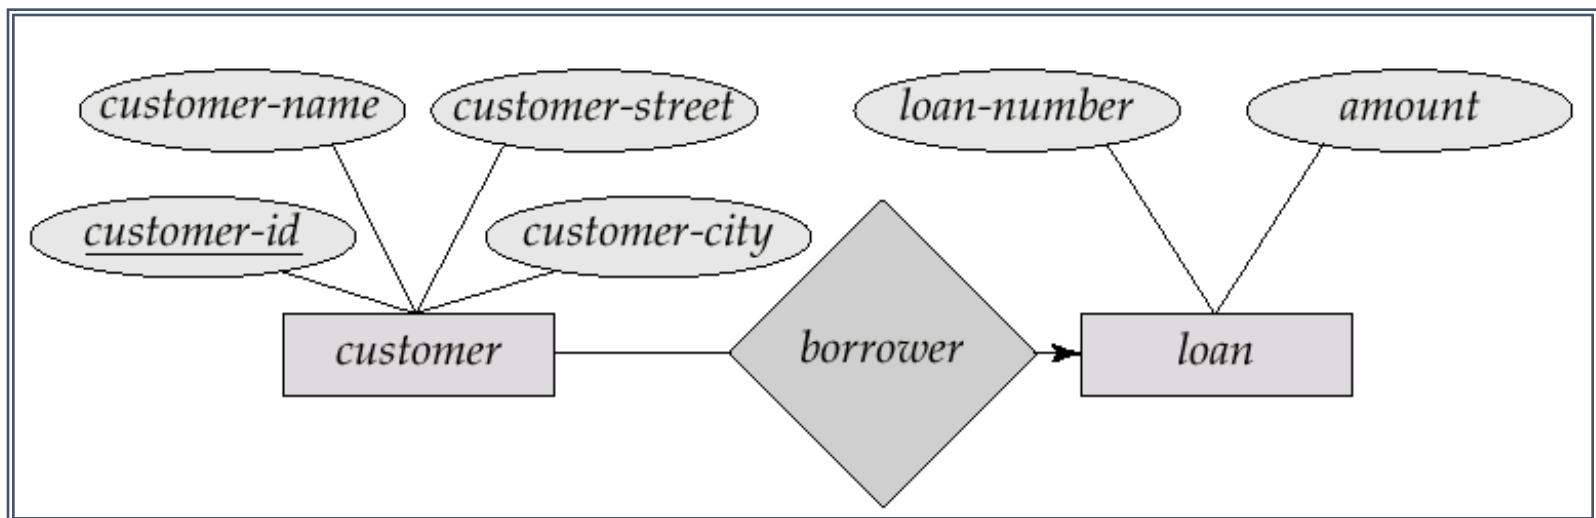
One-To-Many Relationship

- In the one-to-many relationship a loan is associated with at most one customer via *borrower*, a customer is associated with several (including 0) loans via *borrower*



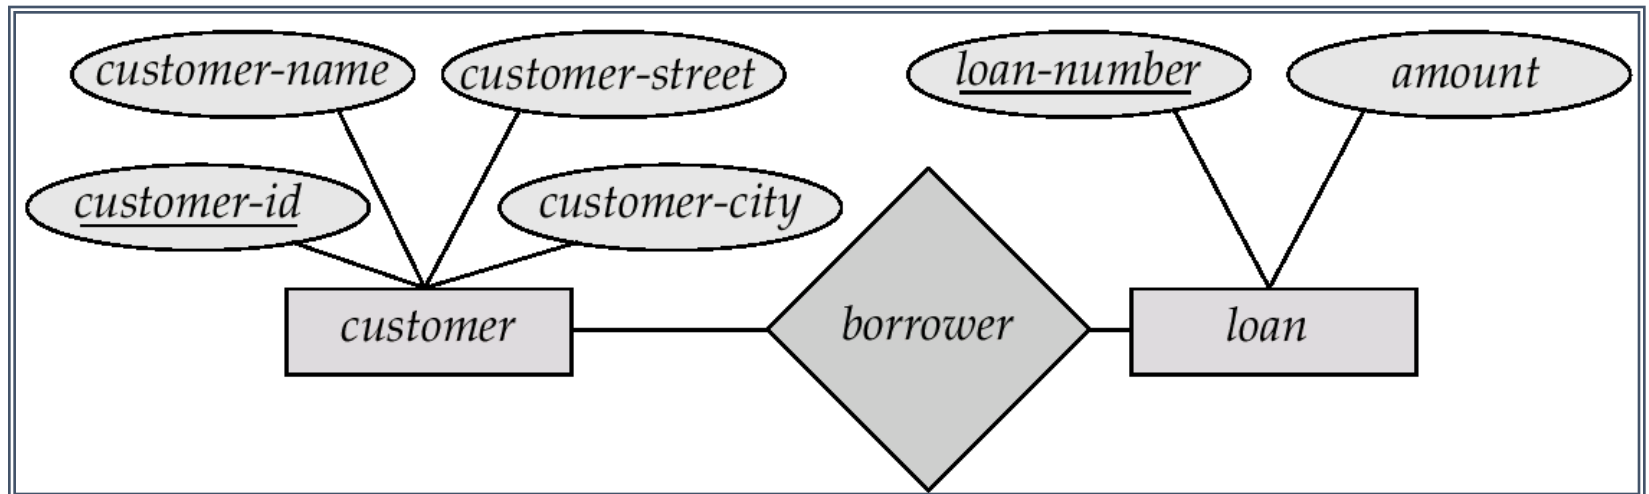
Many-To-One Relationships

- In a many-to-one relationship a loan is associated with several customers via *borrower*, a customer is associated with at most one loan via *borrower*



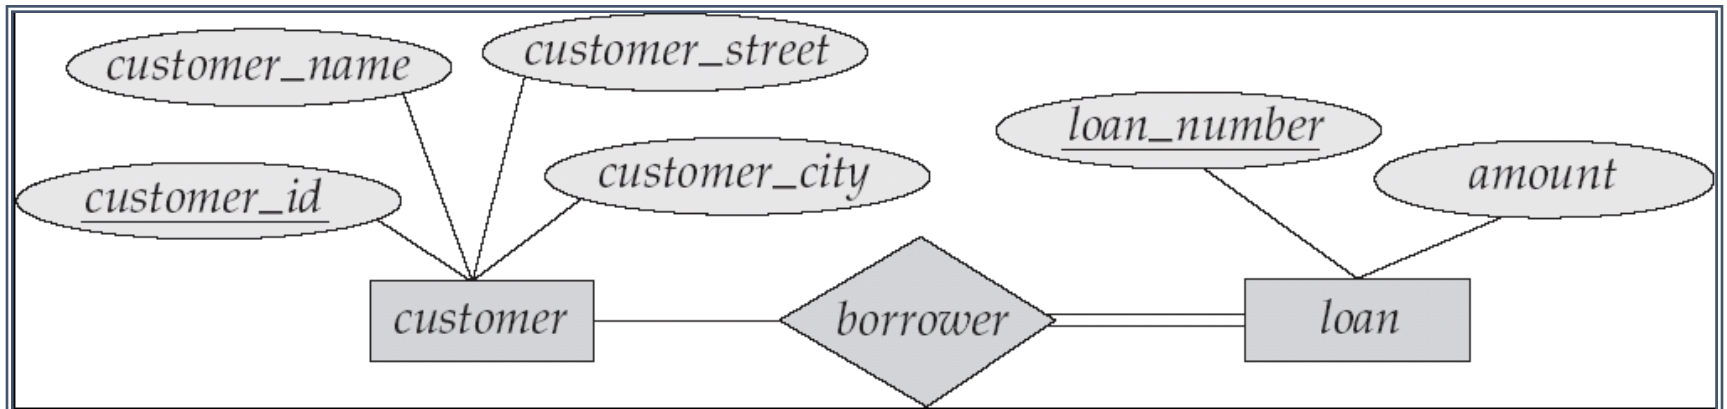
Many-To-Many Relationship

- A customer is associated with several (possibly 0) loans via borrower
- A loan is associated with several (possibly 0) customers via borrower




Participation of an Entity Set in a Relationship Set

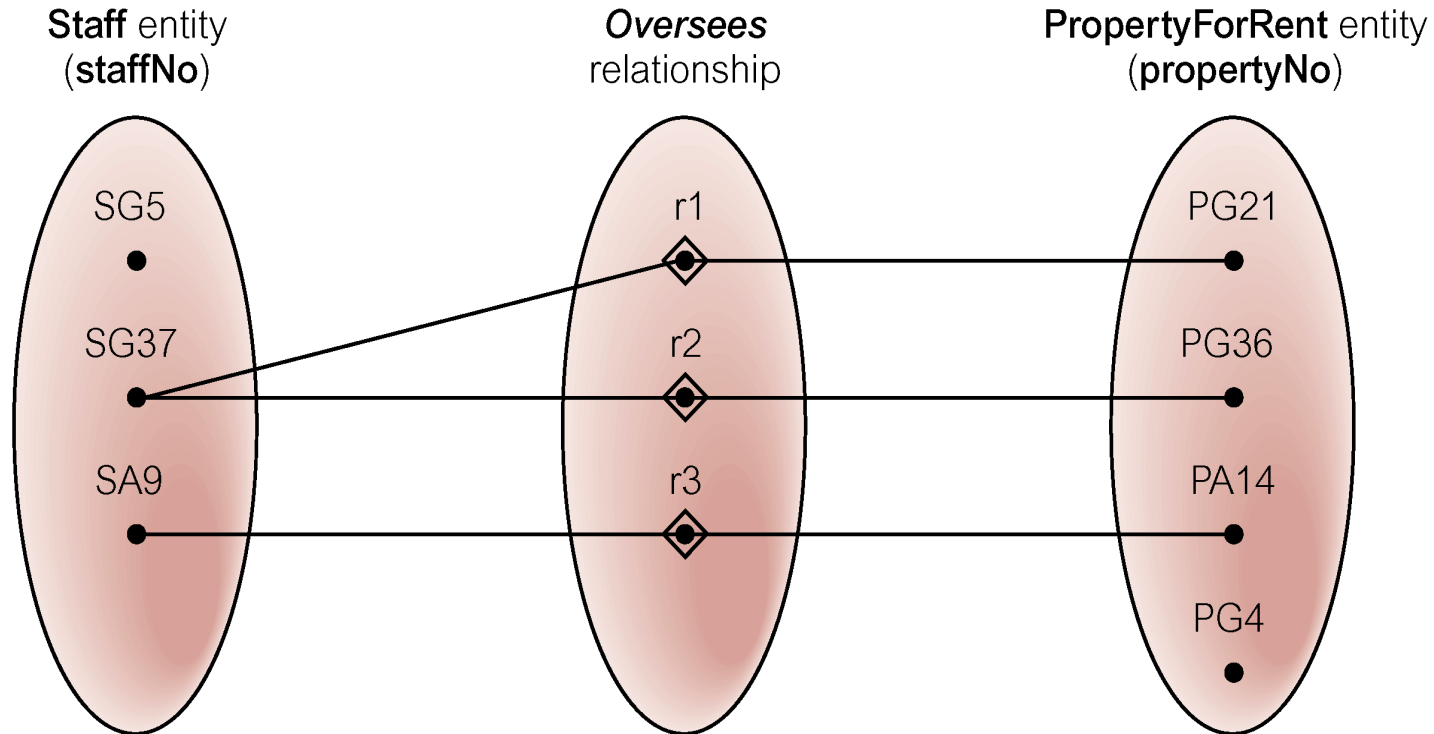
- Total participation (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set
 - E.g. participation of loan in borrower is total
 - ▶ every loan must have a customer associated to it via borrower
- Partial participation: some entities may not participate in any relationship in the relationship set
 - Example: participation of customer in borrower is partial



Semantic net of Staff *Manages* Branch relationship type

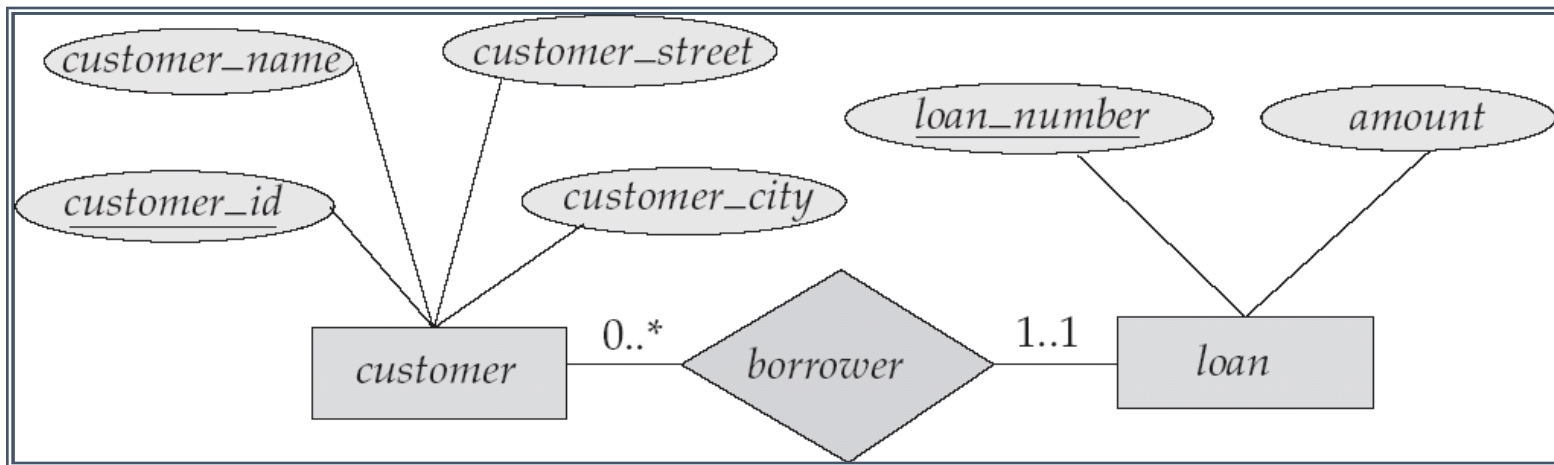


Semantic net of Staff *Oversees* PropertyForRent relationship type



Alternative Notation for Cardinality Limits


- Cardinality limits can also express participation constraints



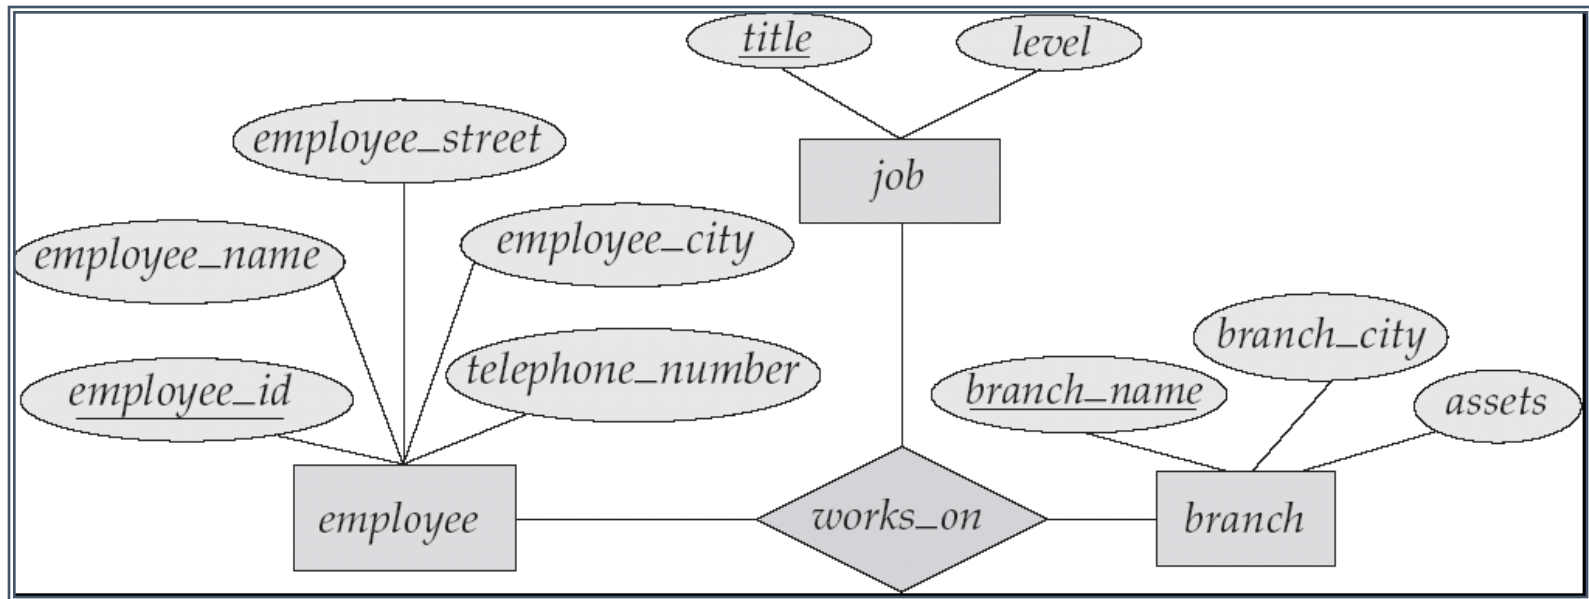
Customer: Loan => *:1

Many-To-One Relationships

- In a many-to-one relationship a loan is associated with several customers via *borrower*, a customer is associated with at most one loan via *borrower*



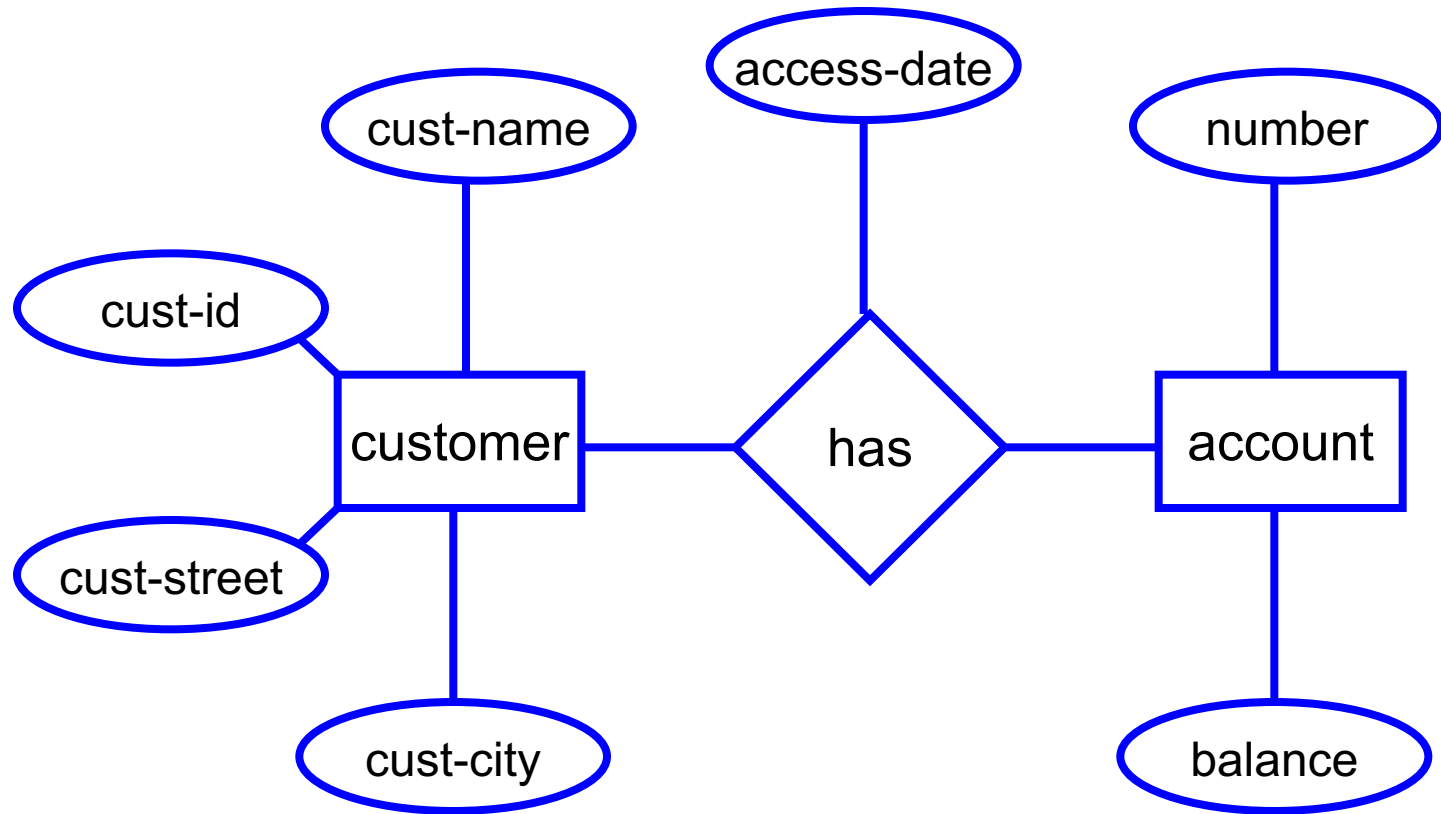
E-R Diagram with a Ternary Relationship



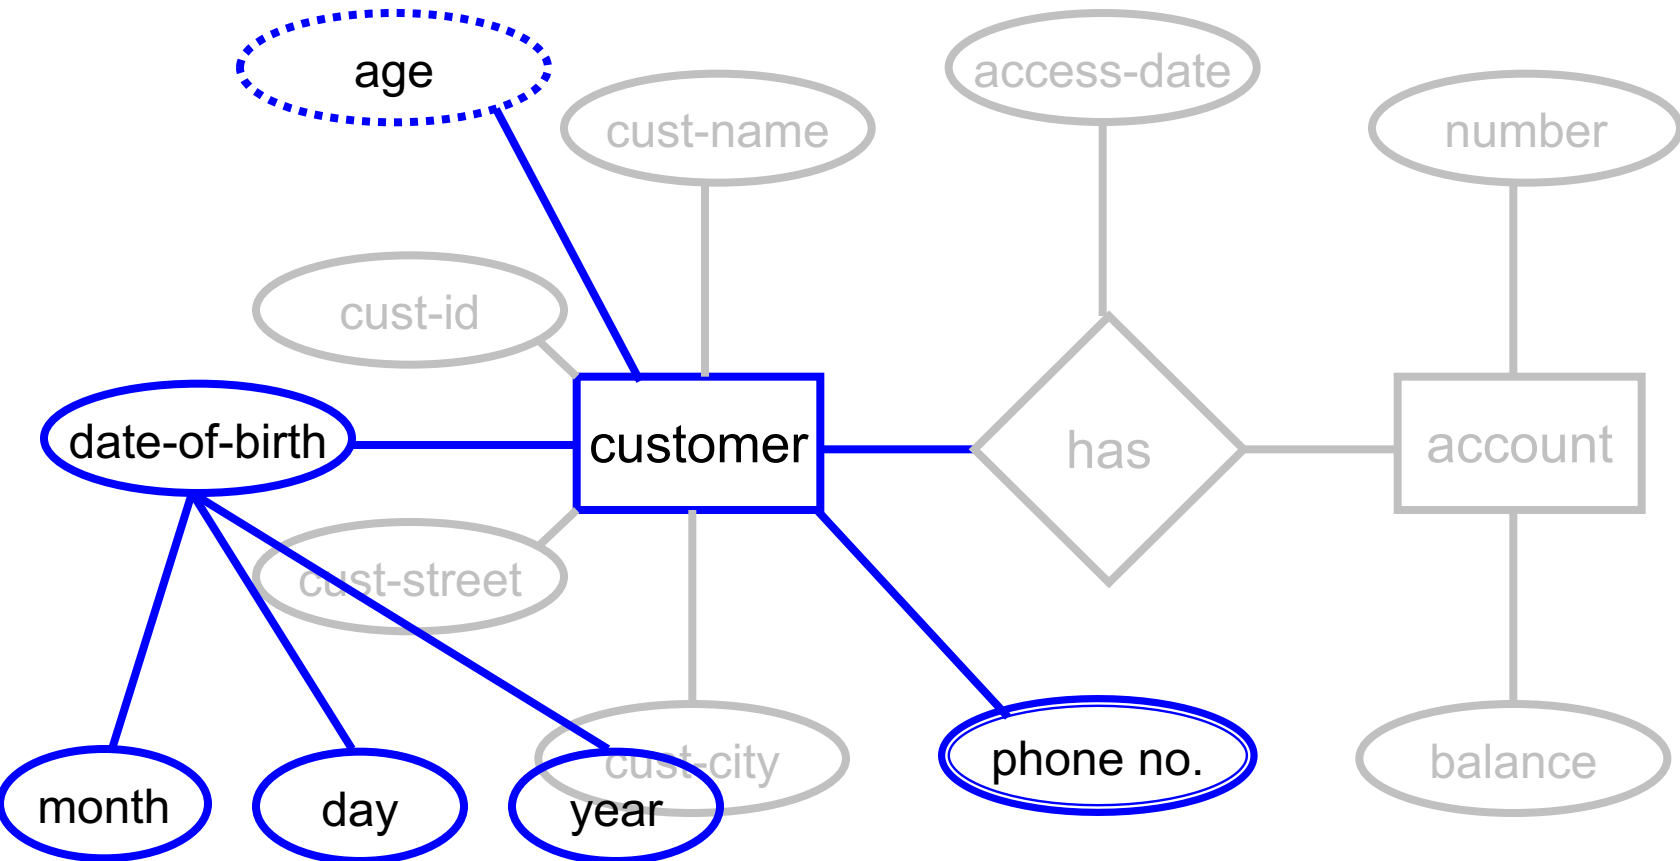
Next: Types of Attributes

- Simple vs Composite
 - Single value per attribute ?
- Single-valued vs Multi-valued
 - E.g. Phone numbers are multi-valued
- Derived
 - If date-of-birth is present, age can be derived
 - Can help in avoiding redundancy, enforcing constraints etc

Types of Attributes



Types of Attributes

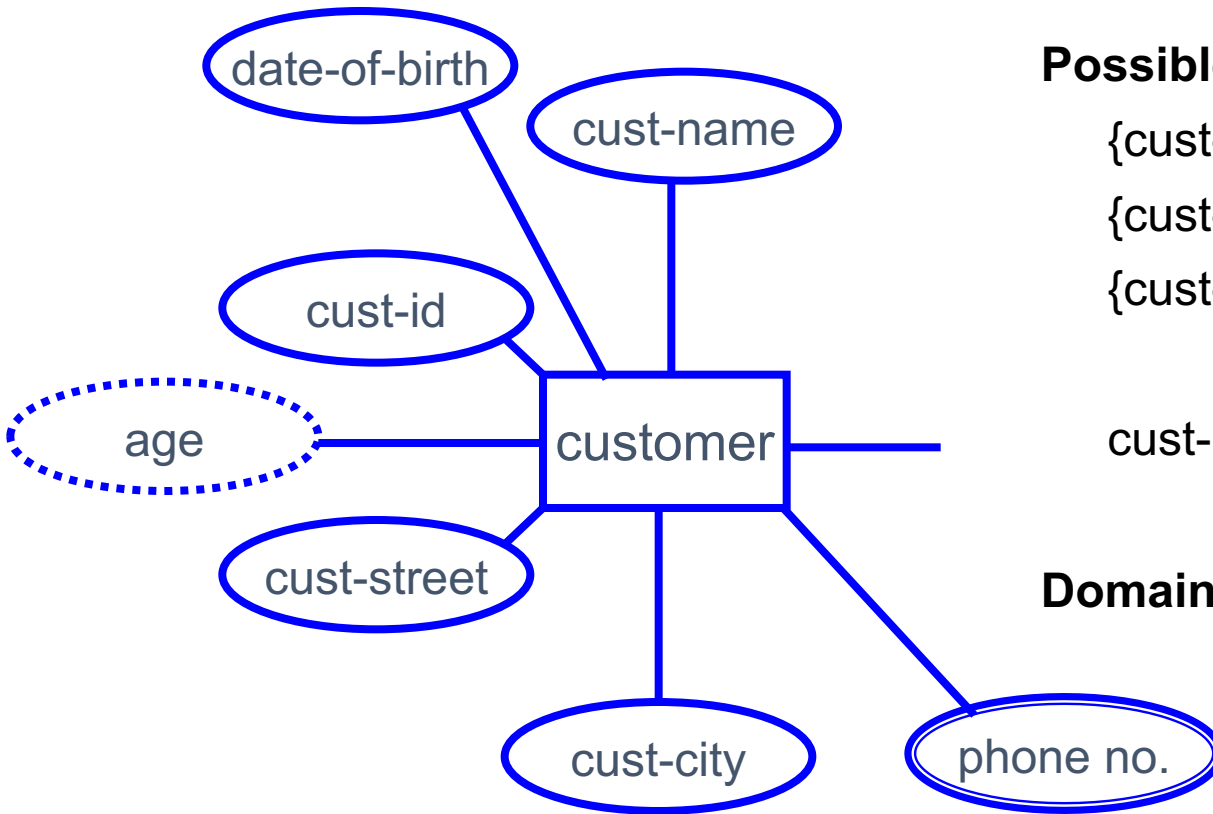


Composite Attribute

Next: Keys

- Key = set of attributes that uniquely identifies an entity or a relationship

Entity Keys



Possible Keys:

{cust-id}

{cust-name, cust-city, cust-street}

{cust-id, age}

cust-name ?? Probably not.

Domain knowledge dependent !!

Entity Keys

- *Superkey*
 - any attribute set that can distinguish entities
- *Candidate key*
 - a **minimal** superkey
 - Can't remove any attribute and preserve key-ness
 - {cust-id, age} not a candidate key since we could have {cust-id} by removing "age" of {cust-id, age}
 - {cust-name, cust-city, cust-street} is
 - assuming cust-name is not unique
- *Primary key*
 - Candidate key chosen as the key by DBA
 - Underlined in the ER Diagram

Entity Sets *customer* and *loan*

id

Name

street

city

Loan ID Value

321-12-3123	Jones	Main	Harrison
-------------	-------	------	----------

019-28-3746	Smith	North	Rye
-------------	-------	-------	-----

677-89-9011	Hayes	Main	Harrison
-------------	-------	------	----------

555-55-5555	Jackson	Dupont	Woodside
-------------	---------	--------	----------

244-66-8800	Curry	North	Rye
-------------	-------	-------	-----

963-96-3963	Williams	Nassau	Princeton
-------------	----------	--------	-----------

335-57-7991	Adams	Spring	Pittsfield
-------------	-------	--------	------------

L-17	1000
------	------

L-23	2000
------	------

L-15	1500
------	------

L-14	1500
------	------

L-19	500
------	-----

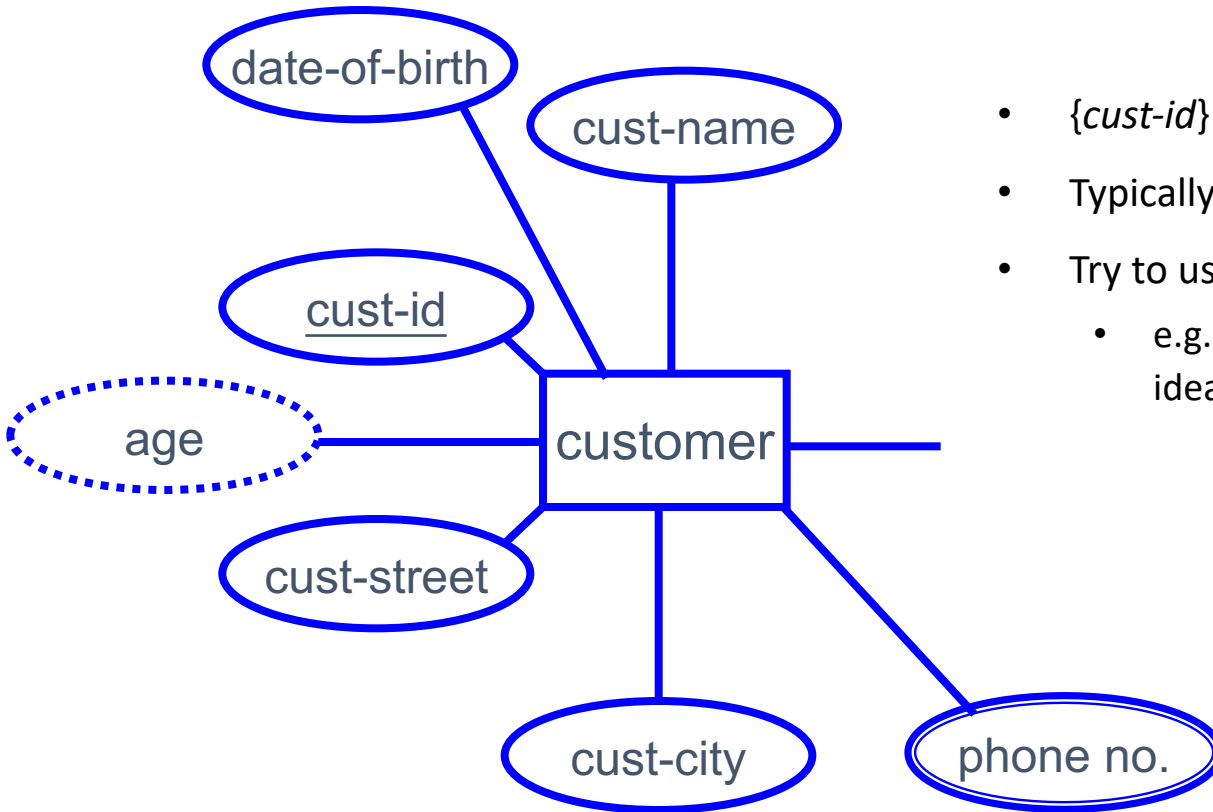
L-11	900
------	-----

L-16	1300
------	------

customer

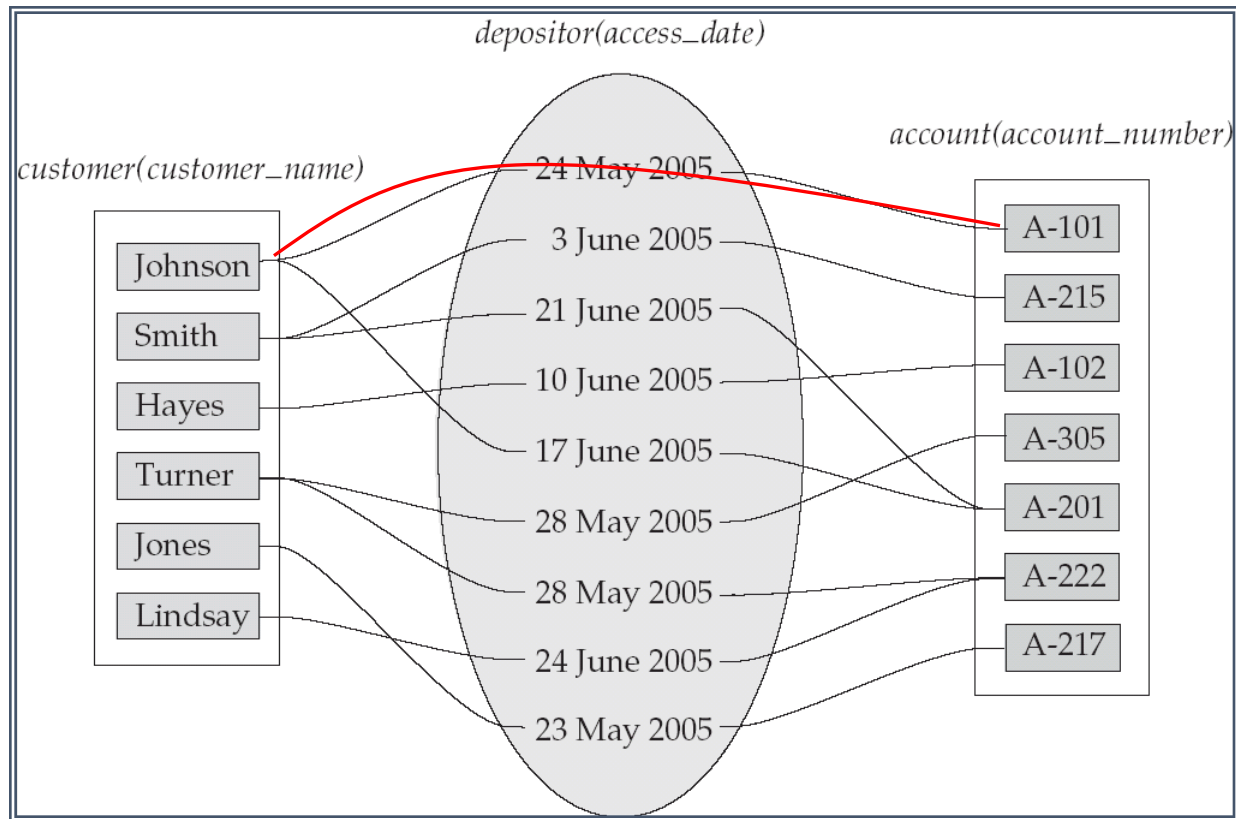
loan

Entity Keys




- {*cust-id*} is a natural primary key
- Typically, SSN forms a good primary key
- Try to use a candidate key that rarely changes
 - e.g. something involving address not a great idea

Relationship Set Keys



Relationship Set Keys

- What attributes are needed to represent a relationship completely and uniquely ?
 - Union of primary keys of the entities involved, and relationship attributes



- $\{cust-id, create-date, account\ number\}$ describes a relationship completely

Example


The table below describes the relationship between customers and accounts

Cust-id	Create-date	Account Number
J121454991	2016/03/03	A101
T423323222	2008/05/30	A103
X123234343	2008/01/01	A345
K987655423	2016/03/03	A521

We need to have one “relationship set key” to uniquely identify the above Data. Thus, one possible key is to use “Cust-id, Create-date, account number”.


Relationship Set Keys

- Is $\{cust-id, create-date, account\ number\}$ a candidate key ?
 - No. since we can remove create-date.
 - In fact, union of primary keys of associated entities is always a superkey




Relationship Set Keys

- Is {cust-id, account-number} a candidate key ?
 - Depends on the cardinality mapping



Relationship Set Keys

- Is {cust-id, account-number} a candidate key ?
 - Depends



- If one-to-one relationship, either {*cust-id*} or {*account-number*} sufficient
 - Since a given *customer* can only have one *account*, she can only participate in one relationship

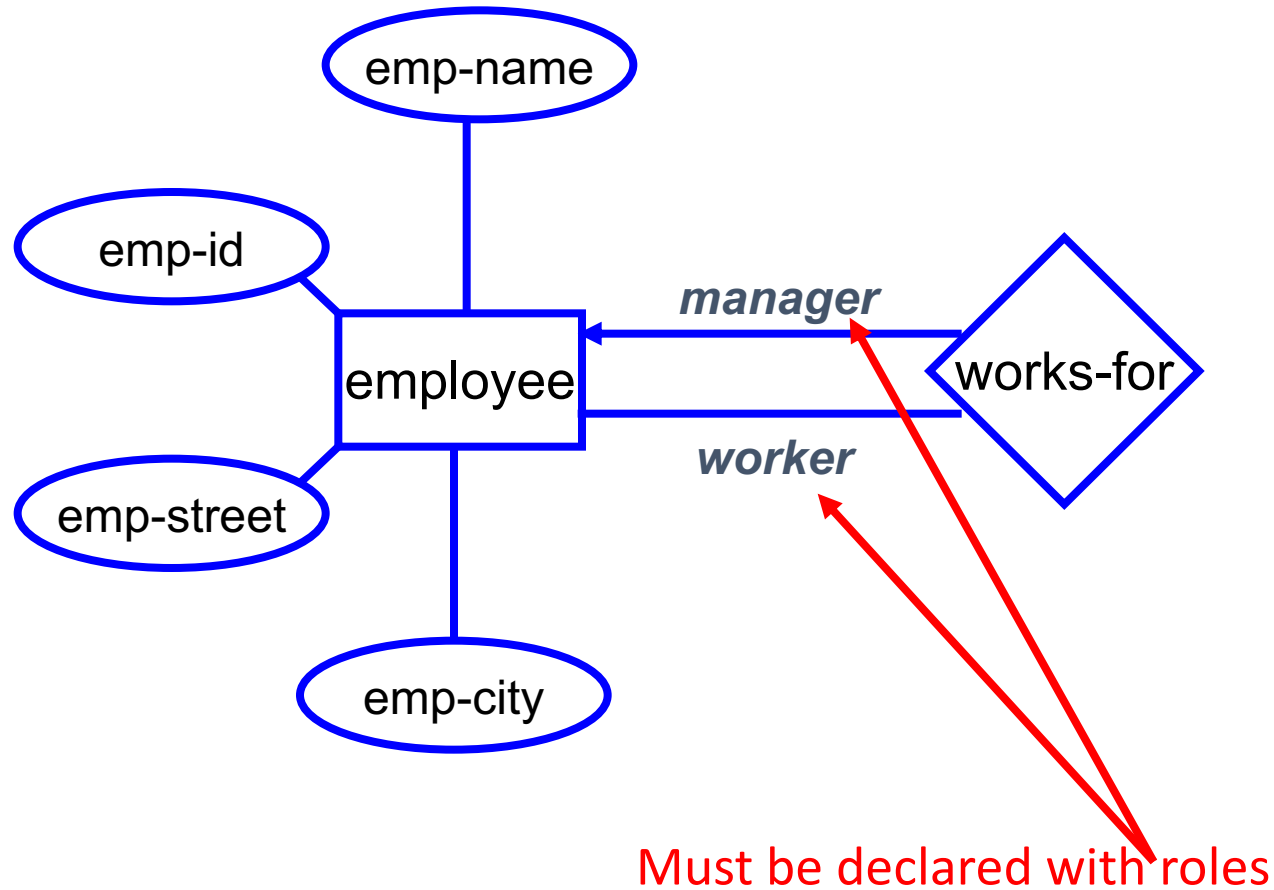
Relationship Set Keys

- General rule for binary relationships
 - one-to-one: primary key of either entity set
 - one-to-many: primary key of the entity set on the **many** side
 - many-to-many: **union of primary keys** of the associate entity sets
- n-ary relationships
 - More complicated rules

Next: Recursive Relationships

- Sometimes a relationship associates an entity set to itself

Recursive Relationships




Next: Weak Entity Sets

- An entity set without enough attributes to have a primary key
- E.g. Transaction Entity
 - Attributes:
 - transaction-number, transaction-date, transaction-amount, transaction-type
 - transaction-number: may not be unique across accounts

Weak Entity Sets

Each account has many transactions (e.g., withdraw) and for
Each account, trans-number is ranged from 1..*




In the above case, different accounts may have the same trans-number

Weak Entity Sets

- A weak entity set must be associated with an identifying or owner entity set
- Account is the owner entity set for Transaction

Weak Entity Sets

Discriminator: A set of attributes that can be used to discriminate in the transaction entity set







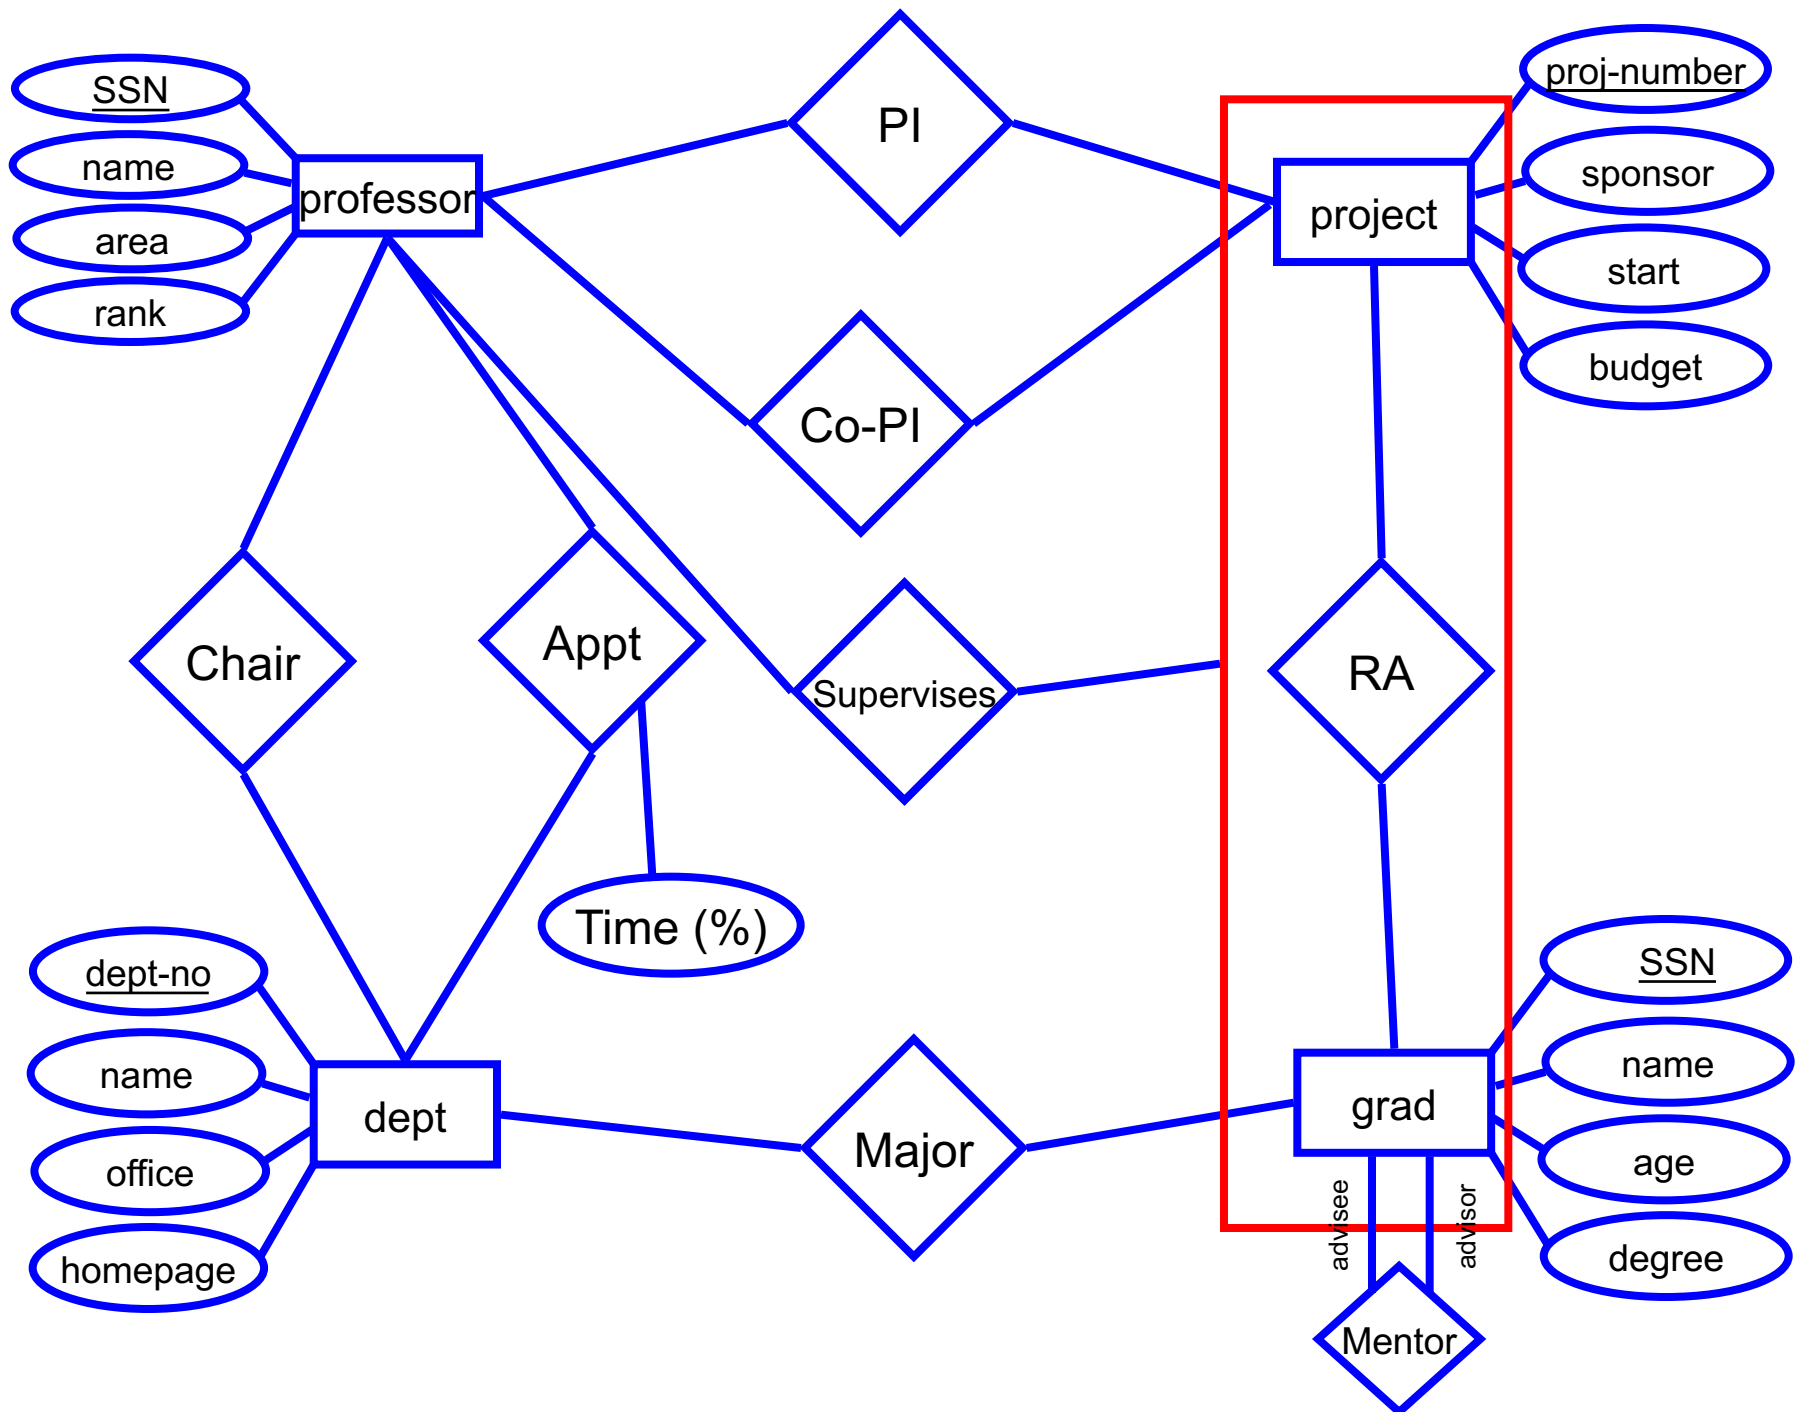
Weak Entity Sets

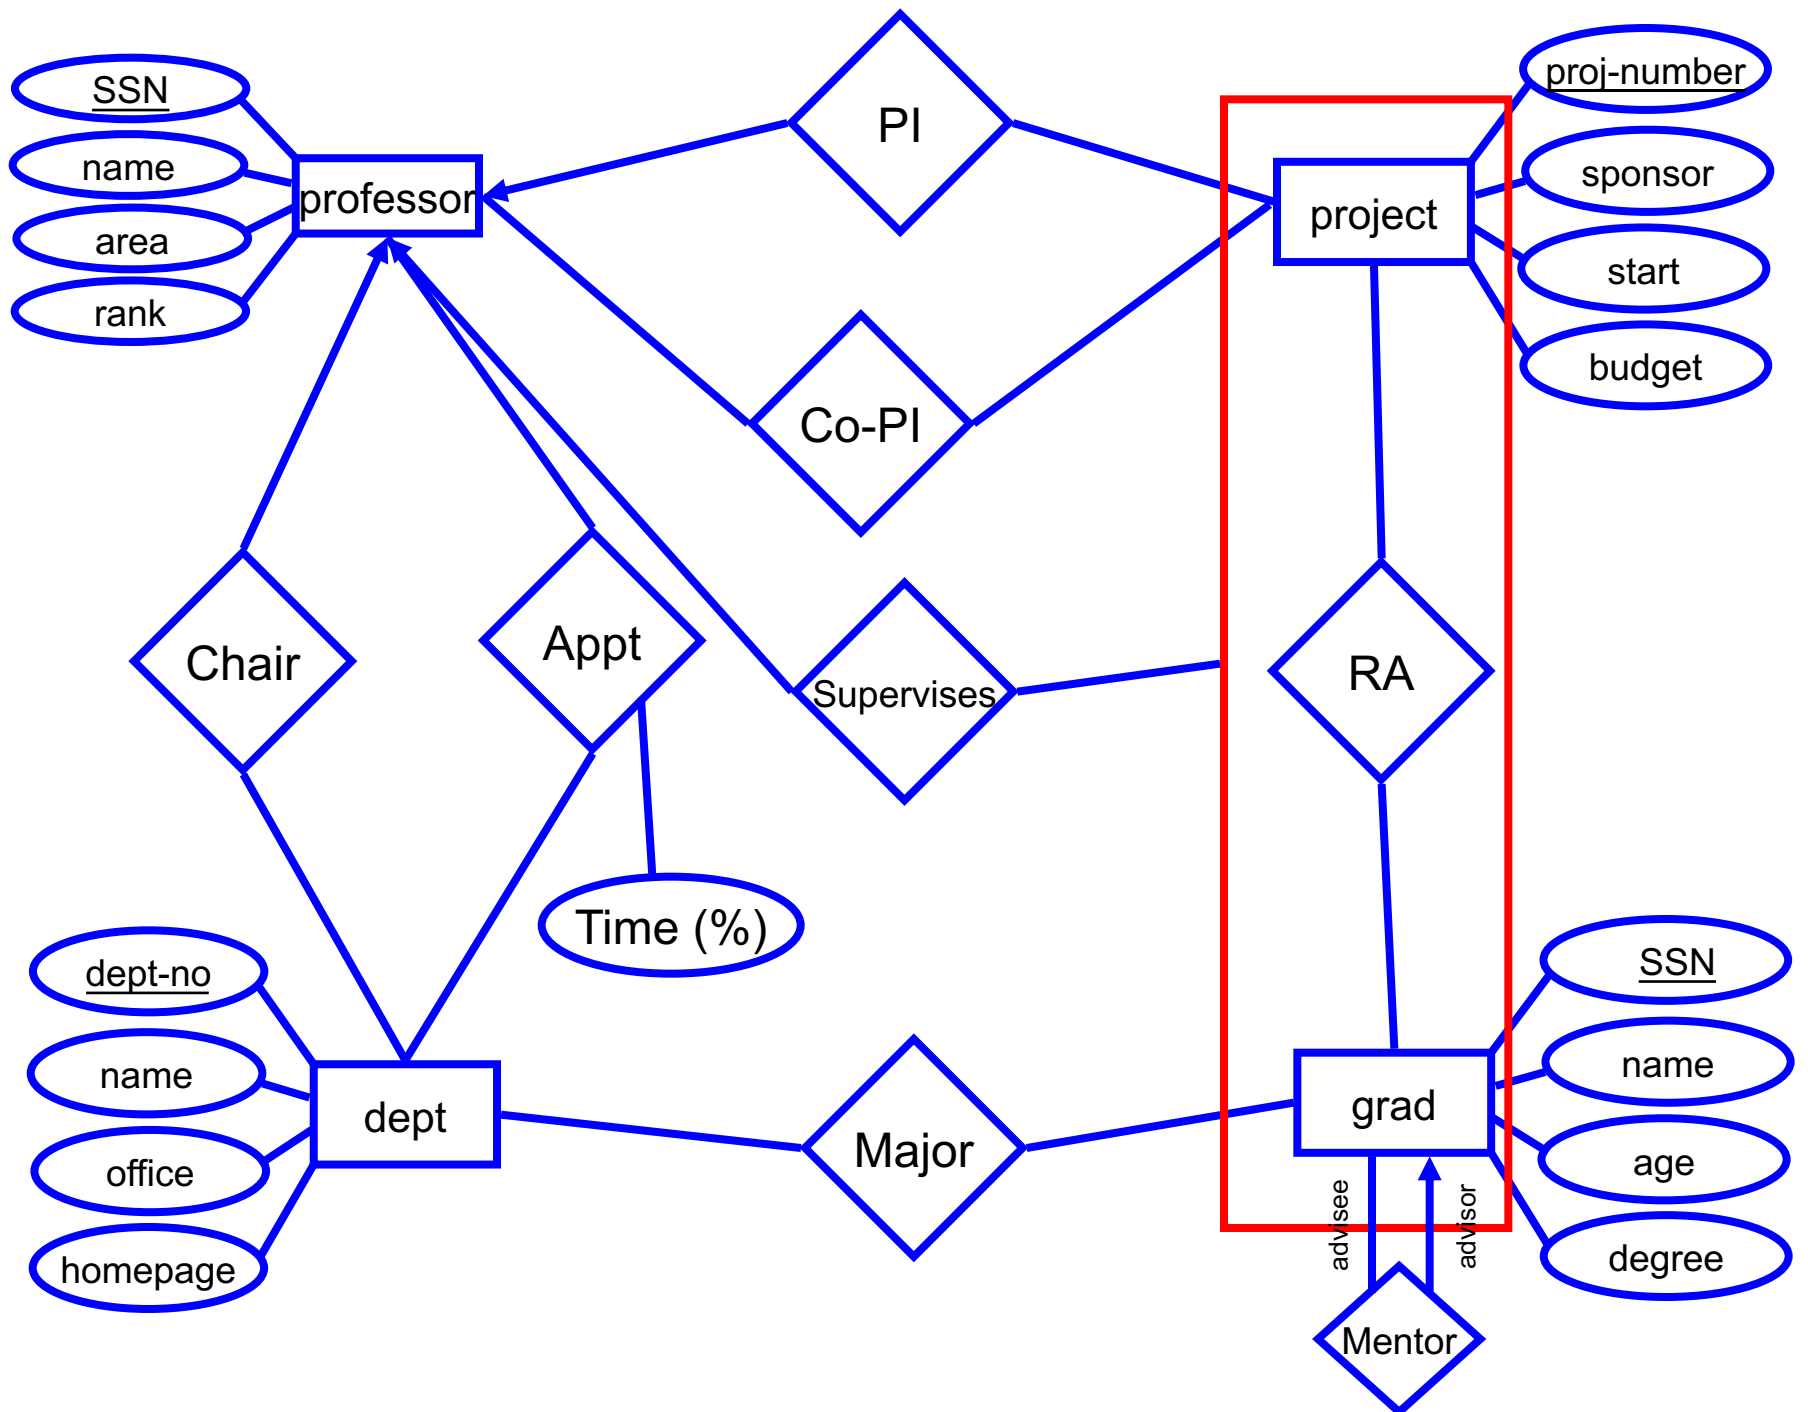
- Primary key:
 - Primary key of the associated strong entity + discriminator attribute set
 - For Transaction:
 - *{account-number, transaction-number}*

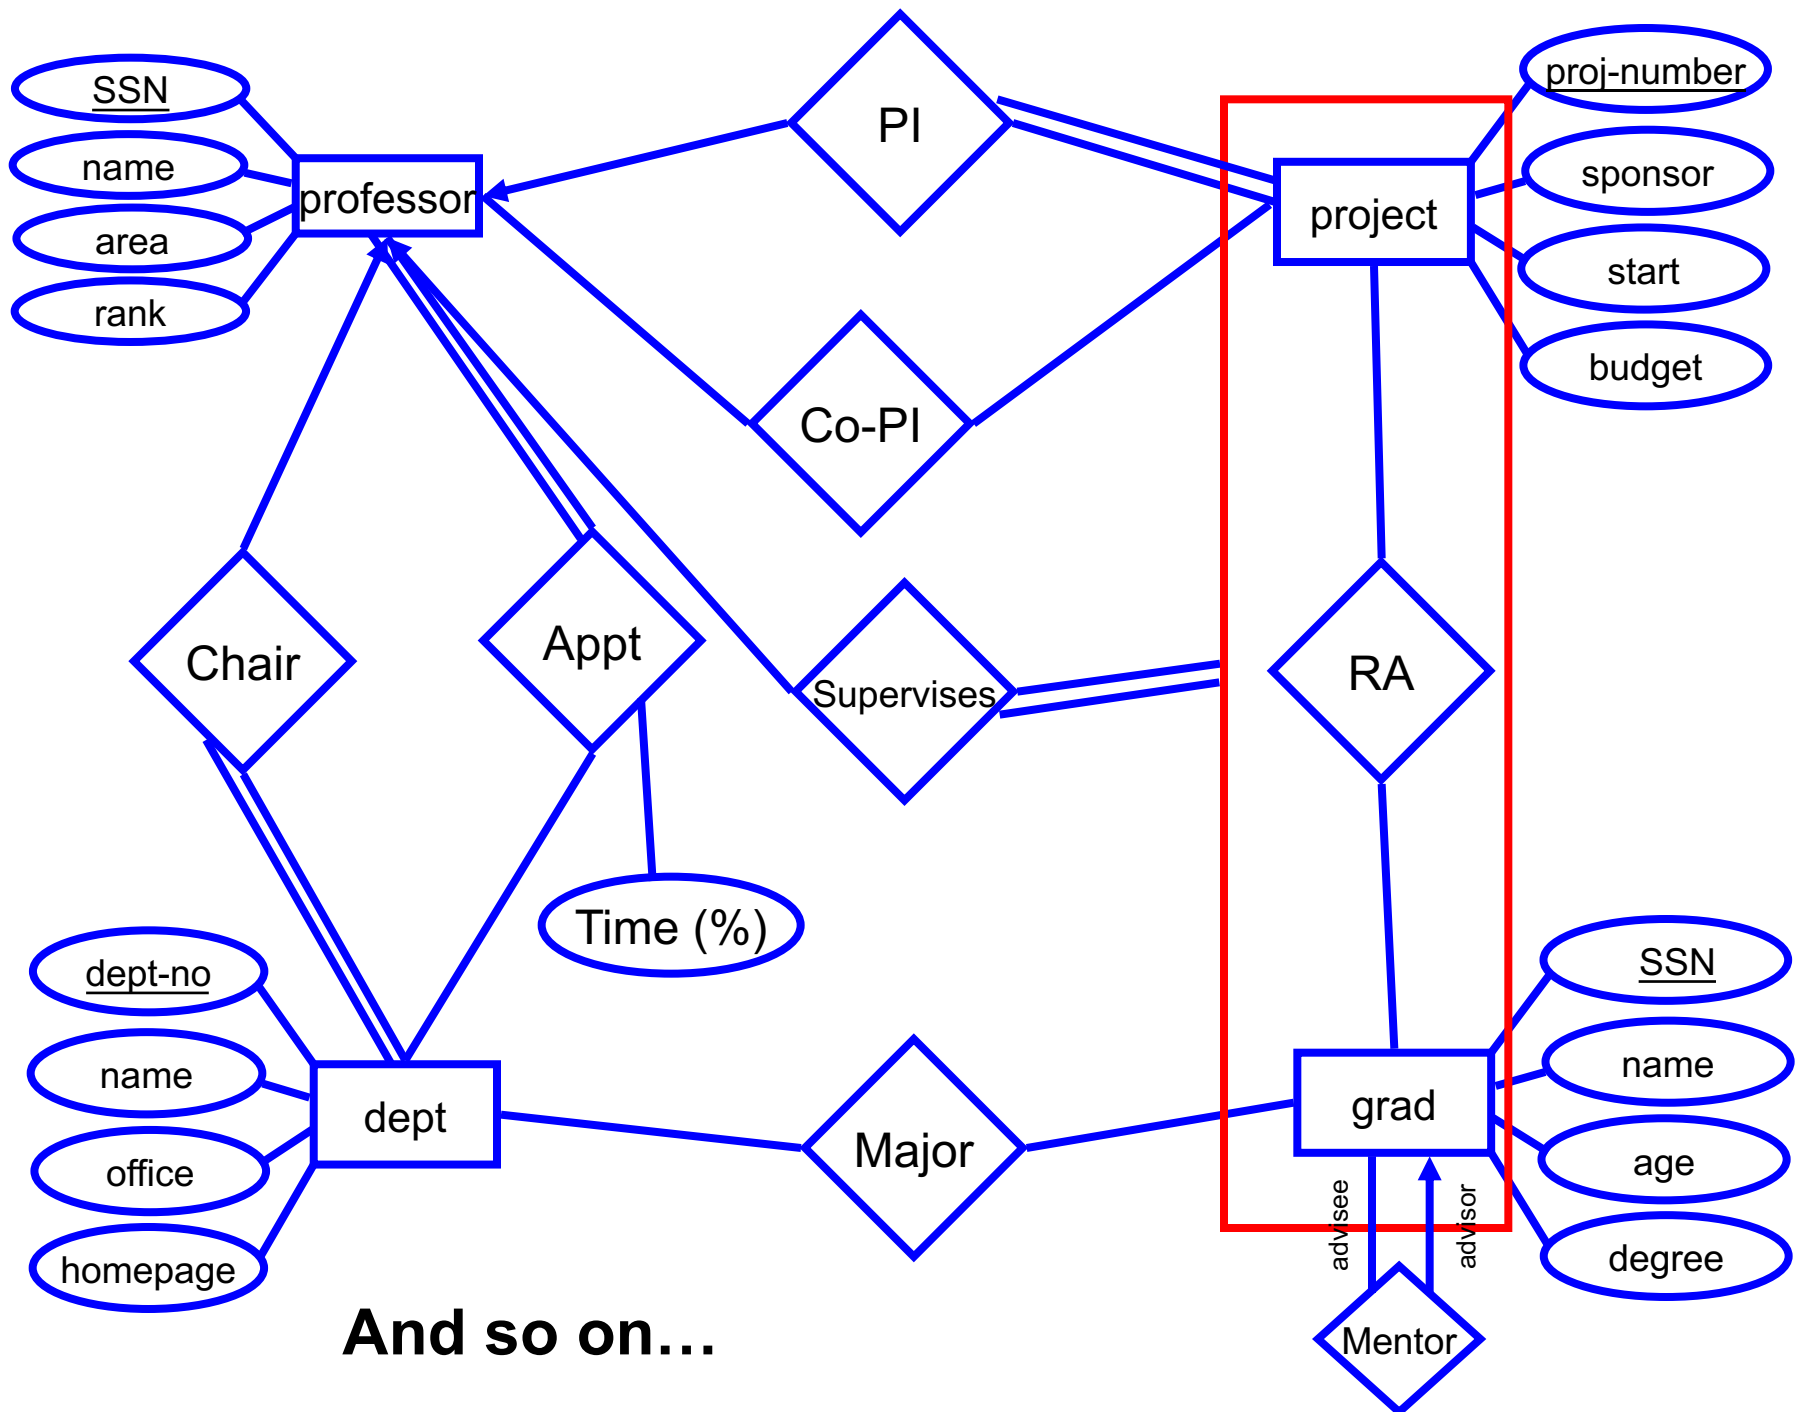
Example Design

- We will model a university database
 - Main entities:
 - Professor
 - Projects
 - Departments
 - Graduate students
 - etc...









Thoughts...

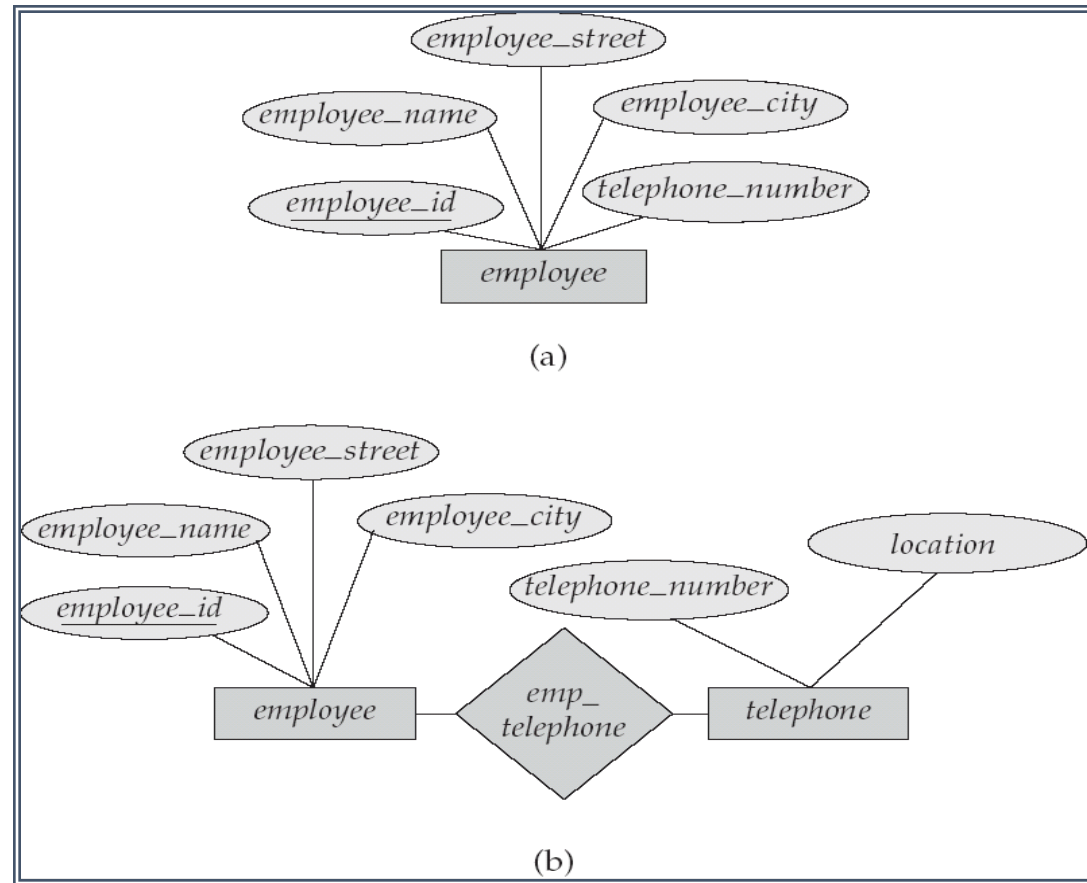
- Nothing about actual data
 - How is it stored ?
- No talk about the query languages
 - How do we access the data ?
- Semantic vs Syntactic Data Models
 - Remember: E/R Model is used for conceptual modeling
 - Many conceptual models have the same properties
- They are much more about representing the knowledge than about database storage/querying

Thoughts...

- Basic design principles
 - Faithful
 - Must make sense
 - Satisfies the application requirements
 - Models the requisite domain knowledge
 - If not modeled, lost afterwards
 - Avoid redundancy
 - Potential for inconsistencies
 - Go for simplicity
- Typically an iterative process that goes back and forth

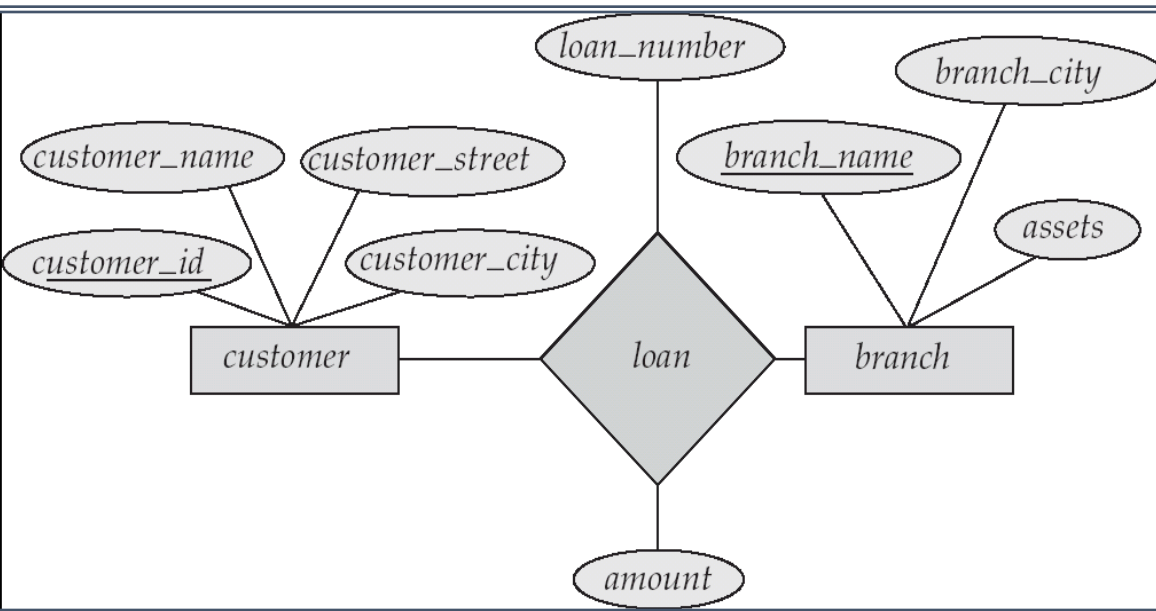
Design Issues

- Entity sets vs attributes
 - Depends on the semantics of the application
 - *If phones have other attributes, it is better to present as an entity*



Design Issues

- Entity sets vs Relationship sets
 - Consider *loan* (Not good to represent many customers hold a loan)



chris	100	300	CS
Tsao	100	300	CS
John	100	300	CS

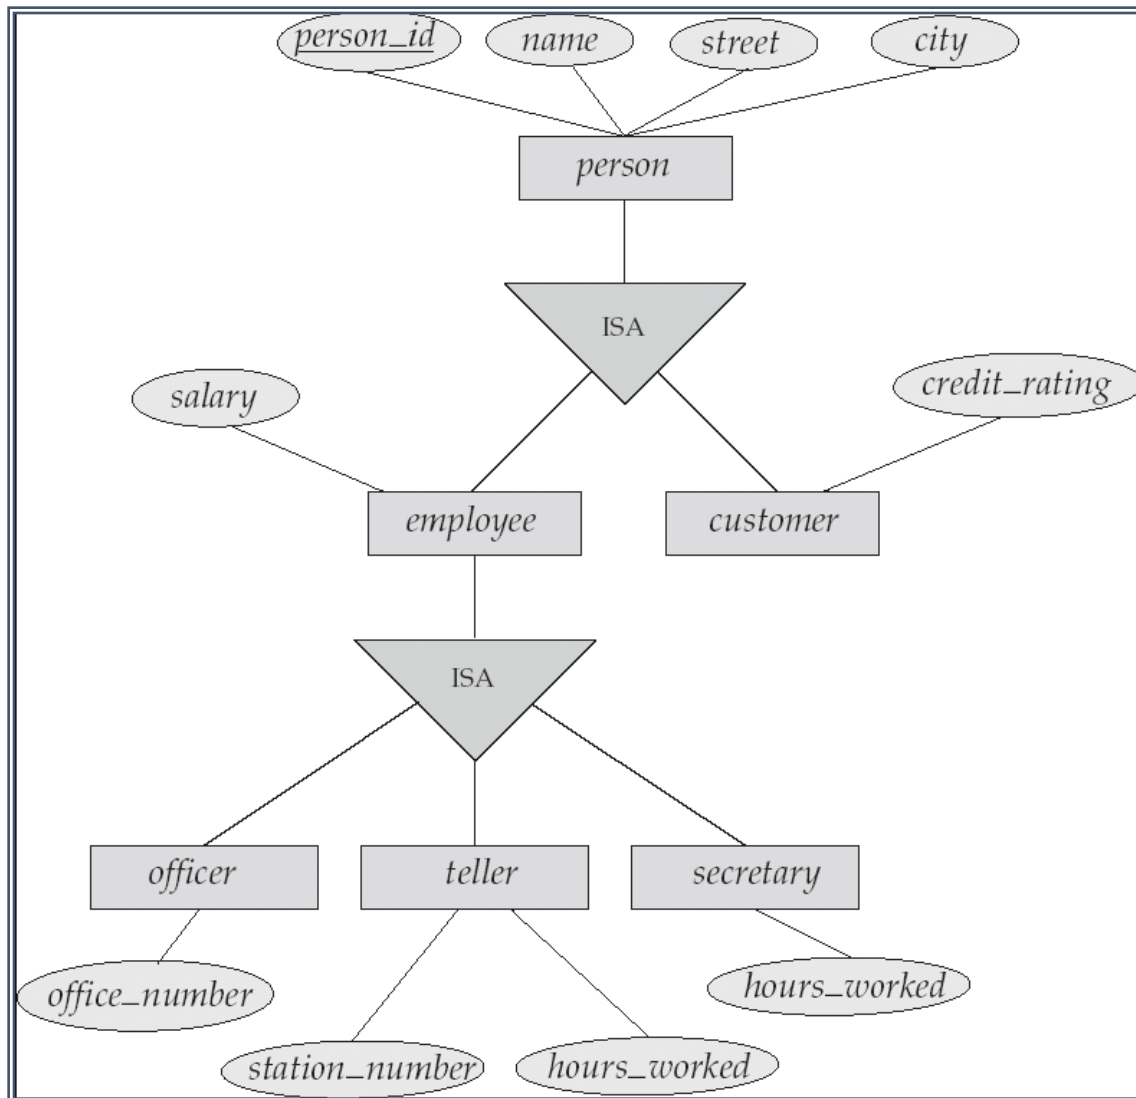
Design Issues

- N-ary vs binary relationships
 - Possible to avoid n-ary relationships, but there are some cases where it is advantageous to use them

Next: Specialization


- Consider entity person:
 - Attributes: name, street, city
- Further classification:
 - customer
 - Additional attributes: customer-id, credit-rating
 - employee
 - Additional attributes: employee-id, salary
- Note similarities to object-oriented programming

Example

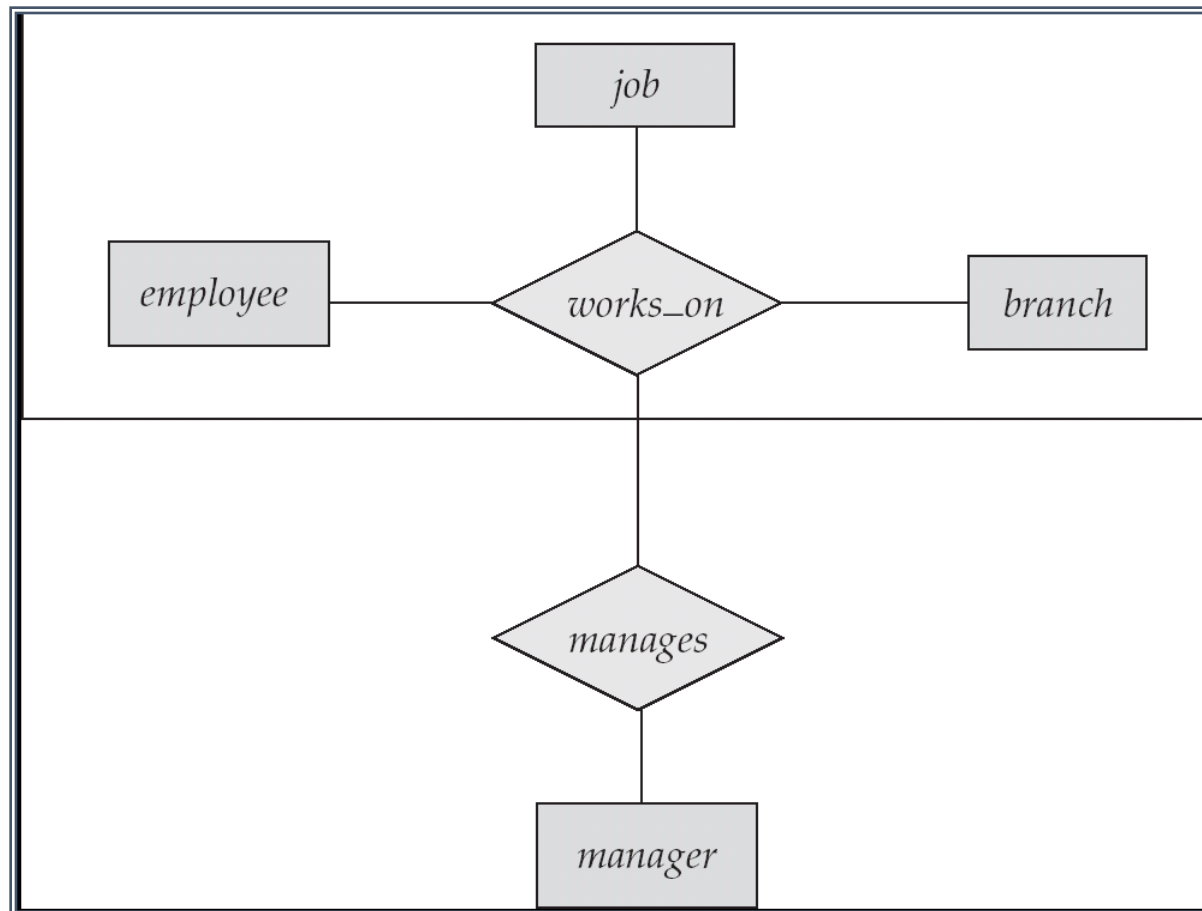


Aggregation

- Suppose we want to record managers for tasks performed by an employee at a branch



Aggregation (cont'd)



Break

- Entity-relationship Model
 - Intuitive diagram-based representation of domain knowledge, data properties etc...
 - Two key concepts:
 - Entities
 - Relationships
 - We also looked at:
 - Relationship cardinalities
 - Keys
 - Weak entity sets

Relational Data Model

Introduced by Ted Codd (late 60's – early 70's)

- *Before = “Network Data Model” (Cobol as DDL, DML)*
- *Very contentious: Database Wars (Charlie Bachman vs. Mike Stonebraker)*

Relational data model contributes:

1. *Separation of logical, physical data models (data independence)*
2. *Declarative query languages*
3. *Formal semantics*
4. *Query optimization (key to commercial success)*

1st prototypes:

- *Ingres → CA*
- *Postgres → Illustra → Informix → IBM*
- *System R → Oracle, DB2*

Key Abstraction: Relation

Account =

bname	acct_no	balance
Downtown	A-101	500
Brighton	A-201	900
Brighton	A-217	500

Terms:

- Tables (aka: Relations)

Why called Relations?

Why Called Relations?

Mathematical relations

Given sets: $R = \{1, 2, 3\}$, $S = \{3, 4\}$

- $R \times S = \{ (1, 3), (1, 4), (2, 3), (2, 4), (3, 3), (3, 4) \}$
- A **relation** on R, S is any subset (\subseteq) of $R \times S$
(e.g: $\{ (1, 4), (3, 4) \}$)

Database relations

Given attribute domains

Branches = $\{ \text{Downtown, Brighton, ...} \}$

Accounts = $\{ \text{A-101, A-201, A-217, ...} \}$

Balances = R

$\text{Account} \subseteq \text{Branches} \times \text{Accounts} \times \text{Balances}$

*$\{ (\text{Downtown}, \text{A-101}, 500),$
 $(\text{Brighton}, \text{A-201}, 900),$
 $(\text{Brighton}, \text{A-217}, 500) \}$*

Relations

Account =

bname	acct_no	balance
Downtown	A-101	500
Brighton	A-201	900
Brighton	A-217	500

Considered equivalent to...

*{ (Downtown, A-101, 500),
(Brighton, A-201, 900),
(Brighton, A-217, 500) }*

*Relational database semantics defined in
terms of mathematical relations*

So...

- That's the basic relational model
- That's it ?
 - What about semantic information ?
 - Relationships between entities ?
 - What about the constraints ?
 - How do we represent one-to-one vs many-to-one relationships ?
 - Those constraints are all embedded in the schema

Keys and Relations

- Recall:
 - Keys: Sets of attributes that allow us to identify entities
 - Very loosely speaking, tuples are viewed entities
- Just as in E/R Model:
 - Superkeys, candidate keys, and primary keys

Keys

- Superkey
 - set of attributes of table for which every row has distinct set of values
- Candidate key
 - Minimal such set of attributes
- Primary key
 - DB Chosen Candidate key
 - Plays a very important role
 - E.g. relations typically sorted by this

Keys

- Also act as integrity constraints
 - i.e., guard against illegal/invalid instance of given schema

e.g., Branch = (bname, bcity, assets)

bname	bcity	assets
Brighton	Brooklyn	5M
Brighton	Boston	3M



Invalid

Keys

- In fact, keys are one of the primary ways to enforce constraints/structure
- Consider a one-to-many relationship e.g.
 - Between customers and accounts
 - The relational model will be:
 - Customers(custid, custname,...)
 - Accounts(accountid, custid, balance,...)
 - Allows for multiple accounts per customer, but **not multiple customers per account**
 - Not possible to store such information (Key is accountid in Accounts table)
- In other words, constraints will lead to less representation power
 - Contrast with:
 - Customers(custid, custname,...)
 - Accounts(accountid, balance,...)
 - CustomerHasAccounts(custid, accountid)

More on Keys

- Determining Primary Keys
 - If relation schema derived from E-R diagrams, we can determine the primary keys using the original entity and relationship sets
 - Otherwise, same way we do it for E-R diagrams
 - Find candidate keys (minimal sets of attributes that can uniquely identify a tuple)
 - Designate one of them to be primary key
- Foreign Keys
 - If a relation schema includes the primary key of another relation schema, that attribute is called the foreign key

Schema Diagram for the Banking Enterprise

